New Upper Bounds for Dense Max 2-SAT by Accumulating Pruning Constraints

12 October 2026
First research draft

Abstract

We study exact algorithms for classes of Max 2-SAT instances whose clause density can grow with the number of variables. The objective is a running time $O^*(2^{(1-\gamma)n})$, for a constant $\gamma > 0$ independent of density, using polynomial space. We organize the analysis around comparisons between assignments: a branch needs to be explored only where it improves on an alternative that has already been covered. Such comparisons produce hard constraints whose accumulation can reduce the search space. We give three principal families: formulas with a persistent relative literal bias; independently signed clauses on any fixed almost-regular interaction graph; and formulas with repeated interactions across a large cut, including literal bicliques with a linear number of defective cross pairs. We also prove a compression bound for many bounded-size blocks that admit two assignments with identical, or nearly identical, outside effects. Conditional comparisons connect these results to bounded caching and generalized clause learning. We distinguish these proved mechanisms from the open structural question: whether failure of constraint accumulation necessarily exposes redundancy that can be exploited at comparable cost.

Status of this draft. This is a paper draft with an explicit research agenda, not a completed novelty or independent proof audit. The core statements below include proofs. The main presentation starts from literal gaps, random formulas and Max-Cut on regular graph families and the stochastic block model. Short combinatorial routines precede the spectral analysis; hidden communities additionally use explicit spectral preprocessing. General accumulation and discovery lemmas, technical membership estimates, and sharper special-case constants are placed in appendices. The final sections record proposed algorithms and open implications without treating them as established theorems. Numerical constants are conservative.

1 Introduction

Max 2-SAT asks for a Boolean assignment satisfying the largest number of clauses of a formula with at most two literals per clause. Exhaustive search takes $O^*(2^n)$ time and polynomial space. Our interest is in structural conditions that give a fixed exponential saving even when the number of binary clauses is superlinear in n.

The space requirement is central. Williams gave an algorithm with running time $O^*(2^{\omega n/3})$ and exponential space, where ω is the matrix multiplication exponent [1]. Thus the question here is not whether unrestricted-space exact algorithms can beat 2^n. Kulikov and Kutzkov obtained polynomial-space algorithms with a fixed exponential saving for every fixed clause density, using information from one branch to constrain another [2]. Degree-dependent algorithms and separator methods give further improvements [3]. A bound whose saving vanishes as the degree grows does not, by itself, give a fixed base below two on a family of growing density.

The central observation is elementary. If an alternative assignment is at least as good for a particular completion, there is no need to explore both assignments on that completion. The difficulty is to make enough such comparisons simultaneously useful, discover them at an affordable cost, and enumerate the surviving search space efficiently.

We pursue two ways in which comparisons become useful. First, comparisons can restrict different aspects of a completion, so their combined effect is strong even when each individual constraint is weak. Second, many assignments can have the same effect on the remaining variables. In this case, comparing the assignments directly can replace repeated search by a smaller collection of representatives. These mechanisms can coexist.

1.1 Input families and the common idea

We first present three kinds of input families, each with a direct description. Their precise qualifications are part of the corresponding theorems.

1. **Literal gaps.** A linear fraction of variables have a fixed relative literal bias and the degrees have bounded spread. The dense specialization includes a linear number of variables with linear gaps. Neighborhood correlations are arbitrary.

2. **Random formulas.** Independently random literal signs work on every fixed graph, including several clauses per pair and independently fixed units. Further models allow vanishing clause probabilities, parity bias, and units chosen after the binary clauses, above a logarithmic degree threshold.

3. **Max-Cut families.** Regular graphs with small nontrivial eigenvalues admit a combinatorial count-and-compare algorithm. In the stochastic block model, a fixed number of hidden communities and comparable edge variances give a fixed saving whenever $nv_{\min} \geq C \log n$. The community labels and edge probabilities need not be supplied.

The common operation is to compare a branch with the opposite value and accumulate hard constraints on the remaining assignment. Literal bias, random signs and regular-graph Max-Cut use combinatorial routines with explicit pseudocode. Spectral estimates belong to their analysis. Hidden communities also require spectral preprocessing in our current algorithm; when their partition and probabilities are supplied, group counts suffice.

We then study a complementary mechanism: repeated effects on the remaining variables permit compression into representatives. This yields bounds for literal bicliques, a linear number of defective cross pairs, repeated interaction profiles, and bounded blocks. These mechanisms can coexist; they are not an exhaustive classification. Compression across a cut does not subsume literal bias, and an exact literal biclique can coexist with almost balanced literal degrees; see Section 6.3. Section 3.8 gives a modeling interpretation of literal bias in terms of asymmetric pairwise conflicts, together with examples and the qualifications needed before claiming practical relevance.

1.2 Related work and scope

The comparison interpretation is closely related to the clause-learning approach of Kulikov and Kutzkov [2]; it should not be presented as a new general principle. Local maxima are also a classical object in exact Max 2-SAT algorithms. Hastings relates their number to exact algorithms and obtains further degree-dependent improvements under additional assumptions [4]. Formula caching has a theoretical literature, including the proof-complexity analysis of Beame, Impagliazzo, Pitassi and Segerlind [5]. Our question is which concrete dense classes admit density-independent savings from these mechanisms with polynomial space. A full comparison of the class statements, rather than only the algorithmic ideas, remains part of the publication audit.

Family	Assumptions to keep visible	Exponential bound
Relative literal bias	Canonical pair form; units allowed; fixed ε, η, R	$2^{(1-\varepsilon^2\eta^2/(3000R))n}$
Weighted random signs	Any fixed graph; bounded independent pair laws; units fixed independently	$(2e^{-c_0/2})^n$ with exponentially high probability
Regular-graph Max-Cut	Regular; $\lambda(G) \le \lambda_0\sqrt{d}$; $d \le (1-\varepsilon)n$; fields arbitrary	c^n, $c(\lambda_0, \varepsilon) < 2$
Stochastic block model	Fixed k, χ; comparable edge variances; $nv_{\min} \ge C \log n$; groups hidden	c^n, $c(k, \chi) < 2$, with probability $1 - o(1)$
Literal biclique	Sides of sizes s, t; exact prescribed cross clauses	$2^{n-\min(s,t)}$ when supplied
Bounded blocks	q disjoint blocks, size at most b; at most r exceptional outside variables each	$2^n(1-2^{-b})^{q/(2r+1)}$

Table 1: Core bounds proved in the draft. All suppress polynomial factors and use polynomial space. Structural detection costs are treated separately; the random-instance probabilities are over the input, not the algorithm.

2 Notation, simplification, and pruning constraints

Let F be an unweighted Max 2-SAT formula on a set V of n variables. In this section and in the literal-gap theorem, the input is *simple*: no clause occurs twice, and no variable occurs twice in a clause. We write $\mathbf{v}$ for a variable and $v, \neg v$ for its two literals. The numbers $d_1^F(v)$ and $d_2^F(v)$ count the unit and binary clauses containing v, respectively. We omit the superscript when the formula is clear. Unit clauses created by simplification retain their multiplicities. The number of binary clauses is m_2, and their density is $\rho = m_2/n$. The notation O^* suppresses factors polynomial in the input size.

A *partial assignment* α assigns values to a subset of V. A variable outside that subset is *unassigned*; a complete assignment extending α is a *completion* of α. For a complete assignment x, let $F(x)$ be its number of satisfied clauses, and let $\mathrm{OPT}(F) = \max_x F(x)$. Restricting F by α removes satisfied clauses and false literals. We record the removed satisfied clauses in a separate score s. Thus, for every completion, its original score equals s plus its score in the residual formula. An empty soft clause contributes zero and is discarded.

The input clauses are soft: an optimal assignment may leave some of them unsatisfied. A pruning constraint is an additional restriction used to reduce the search space. We write C for one such constraint and $\mathcal{C}$ for a collection of them. Our constraints have the form $\sum_i \ell_i \le b$, where every ℓ_i is a Boolean literal and repeated occurrences are counted separately.

2.1 Simplification rules and preprocessing

Lemma 2.1 (Unit clause cancellation)**.** *Removing one copy of* (v) *and one copy of* $(\neg v)$ *decreases the score of every complete assignment by exactly one.*

Proof. Exactly one of the two unit clauses is satisfied by every assignment. □

Lemma 2.2 (Simplifying two clauses on the same variable pair)**.** *Let* a, b *be literals on distinct variables. Replacing* $(a \vee b), (a \vee \neg b)$ *by the unit clause* (a) *decreases the score of every complete assignment by exactly one.*

Proof. If $a = 1$, the two original clauses contribute two and the new unit contributes one. If $a = 0$, exactly one original clause is satisfied and the new unit contributes zero. In both cases the difference is one. □

Lemma 2.3 (Unit clause domination)**.** *Suppose $d_1(\neg v) = 0$ and $d_1(v) \geq d_2(\neg v)$. For every fixed assignment of the other variables, setting $v = 1$ satisfies at least as many clauses as setting $v = 0$.*

Proof. Changing v from zero to one gains all $d_1(v)$ positive units. Clauses containing v cannot lose satisfaction. The only possible losses are among the $d_2(\neg v)$ binary clauses containing $\neg v$, since there are no units $(\neg v)$. The net gain is therefore at least $d_1(v) - d_2(\neg v) \geq 0$. □

In particular, a pure literal, whose opposite has no occurrences, can be set to one. When pruning constraints are present, such an assignment must also respect them. The procedure below uses domination only when the literal being set to one occurs in no current constraint.

Algorithm 1 PREPROCESS(F)

$s \leftarrow 0$
while F contains $(a \vee b), (a \vee \neg b)$ on distinct variables **do**
  Replace these clauses by (a) using Lemma 2.2.
  $s \leftarrow s + 1$
end while
for each variable **v do**
  $t \leftarrow \min\{d_1(v), d_1(\neg v)\}$
  Remove t copies of each opposite unit; set $s \leftarrow s + t$.
end for
return (F, s)

The score identities justify PREPROCESS for every assignment, not only for an optimum. After it terminates, a variable pair occurs in at most two binary clauses. If there are two, their literals are complementary in both positions: otherwise the two clauses would share a literal and the loop could apply again. We call this form *canonical*. Assigning variables creates only units, so the binary clauses do not need to be preprocessed again at later nodes.

2.2 Constructing and enforcing a pruning constraint

Lemma 2.4 (Pruning constraints from comparison)**.** *In a residual formula F, let v be either literal of an unassigned variable. List the partner literals in clauses containing v as $x_1, \ldots, x_{d_2(v)}$, and those in clauses containing $\neg v$ as $y_1, \ldots, y_{d_2(\neg v)}$. Define*

$$C: \quad \sum_{i=1}^{d_2(v)} \neg x_i + \sum_{j=1}^{d_2(\neg v)} y_j \leq d_2(\neg v) - d_1(v) + d_1(\neg v) - 1. \tag{1}$$

For every complete assignment of the other variables, C holds if and only if setting $v = 0$ has strictly larger score than setting $v = 1$.

Proof. Keep all other variables fixed. A unit (v) contributes one to the gain of $v = 1$ over $v = 0$, and a unit $(\neg v)$ contributes minus one. A clause $(v \vee x_i)$ contributes $1 - x_i = \neg x_i$ to this gain. A clause $(\neg v \vee y_j)$ contributes $-(1 - y_j) = y_j - 1$. All clauses not containing the variable of v have the same score in both assignments. Their contributions cancel. The exact gain is therefore

$$d_1(v) - d_1(\neg v) - d_2(\neg v) + \sum_i \neg x_i + \sum_j y_j.$$

Setting $v = 0$ is strictly better precisely when this integer is at most -1, which is the stated constraint. □

To restrict a constraint after an assignment, remove every false literal occurrence. Remove every true literal occurrence and decrease its bound by one. Also, a complementary pair $w, \neg w$ can be removed while decreasing the bound by one, since $w + \neg w = 1$. A negative bound proves infeasibility. After cancellation, if a literal occurs more often than the bound, it must be false in every satisfying completion.

Lemma 2.5 (Infeasibility by summation)**.** *Let $\mathcal{C}$ consist of constraints $C : \sum_{i=1}^{r_C} \ell_i \leq b_C$, with literal occurrences counted separately. Add the constraints, and repeatedly remove complementary literal pairs from the sum, decreasing its bound by one per pair. If the resulting bound is negative, no assignment satisfies all constraints in $\mathcal{C}$.*

Proof. An assignment satisfying every constraint also satisfies their sum. Cancelling a pair preserves the inequality because its contribution is exactly one for every assignment. The remaining left-hand side is nonnegative, so it cannot satisfy a negative bound. □

This is a sufficient infeasibility test. Passing it does not imply that the collection has a common satisfying assignment. The summed constraint is temporary; it is not added to $\mathcal{C}$.

In SIMPLIFY, the instruction *assign* $v = p$ records the corresponding variable value in α, updates the residual soft formula and score s, and substitutes the value in every constraint. All choices within preprocessing and simplification follow fixed orders.

Algorithm 2 SIMPLIFY$(F, \mathcal{C})$

$s \leftarrow 0$; $\alpha \leftarrow \emptyset$
while true **do**
  Cancel opposite units by Lemma 2.1, adding their constant to s.
  Cancel complementary literal pairs within each $C \in \mathcal{C}$.
  Let r_C be its remaining number of occurrences; remove C if $b_C \geq r_C$.
  if some $b_C < 0$, or Lemma 2.5 proves infeasibility **then**
    return INFEASIBLE
  end if
  if some literal v occurs more than b_C times in some $C \in \mathcal{C}$ **then**
    Assign $v = 0$, updating $F, \mathcal{C}, s, \alpha$.
  else if an unassigned literal v satisfies $d_1(\neg v) = 0$, $d_1(v) \geq d_2(\neg v)$, and occurs in no $C \in \mathcal{C}$ **then**
    Assign $v = 1$, updating $F, \mathcal{C}, s, \alpha$.
  else
    return $(F, \mathcal{C}, s, \alpha)$
  end if
end while

The last assignment rule preserves feasibility of the current constraints: setting $v = 1$ cannot increase any left-hand side, because only $\neg v$ can occur there. Lemma 2.3 shows that the soft score does not decrease. Thus this rule preserves the best score subject to the current constraints. An empty collection has sum $0 \leq 0$.

3 The literal gap algorithm and its analysis

3.1 Intuition

Suppose a linear number of variables have a substantial difference between the frequencies of their two literals. For the initial intuition, ignore unit clauses. If one literal occurs much more often than its negation, setting it to one satisfies more clauses on average over assignments to the other variables. The other value may still be better for particular assignments, so we cannot simply discard it. Instead, the pruning constraint of Lemma 2.4 describes exactly when that other value is strictly better. A large gap makes this a substantial requirement on the neighboring literals. Here a large gap means a constant fraction of the variable's degree; the degree itself need not be linear in n.

Consider a path in the branching tree on which each branching value is chosen with probability $1/2$. This random path is only a way of examining the search tree; the algorithm is deterministic. Whenever we branch on a variable that still has a substantial gap, the nonpreferred value adds a constraint whose bound is well below the expected number of true literal occurrences. For example, $\sum_{i=1}^{r} \ell_i \leq r/4$ allows only half the expected number $r/2$ of true occurrences under a uniform assignment. Each occurrence contributes either zero or one; its expectation is one half. Taking the nonpreferred value at many such steps therefore accumulates requirements that the remaining assignment must meet.

The constraints need not have disjoint neighborhoods, and may even be identical. The question is whether a few shared variables can satisfy many requirements at once. Each branching variable creates at most one constraint, using its still-present binary clauses. Consequently, a fixed variable can occur across the accumulated constraints at most as many times as it occurs in binary clauses after preprocessing. Small degrees limit overlap in this way. Larger degrees allow more overlap, but a relative gap at a branching variable also gives a stronger constraint. The theorem controls the concentration of these occurrence counts so that a small set of variables cannot account for all the accumulated requirements.

This suggests the source of the savings: on most paths, enough restrictive constraints accumulate for their sum to become infeasible while many variables remain unassigned. The infeasibility test of Lemma 2.5 then rejects the whole remaining subtree. Simplification provides a second source of savings: when it assigns a variable without branching, only one value needs to be explored. Both effects are useful. The analysis must also account for gaps changing as neighboring variables are assigned; it does not assume that every initial gap persists along every path.

3.2 The algorithm

The algorithm first applies Preprocess, then searches recursively. At each node, Simplify applies reduction rules and checks the accumulated constraints. If branching is still necessary, the algorithm explores both values of the next variable and adds a pruning constraint only to its nonpreferred child. The procedures from Section 2 perform the simplification and propagation; we only need to specify the variable order and the preferred value.

To include unit clauses in the preference calculation, define the *effective gap* of a variable $\mathbf{v}$ in a residual formula H by

$$g_{\mathbf{v}}(H) = d_2^H(v) - d_2^H(\neg v) + 2\big(d_1^H(v) - d_1^H(\neg v)\big). \tag{2}$$

Without units this is the difference of the two literal frequencies. A unit contributes its full preference irrespective of the other variables, whereas a binary clause contributes its preference only when its partner literal is false. Under uniformly assigned remaining variables, the latter event has probability one half. Thus $g_{\mathbf{v}}(H)/2$ is the expected score gain from setting $v = 1$ rather than $v = 0$.

Fix an order on the input variables. At every branching node, choose the first unassigned variable **u** in this order. Its *preferred literal* is

$$v = \begin{cases} u, & d_2(u) + 2d_1(u) \geq d_2(\neg u) + 2d_1(\neg u), \\ \neg u, & \text{otherwise.} \end{cases} \tag{3}$$

All counts here are taken in the current formula, after SIMPLIFY. Setting $v = 1$ is the preferred choice; ties prefer the positive literal. The algorithm calculates this preference directly from clause counts, without sampling assignments. The fixed order does not have to put variables with large gaps first.

The global variables M and $\alpha^\star$ store the best complete score and a corresponding assignment. The parameters $F, \mathcal{C}, \alpha, s$ belong to the current node: its residual formula, pruning constraints, partial assignment, and accumulated satisfied-clause contribution. Variables remain in the unassigned set even if they disappear from F.

Algorithm 3 LITERALGAP(F)

Fix an order on the input variables.
$M \leftarrow -\infty$; $\alpha^\star \leftarrow \emptyset$ *// Global best score and complete assignment.*
$(F, s) \leftarrow$ PREPROCESS(F)
GAPSEARCH($F, \emptyset, \emptyset, s$)
return $(M, \alpha^\star)$

Algorithm 4 GAPSEARCH($F, \mathcal{C}, \alpha, s$)

$R \leftarrow$ SIMPLIFY($F, \mathcal{C}$)
if $R =$ INFEASIBLE **then**
  return
end if
$(F, \mathcal{C}, t, \beta) \leftarrow R$; $s \leftarrow s+t$; $\alpha \leftarrow \alpha \cup \beta$ *// Include the assignments and score obtained by simplification.*
if all input variables are assigned **then**
  If $s > M$, set $M \leftarrow s$ and $\alpha^\star \leftarrow \alpha$.
  return
end if
Let **u** be the first unassigned variable in the fixed order.
Choose its preferred literal v by (3).
Construct C from F and v by Lemma 2.4.
for $p \in [1, 0]$ **do**
  Make separate copies $(F_p, \mathcal{C}_p, \alpha_p, s_p)$ of the current state.
  if $p = 0$ **then**
    $\mathcal{C}_p \leftarrow \mathcal{C}_p \cup \{C\}$
  end if
  Assign $v = p$, updating the copied formula, constraints, assignment, and score.
  GAPSEARCH($F_p, \mathcal{C}_p, \alpha_p, s_p$)
end for

The algorithm does not need to identify a set of variables with large gaps or know the parameters in the running-time theorem. It applies the same rules to every simple input formula. The theorem's assumptions determine when these rules give an exponential saving.

3.3 Correctness

A single pruning constraint follows from comparing the two values of one variable. We must also show that many such comparisons, together with simplification, preserve at least one optimum. The following lemma does this by choosing the earliest optimum in a fixed ordering of the full decision tree and showing that none of the rules removes its path. This establishes correctness independently of the assumptions used to bound the running time.

Lemma 3.1 (Correctness of the literal gap algorithm)**.** *For every simple input formula,* LiteralGap *returns its optimum score and an assignment attaining that score.*

Proof. Preprocessing and unit cancellation preserve every assignment's score when their recorded constants are included. It remains to show that the search retains an optimum despite the pruning constraints and assignments made by Simplify.

For this purpose, represent each assignment made by Simplify as a decision with one retained child. Add the other child temporarily, and also restore subtrees rejected by infeasibility tests. Complete these restored subtrees by arbitrary binary decisions until every variable is assigned. This gives a full decision tree whose leaves are all complete assignments. At an ordinary branching node, place the preferred child first. At an assignment made by Simplify, place the retained value first. This defines an order on all complete assignments: list the leaves from left to right. Choose the earliest optimum x^* in this order.

We show that the actual algorithm retains the path of x^*. Suppose it has reached a node on that path. At an ordinary branch where x^* takes the nonpreferred value $v = 0$, change only this variable to $v = 1$. The resulting assignment lies in the earlier child. Its score must be strictly smaller than that of x^*: a larger score contradicts optimality, and an equal score contradicts the choice of the earliest optimum. By Lemma 2.4, x^* therefore satisfies the newly added constraint. Inductively, it satisfies every constraint generated on its path.

Consequently, neither an individual constraint nor their sum can certify infeasibility on this path. An assignment forced by a constraint agrees with x^*. If a domination rule instead assigns $v = 1$ and x^* had $v = 0$, Lemma 2.3 would give an assignment with at least the same score in the earlier child, again contradicting the choice of x^*. Removing an already satisfied constraint and cancelling complementary literals preserve satisfaction. Thus every operation retains this optimum.

The algorithm eventually visits its complete assignment and records its original score. Every other score recorded also belongs to a complete assignment of the original formula, so none exceeds the optimum. □

3.4 Running-time guarantee

The following theorem makes the two requirements from the intuition precise: many variables have a gap proportional to their binary degree, and the total occurrence counts are not too concentrated on a few variables. The first supplies strong pruning constraints. The second bounds how much shared variables can affect their sum and how much bias can remain in a small unassigned set. All these conditions concern the formula immediately after preprocessing.

Theorem 3.2 (Relative literal bias)**.** *Fix constants $\varepsilon, \eta \in (0, 1]$ and $R \geq 1$. Let F be a simple unweighted Max 2-SAT formula on a set V of n variables, and let $(F_0, s_0) = \textsc{Preprocess}(F)$. For each variable $\mathbf{v}$, define*

$$D_{\mathbf{v}} = d_2^{F_0}(v) + d_2^{F_0}(\neg v), \qquad u_{\mathbf{v}} = d_1^{F_0}(v) + d_1^{F_0}(\neg v),$$

$$g_{\mathbf{v}} = d_2^{F_0}(v) - d_2^{F_0}(\neg v) + 2\big(d_1^{F_0}(v) - d_1^{F_0}(\neg v)\big).$$

Suppose there is a set $S \subseteq V$ with $|S| \geq \eta n$ such that $|g_{\mathbf{v}}| \geq \varepsilon D_{\mathbf{v}}$ for every $\mathbf{v} \in S$. Writing $d = |S|^{-1} \sum_{\mathbf{v} \in S} D_{\mathbf{v}}$, suppose also that

$$d \geq 1, \qquad \sum_{\mathbf{v} \in V} (D_{\mathbf{v}} + u_{\mathbf{v}})^2 \leq Rnd^2. \tag{4}$$

Then an optimal assignment for F can be computed deterministically in polynomial space and time $O^*(2^{(1-\gamma)n})$, where $\gamma = \varepsilon^2 \eta^2 / (3000R)$.

The average degree d may grow with n. For fixed ε, η, R, the saving γ remains a positive constant independent of d. Condition (4) allows unequal degrees; it controls their squared sum rather than imposing a common degree or requiring disjoint neighborhoods.

3.5 Proof of the running-time bound

The proof measures the fraction of paths that survive to a depth at which a small constant fraction of the variables is still unassigned. A path includes assignments made by Simplify, as well as ordinary branching decisions. The argument proceeds as follows.

1. Paths with many assignments made by simplification are already few: at each such step only one of the two values is retained. This is the role of Lemma 3.3 in Step 2 below.

2. On the other paths, a substantial part of the initial total gap must be encountered at ordinary branching steps. A small set of variables cannot contain most of it, and random changes caused by assigning neighbors are unlikely to remove it. Step 3 combines the degree bound with Lemmas 3.4 and 3.7.

3. These branching steps generate constraints whose combined requirements are unlikely to be met by the assigned values and the few remaining variables. Step 4 uses the exact sum identity of Lemma 3.5 and the remaining estimates of Lemma 3.7 to prove rejection by the sum test.

4. An exponentially small surviving fraction at the chosen depth gives the search-tree bound in Step 5. Shallower levels have few nodes because they are still a linear distance from the leaves.

The main proof gives these steps and their quantitative bounds. The supporting lemmas, including the martingale calculations, follow it.

Proof of Theorem 3.2. We prove the bound by analyzing LiteralGap (Algorithms 3 and 4). Its correctness follows from Lemma 3.1. Depth-first search stores a polynomial-size state at each of at most n assignment levels, so its space use is polynomial. Each call to Simplify performs at most n assignments and polynomial work between assignments. We therefore need to bound the size of the search tree.

1. Choose an analysis depth. Put $\theta = \varepsilon\eta$, $B_{\mathbf{v}} = D_{\mathbf{v}} + u_{\mathbf{v}}$, and

$$\delta = \frac{\theta^2}{256R}, \qquad T = \lceil (1-\delta)n \rceil.$$

Expose assignments one variable at a time, including those made within Simplify. Call the latter *forced steps*; call decisions with two recursive children *branching steps*. Restore missing children and pruned subtrees as in the correctness proof, obtaining a full binary tree. Follow a random path in this tree, choosing each value by an independent fair coin. After a path is pruned, continue it arbitrarily by selecting unassigned variables, using the same preferred-value and constraint-generation rules but ignoring rejection tests. This continuation is used only in the analysis.

Let π_t be the probability that the random path is retained through t assignments, including all infeasibility checks at that level. Each length-t path in the full tree has probability 2^{-t}, so the number of retained states at that level is $2^t\pi_t$. Rejected children and the work inside simplification contribute only polynomial factors. It suffices to bound $\max_{0\le t\le n} 2^t\pi_t$.

2. Set aside paths with many forced steps. Let R_T be the number of forced steps among the first T assignments. By Lemma 3.3,

$$\mathbb{P}(\text{the path survives to } T \text{ and } R_T > \delta n) \le 2^{-\delta n}.$$

On the remaining paths, at most δn variables have been assigned by simplification. We show that ordinary branching then creates enough constraints to reject almost all such paths by depth T.

3. Account for the initial gaps. We first bound how much gap can be confined to a small set of variables. We will use this bound both for variables still unassigned at depth T and for those assigned by simplification. For every set $W \subseteq V$ of size at most δn, the Cauchy–Schwarz inequality gives

$$\begin{aligned}\sum_{\mathbf{v}\in W} B_{\mathbf{v}} &\le \left(\sum_{\mathbf{v}\in W} 1^2\right)^{1/2}\left(\sum_{\mathbf{v}\in W} B_{\mathbf{v}}^2\right)^{1/2}\\ &\le \sqrt{\delta n}\sqrt{Rnd^2} = \sqrt{\delta R}\, nd = \frac{\theta nd}{16}.\end{aligned}$$

The total initial absolute gap of variables in S is at least

$$\Phi_0 := \sum_{\mathbf{v}\in S} |g_{\mathbf{v}}| \ge \varepsilon \sum_{\mathbf{v}\in S} D_{\mathbf{v}} = \varepsilon |S| d \ge \theta nd.$$

Let Φ_T be their total absolute gap among the variables still unassigned at depth T. There are at most δn unassigned variables. Lemma 3.4 bounds every residual absolute gap by $2B_{\mathbf{v}}$, so $\Phi_T \le \theta nd/8$. The same argument bounds the sum of the gaps of variables in S removed at forced steps by $\theta nd/8$ whenever $R_T \le \delta n$.

We must also account for the changes in a variable's gap before it is selected. Let H_T sum the absolute gaps of variables in S at the moments they are assigned, including forced steps. Let G_T be the corresponding sum over branching steps only. The first estimate of Lemma 3.7, with $a = \theta nd/4$, gives

$$\mathbb{P}\big(H_T < \Phi_0 - \Phi_T - \theta nd/4\big) \le e^{-\theta^2 n/(128R)}.$$

Outside this exceptional event, on every path with $R_T \le \delta n$,

$$G_T \ge \theta nd - \frac{\theta nd}{8} - \frac{\theta nd}{4} - \frac{\theta nd}{8} = \frac{\theta nd}{2}.$$

4. Show that the summed constraints are infeasible. We separate the total requirement at constraint creation, the effect of values assigned later, and the best contribution still available from unassigned variables. This expresses precisely what the sum test must compare. For a constraint C, let r_C^0 and b_C^0 refer to its form when it is created. Its initial restriction relative to the expected number of true literal occurrences is $r_C^0/2 - b_C^0$. Let A_T be the sum of these quantities over constraints created by depth T. Let Q_T be the total number of true occurrences subsequently assigned in these constraints, minus one half the number of occurrences assigned. Finally, let

$$U_T = \frac{1}{2} \sum_{\mathbf{w}\ \text{unassigned}} |m_{\mathbf{w}}^+ - m_{\mathbf{w}}^-|,$$

where $m^+_{\mathbf{w}}$ and $m^-_{\mathbf{w}}$ count the occurrences of w and $\neg w$ across those constraints. These quantities count occurrences, not distinct literals.

Lemma 3.5 shows that the minimum left-hand side minus right-hand side of their sum is

$$Z_T = A_T + Q_T - U_T. \tag{5}$$

In particular, a surviving path must satisfy $Z_T \le 0$. The occurrence bound $m^+_{\mathbf{w}} + m^-_{\mathbf{w}} \le D_{\mathbf{w}}$ and the Cauchy–Schwarz estimate above give

$$U_T \le \frac{1}{2} \sum_{\mathbf{w} \text{ unassigned}} D_{\mathbf{w}} \le \frac{1}{2} \sum_{\mathbf{w} \text{ unassigned}} B_{\mathbf{w}} \le \frac{\theta n d}{32}.$$

The remaining estimates of Lemma 3.7, with $b = c = \theta n d/32$, show that

$$A_T \ge G_T/4 - \theta n d/32, \qquad Q_T \ge -\theta n d/32$$

except on events with probabilities at most $e^{-\theta^2 n/(4096R)}$ and $e^{-\theta^2 n/(512R)}$, respectively. If none of the three exceptional events occurs and $R_T \le \delta n$, then

$$Z_T \ge \frac{\theta n d}{8} - \frac{\theta n d}{32} - \frac{\theta n d}{32} - \frac{\theta n d}{32} = \frac{\theta n d}{32} > 0.$$

The sum test therefore rejects this path.

5. Convert the survival probability to a time bound. Combining the forced-step bound and the three exceptional probabilities,

$$\pi_T \le 2^{-\delta n} + 3e^{-\theta^2 n/(4096R)} \le 4e^{-\theta^2 n/(4096R)}.$$

For $t \le T$, $2^t \pi_t \le 2^T$. For $t \ge T$, survival to depth t requires survival to depth T, so $2^t \pi_t \le 2^n \pi_T$. Consequently,

$$\max_t 2^t \pi_t \le \max\left\{2^{(1-\delta)n+1}, 4 \cdot 2^{(1-\theta^2/(4096R\ln 2))n}\right\}.$$

Both exponents give at least the saving $\theta^2/(3000R)$, since $1/256 > 1/3000$ and $1/(4096 \ln 2) > 1/3000$. Polynomial work per state completes the proof. □

3.6 Supporting lemmas for the random-path analysis

Throughout this subsection, use the full random decision tree and notation from the proof. Let $\mathcal{U}_t$ be the unassigned variables after t steps, and let $\mathcal{H}_t$ denote the information revealed by those steps. The next variable and whether the step is forced are determined before its value is revealed. Write $\mathbf{x}_t$ for the variable assigned at step t. Let $g_{\mathbf{v}}(t)$ be the effective gap in the residual soft formula after t assignments. Opposite-unit cancellation does not change this gap.

The first estimate handles paths with many assignments made by Simplify. Each such step requires the random path to follow a specified value, even though the position and the required value of the step may depend on the earlier choices. This is the estimate used to set those paths aside in Step 2 of the main proof.

Lemma 3.3 (Forced assignments on a random path)**.** *For every real $r \ge 0$, the probability that the path survives to depth T and has more than r forced steps is at most 2^{-r}.*

Proof. Let I_t indicate that all forced assignments encountered up to step t were followed, and let R_t count those steps. Put $W_t = 2^{R_t} I_t$. Initially $W_0 = 1$. At a branching step, $W_t = W_{t-1}$. At a forced step with $I_{t-1} = 1$, one fair choice doubles W_{t-1} and the other gives zero. If $I_{t-1} = 0$, both give zero. In all cases $\mathbb{E}[W_t \mid \mathcal{H}_{t-1}] = W_{t-1}$, and thus $\mathbb{E}W_T = 1$. Survival with $R_T > r$ implies $I_T = 1$ and $W_T > 2^r$. Markov's inequality gives the claimed probability bound. □

For the remaining paths we need to relate the initial gaps to the gaps seen when variables are selected. The next lemma supplies two elementary facts for Step 3: an occurrence bound on every residual gap, and a description of how assigning one neighbor changes it. The second fact will also control gap fluctuations in Lemma 3.7.

Lemma 3.4 (Changes in the effective gap)**.** *For every unassigned variable* $\mathbf{v}$*,* $|g_{\mathbf{v}}(t)| \leq 2(D_{\mathbf{v}} + u_{\mathbf{v}})$*. Let* $c_{\mathbf{xv}}$ *be the number of binary clauses on the pair* $\{\mathbf{x}, \mathbf{v}\}$ *in* F_0*. When an unassigned neighbor* $\mathbf{x}$ *is assigned a fair value, the change in* $g_{\mathbf{v}}$ *is* $+c_{\mathbf{xv}}$ *or* $-c_{\mathbf{xv}}$ *with equal probability.*

Proof. Every binary clause incident to $\mathbf{v}$ either remains binary, disappears, or becomes a unit. Its absolute contribution to the effective gap is therefore at most two. Each original uncancelled unit also contributes at most two. This proves the first assertion.

For a clause $(v \vee x)$, the contribution to $g_{\mathbf{v}}$ changes from one to zero when $x = 1$, and from one to two when $x = 0$. Thus the two changes are minus one and plus one. Replacing either literal by its negation reverses the corresponding signs. If two clauses survive on the same pair after Preprocess, they are complementary in both positions, and their changes have the same sign for either assignment of $\mathbf{x}$. They therefore add to $\pm c_{\mathbf{xv}}$. Cancelling opposite units changes neither effective gap. □

Step 4 needs a numerical condition that guarantees rejection by the algorithm's sum test. The following identity separates the requirements introduced by the constraints from the contributions of assigned and unassigned variables. It also justifies retaining all created constraints in the analysis, even when Simplify has already discarded some as satisfied.

Lemma 3.5 (The minimum of the summed constraints)**.** *With* A_T, Q_T, U_T *as defined in the main proof, the minimum difference between the left- and right-hand sides of the sum is* $A_T + Q_T - U_T$*. For each constraint* C*, let* g_C *be the absolute effective gap of its branching variable when the constraint is created. Then*

$$r_C^0/2 - b_C^0 = g_C/2 + 1, \qquad C : \sum_{i=1}^{r_C^0} \ell_i \leq (r_C^0 - g_C)/2 - 1, \tag{6}$$

Also $m_{\mathbf{w}}^+ + m_{\mathbf{w}}^- \leq D_{\mathbf{w}}$*. These facts imply that survival to depth* T *requires* $A_T + Q_T - U_T \leq 0$*, even when* Simplify *has removed already satisfied constraints.*

Proof. At creation, $r_C^0 = d_2(v) + d_2(\neg v)$, and Lemma 2.4 gives $b_C^0 = d_2(\neg v) - d_1(v) + d_1(\neg v) - 1$. Subtraction yields

$$r_C^0/2 - b_C^0 = \frac{d_2(v) - d_2(\neg v) + 2(d_1(v) - d_1(\neg v))}{2} + 1.$$

The expression in the numerator is nonnegative for the preferred literal and equals the absolute effective gap of its variable.

For the analysis, keep a copy of every constraint in its creation form. Let a_C be the number of its occurrences assigned by depth T, and let j_C be the number of those occurrences made true. Its remaining bound is $b_C^0 - j_C$. Summing over constraints, the remaining left-hand side is

$$\sum_{\mathbf{w} \in \mathcal{U}_T} (m_{\mathbf{w}}^+ w + m_{\mathbf{w}}^- \neg w).$$

For each variable its minimum contribution is $\min(m_{\mathbf{w}}^+, m_{\mathbf{w}}^-)$, obtained by choosing the better of its two values. Since $\min(p, q) = (p + q - |p - q|)/2$ and $\sum_{\mathbf{w}}(m_{\mathbf{w}}^+ + m_{\mathbf{w}}^-) = \sum_C (r_C^0 - a_C)$,

the minimum of the sum minus its bound is

$$\sum_{\mathbf{w}} \min(m^+_{\mathbf{w}}, m^-_{\mathbf{w}}) - \sum_C (b^0_C - j_C) = \sum_C (r^0_C/2 - b^0_C) + \sum_C (j_C - a_C/2) - \tfrac{1}{2}\sum_{\mathbf{w}} |m^+_{\mathbf{w}} - m^-_{\mathbf{w}}|$$
$$= A_T + Q_T - U_T.$$

Cancellation in Lemma 2.5 computes this same minimum.

Each other variable creates at most one constraint, containing exactly the still-present binary occurrences incident to it. Thus a fixed $\mathbf{w}$ contributes at most $\sum_{\mathbf{v}} c_{\mathbf{vw}} = D_{\mathbf{w}}$ occurrences across the constraints. This proves the occurrence bound without any assumption that neighborhoods are disjoint.

Finally, a constraint discarded by SIMPLIFY has nonpositive left-hand side minus right-hand side for every later completion of the current partial assignment. Removing it can only increase the value of the remaining summed expression, pointwise. A positive minimum for the sum of all the recorded constraints therefore implies a positive minimum for the live sum tested by the algorithm. Such a node is rejected. □

The probabilistic estimates must allow each constraint to depend on earlier branching choices. The next lemma bounds cumulative deviations without requiring independence between steps: it uses conditional mean zero and bounds on the size of each change. We apply it in Lemma 3.7; the main proof only needs the three resulting estimates.

Lemma 3.6 (A bounded-increment martingale estimate)**.** *Let $M_0 = 0, M_1, \dots, M_T$ be random variables, with M_t determined by the information $\mathcal{H}_t$ revealed through step t, such that $\mathbb{E}[M_t \mid \mathcal{H}_{t-1}] = M_{t-1}$ for every t. Suppose $|M_t - M_{t-1}| \le b_t$, where b_t is known from $\mathcal{H}_{t-1}$, and $\sum_{t=1}^T b_t^2 \le B$ on every path for a constant $B > 0$. Then, for every $a > 0$,*

$$\mathbb{P}(M_T \le -a) \le \exp(-a^2/(2B)).$$

Proof. The displayed conditional-expectation identity is the martingale property: the expected next value, given everything revealed so far, is the current value. Hence $Y_t = M_t - M_{t-1}$ has conditional mean zero. Convexity of the exponential on $[-b_t, b_t]$ gives, for $\lambda > 0$,

$$\mathbb{E}[e^{-\lambda Y_t} \mid \mathcal{H}_{t-1}] \le \tfrac{1}{2}(e^{\lambda b_t} + e^{-\lambda b_t}) \le e^{\lambda^2 b_t^2/2}.$$

For the first inequality, bound the exponential by the line joining its endpoint values and use $\mathbb{E}[Y_t \mid \mathcal{H}_{t-1}] = 0$. The second follows by comparing the power series of $\cosh z$ and $e^{z^2/2}$, using $(2k)! \ge 2^k k!$. It follows by successive conditional expectations that

$$\mathbb{E}\exp\left(-\lambda M_T - \frac{\lambda^2}{2}\sum_{t=1}^T b_t^2\right) \le 1.$$

On $M_T \le -a$, the random variable inside this expectation is at least $\exp(\lambda a - \lambda^2 B/2)$. Markov's inequality bounds the probability by $\exp(-\lambda a + \lambda^2 B/2)$. Taking $\lambda = a/B$ proves the claim. □

There are three ways the accumulated constraints could fail to give the required pruning. Too much initial gap could disappear before its variable is selected; branching could introduce too little restriction relative to the gaps it encounters; or later assigned values could be unusually favorable to the constraints. The following estimates bound these three events in that order. Step 3 uses the first estimate, and Step 4 uses the other two together with Lemma 3.5.

Lemma 3.7 (Three estimates for a random path)**.** *Assume $d \geq 1$ and $\sum_{\mathbf{v}} B_{\mathbf{v}}^2 \leq Rnd^2$. For every $a, b, c > 0$, the quantities in the main proof satisfy*

$$\mathbb{P}(H_T < \Phi_0 - \Phi_T - a) \leq e^{-a^2/(8Rnd^2)},$$
$$\mathbb{P}(A_T < G_T/4 - b) \leq e^{-b^2/(4Rnd^2)},$$
$$\mathbb{P}(Q_T < -c) \leq e^{-2c^2/(Rnd^2)}.$$

Proof. **The total gap encountered.** At step t, let $h_t = |g_{\mathbf{x}_t}(t-1)|$ if $\mathbf{x}_t \in S$ and $h_t = 0$ otherwise. Removing the selected variable reduces the potential Φ by h_t. Let Y_t be the change in the total absolute gap of the other unassigned variables of S caused by assigning its value. Thus

$$\Phi_t = \Phi_{t-1} - h_t + Y_t.$$

For a remaining variable with current signed gap g and pair multiplicity q, Lemma 3.4 shows that its expected new absolute gap is $(|g+q|+|g-q|)/2 \geq |g|$, by the triangle inequality. Adding these expectations does not require the changes for different neighbors to be independent. Consequently $z_t := \mathbb{E}[Y_t \mid \mathcal{H}_{t-1}] \geq 0$.

Define

$$M_t = \sum_{j=1}^{t} (Y_j - z_j), \qquad M_0 = 0.$$

The quantities z_t are determined before the next value is revealed, and $\mathbb{E}[Y_t - z_t \mid \mathcal{H}_{t-1}] = 0$. Therefore $\mathbb{E}[M_t \mid \mathcal{H}_{t-1}] = M_{t-1}$, which proves that M is a martingale. Each absolute gap changes by at most its pair multiplicity, so $|Y_t| \leq \sum_{\mathbf{v}} c_{\mathbf{x}_t \mathbf{v}} = D_{\mathbf{x}_t}$ and $|Y_t - z_t| \leq 2D_{\mathbf{x}_t}$. Each variable is assigned once; hence the sum of the squared increment bounds is at most $4\sum_{\mathbf{v}} D_{\mathbf{v}}^2 \leq 4Rnd^2$. Summing the potential identity gives

$$H_T = \Phi_0 - \Phi_T + \sum_{t=1}^{T} z_t + M_T.$$

Since $z_t \geq 0$, the first claimed event implies $M_T < -a$. Lemma 3.6 proves the first estimate.

The restrictions introduced at branching steps. At a branching step let $k_t = |g_{\mathbf{x}_t}(t-1)|/2 + 1$. The increment X_t of A_t is k_t in the nonpreferred child and zero in the preferred child. Each has conditional probability one half, so $\mathbb{E}[X_t \mid \mathcal{H}_{t-1}] = k_t/2$. At a forced step put $X_t = 0$ and its conditional mean equal to zero. The centered sum of these increments is a martingale. At a branching step, its absolute increment is $k_t/2$, and

$$(k_t/2)^2 \leq (B_{\mathbf{x}_t} + 1)^2/4 \leq (B_{\mathbf{x}_t}^2 + 1)/2 \leq B_{\mathbf{x}_t}^2 + 1.$$

The sum of these squared bounds is at most $\sum_{\mathbf{v}}(B_{\mathbf{v}}^2 + 1) \leq Rnd^2 + n \leq 2Rnd^2$. The sum of the conditional means is at least $G_T/4$, because G_T counts only the branching steps whose variables belong to S. Thus $A_T < G_T/4 - b$ implies a centered deviation below $-b$. Lemma 3.6 gives the second estimate.

The effect of subsequent assignments. Immediately before assigning $\mathbf{x}_t$, let p and q be the numbers of positive and negative occurrences of this variable in previously created constraints. Its contribution to $Q_t - Q_{t-1}$ is $(p-q)/2$ when its value is one, and $(q-p)/2$ when its value is zero. These outcomes have equal conditional probabilities, so the conditional mean is zero. A newly created constraint contains only variables still unassigned and makes no additional contribution to Q_t at this step. Hence Q is a martingale starting at zero. Its increment is at most $(p+q)/2 \leq D_{\mathbf{x}_t}/2$ in absolute value, by Lemma 3.5. The squared bounds sum to at most $Rnd^2/4$. Applying Lemma 3.6 proves the third estimate. □

3.7 Dense gaps and the scope of the analysis

When the initial gaps are linear in n, the same analysis gives a simpler input condition. Simplicity already bounds each variable's total number of occurrences by $O(n)$, so no additional assumption on the degree sequence is needed. The following corollary states this condition directly for the original formula; its proof explains why preprocessing preserves the quantities needed by the argument.

Corollary 3.8 (Linear effective gaps in a simple formula)**.** *Fix $\varepsilon, \eta \in (0, 1]$. Let F be a simple unweighted Max 2-SAT formula on n variables. Suppose at least ηn variables satisfy*

$$|d_2^F(v) - d_2^F(\neg v) + 2(d_1^F(v) - d_1^F(\neg v))| \geq \varepsilon n.$$

Then an optimal assignment for F can be computed deterministically in polynomial space and time $O^*(2^{(1-\varepsilon^2\eta^2/48000)n})$.

Proof. Before preprocessing, each variable has at most $4(n-1)$ binary occurrences and at most two unit occurrences. The replacement in PREPROCESS preserves its effective gap: at the shared literal, two signed binary occurrences are replaced by a unit with coefficient two; at the other variable, opposite signed occurrences cancel. Unit cancellation also preserves the gap. Neither operation increases the sum of binary and unit occurrences at any variable. Thus, in F_0,

$$B_{\mathbf{v}} = D_{\mathbf{v}} + u_{\mathbf{v}} \leq 4(n-1) + 2 \leq 4n, \qquad \sum_{\mathbf{v}} B_{\mathbf{v}}^2 \leq 16n^3.$$

For the set S of biased variables, the initial total absolute gap is at least $\varepsilon\eta n^2$.

The preceding proof uses the relative-bias hypotheses only through the two inequalities $\Phi_0 \geq \theta n d$ and $\sum_{\mathbf{v}} B_{\mathbf{v}}^2 \leq R n d^2$, with $d \geq 1$. Use that proof with scale $d = n$, $\theta = \varepsilon\eta$, and $R = 16$. Here d is a scale, rather than an average residual degree. All estimates, including those for forced assignments, then apply and give the stated exponent. In particular, the argument remains valid if preprocessing converts most of the original binary bias into unit clauses. □

Without units, the hypothesis is exactly a gap between literal frequencies. For a simple input with units, the gap counting units and binary clauses equally differs from the effective gap by at most one. A linear gap in that convention therefore gives a linear effective gap with any fixed slightly smaller coefficient for sufficiently large n. No restriction on the degree sequence is needed for this dense corollary.

Variable selection and multiplicities. The algorithm uses a fixed branching order to make its specification unambiguous. The probability estimates also apply to a deterministic selection rule based on the current state: the selected variable is known before its fair value is revealed, and each variable is still assigned once. This observation is a separate extension of the stated algorithm. Similarly, the analysis after preprocessing uses the occurrence counts and their squared-sum bound, rather than simplicity itself. It therefore extends to explicitly repeated clauses whenever those quantitative bounds hold. Simplicity supplies the automatic bound used in the dense corollary.

Order of the two children. The preferred value specifies which child receives the pruning constraint. The correctness proof uses a preferred-first order only to select an optimum that survives. Actual exploration can visit the two children in either order, provided the state-based choices and pruning rules remain the same. No conclusion here relies on a random choice by the algorithm.

3.8 Interpretation and possible applications

Literal bias has a direct interpretation in a model of binary decisions: one alternative participates in more potential conflicts than the other. This gives a natural way to look for applications of the theorem. It does not establish that a particular application produces the required density or gap sizes. We distinguish exact encodings and class membership from candidate sources of practical instances.

Two alternatives and asymmetric conflicts. Suppose object i has two alternatives, represented by $x_i = 0$ and $x_i = 1$. These could be two placements, routes, or configurations. A conflict between the choices $x_i = a$ and $x_j = b$ is represented by the soft clause

$$(x_i \neq a) \vee (x_j \neq b),$$

where $x_i \neq 0$ denotes the literal x_i and $x_i \neq 1$ denotes $\neg x_i$. The clause is violated exactly when the conflicting choices are both made. Thus minimizing the number of pairwise conflicts is precisely unweighted Max 2-SAT when every conflict has unit cost. With at most one copy of each conflicting pair of alternatives, this produces a simple formula.

Let c_i^a count the conflicts involving alternative a of object i. Before pair reduction there are no units, and

$$d_2(\neg x_i) - d_2(x_i) = c_i^1 - c_i^0.$$

Consequently, if at least ηn objects satisfy $|c_i^1 - c_i^0| \geq \varepsilon n$, Corollary 3.8 applies. One possible source is a placement problem where one candidate position lies in a congested region and the other lies in a less congested region. This is a modeling example, not an assertion about measured placement instances. The more conflict-prone alternative need not be globally inferior: its actual cost depends on the choices made for the other objects. The pruning constraint describes the completions on which that alternative can still be useful. For smaller relative gaps, the degree-spread condition in (4) must also be checked.

Same-side and different-side preferences. A related model asks for a partition into two labeled sides, with soft preferences that selected pairs lie on the same or on different sides. This is the two-cluster version of correlation clustering. Signed-network models and correlation clustering also arise in the study of social balance and link classification [6]. The exact unit-cost encodings are

preference	clauses	number satisfied
$x_u \neq x_v$	$(x_u \vee x_v),\ (\neg x_u \vee \neg x_v)$	$1 + [x_u \neq x_v]$
$x_u = x_v$	$(x_u \vee \neg x_v),\ (\neg x_u \vee x_v)$	$1 + [x_u = x_v]$.

Both encodings contribute one occurrence of each polarity at each endpoint. Thus adding same-side preferences to Max-Cut does not create literal bias: every variable remains balanced, regardless of the mixture of preferences. Complementing every variable preserves the score. These instances may have other useful structure, but the gap theorem does not follow from the presence of agreement preferences alone.

Fixed labels or preferences for a particular side can break this symmetry. For example, substitute the known labels of a set of anchor vertices. Each preference between an anchor and an unassigned vertex becomes a unit preference plus a constant. If h_v is the number of resulting units favoring $x_v = 1$ minus the number favoring $x_v = 0$, then the effective gap of the residual formula is $g_{\mathbf{v}} = 2h_v$: interactions between unassigned vertices still contribute balanced complementary pairs. A few unweighted anchors cannot generate linear gaps; a linear net preference requires linearly many such contributions.

There is a precise sufficient condition for this anchored model. Let N vertices remain unassigned, let there be at most aN fixed anchors for a constant a, and allow at most one unit-cost preference per vertex pair. If at least ηN unassigned vertices satisfy $|h_v| \geq \beta N$ for constants $\eta, \beta > 0$, the same bias argument gives $O^*(c^N)$ time for some $c < 2$. Indeed, $D_{\mathbf{v}} \leq 2(N-1)$, $u_{\mathbf{v}} \leq aN$, and hence

$$\sum_v (D_{\mathbf{v}} + u_{\mathbf{v}})^2 \leq (2+a)^2 N^3, \qquad \sum_{v \in S} |g_{\mathbf{v}}| \geq 2\beta\eta N^2.$$

Apply the proof of Corollary 3.8 with scale $d = N$, $R = (2+a)^2$, and $\theta = \min\{1, 2\beta\}\eta$. The original labels are supplied data, rather than values chosen by the algorithm. Repeated units created by substituting them are retained. This is a concrete family covered by the proof; whether a given labeled network has this much net anchor information is an empirical question.

Selecting compatible items. Another exact encoding starts from a simple conflict graph $G = (V, E)$. Introduce a unit clause (x_v) for each item and a binary clause $(\neg x_u \vee \neg x_v)$ for each conflict $uv \in E$. If S is the set of selected items, the score is

$$F_G(S) = |E| + |S| - |E(G[S])|.$$

If S contains an edge, removing either endpoint loses one unit reward but eliminates at least one violated conflict clause. The score therefore does not decrease. Repeating this operation produces an independent set with at least the original score. Conversely, an independent set S has score $|E| + |S|$. It follows that

$$\mathrm{OPT}(F_G) = |E| + \alpha(G),$$

where $\alpha(G)$ is the maximum independent-set size. An optimal assignment can be converted to an optimal compatible selection by the same procedure. All clauses have unit weight; there are no additional hard constraints.

Here $d_2(\neg x_v) = \deg_G(v)$, $d_2(x_v) = 0$, and there is one positive unit, so $g_{\mathbf{v}} = 2 - \deg_G(v)$. If a linear number of vertices have degree at least κn for a constant $\kappa > 0$, their absolute effective gaps are at least $\kappa n/2$ for sufficiently large n. Thus the encoding belongs to Corollary 3.8. This example illustrates the meaning of the class; it is not offered as a new improvement for maximum independent set.

Documented applications and the remaining empirical question. Miyashiro and Matsui [7] formulate sports scheduling with minimum home/away breaks as Max 2-SAT. One eliminates a variable from each game to enforce that exactly one team plays at home; the remaining objective is a collection of soft binary clauses. This establishes an application of Max 2-SAT, but the formulation does not by itself establish our dense-gap assumptions.

Erlebach, Hall and Schank [8] use Max 2-SAT when inferring Internet customer–provider relationships from routing paths. Variables describe edge orientations, and two-literal clauses forbid incompatible orientations along consecutive edges of an observed path. For length-two paths there is one clause per path, so maximizing satisfied clauses exactly maximizes the number of valid paths. For longer paths this correspondence is not generally one-to-one; their work analyzes the resulting approximation. The asymmetric interpretation of the variables makes these formulas a candidate for examining literal bias, not an established member of our class.

A useful empirical study would measure the gaps and degree distribution before and after canonical reduction in such encodings. For the dense corollary, the relevant evidence is a linear number of linear effective gaps in a simple input. For the relative-bias theorem, one must also

measure the sum of squared degrees and unit multiplicities in (4). Repeated observations and unequal conflict costs must be accounted for explicitly; they cannot be discarded merely to make an instance simple. Nor can general application-specific hard constraints be treated as the auxiliary pruning constraints of our algorithm, since replacement assignments must respect genuine hard constraints. These checks would establish whether the proved classes explain practical instances, beyond the natural interpretations given here.

4 Random formulas

4.1 Intuition

The literal-gap result in Section 3 uses a substantial initial preference for one value of many variables. Random formulas need not have such a preference. For intuition, consider a fixed graph in which every vertex has degree d, and put one clause on each edge by choosing its two literal signs independently and uniformly. Assume for this explanation that there are no initial unit clauses. The difference between the two literal frequencies of a variable is typically of order $\sqrt{d}$, rather than a constant fraction of d. Thus, for growing d, we cannot directly apply the relative literal-bias theorem. Nevertheless, these smaller gaps still provide useful pruning constraints.

To see why, follow a path in the branching tree and consider the constraints created along it. As in the previous section, the preferred value of the next variable is the one that satisfies more clauses on average when the remaining variables are assigned uniformly. When the path takes the other value, we add a constraint requiring that this choice be strictly better than the preferred value for the eventual completion. This constraint requires the remaining assignment to reverse the current preference. In the random model, a typical preference has size of order $\sqrt{d}$. The resulting individual constraint is therefore weaker than one arising from a gap proportional to d.

It is useful to relate this to ordinary SAT clauses. A clause on k distinct variables forbids exactly one of their 2^k assignments, so a shorter clause excludes a larger fraction of the possibilities. Our pruning constraints can forbid many assignments, and their strength depends on both their length and their bound. For illustration, consider

$$z_1 + \cdots + z_k \leq \frac{k}{2} - r,$$

where the z_i are distinct Boolean variables and $r > 0$. Under a uniform assignment, the sum has mean $k/2$ and standard deviation $\sqrt{k}/2$. Satisfying the constraint therefore requires a downward deviation of at least r from the mean. We call r its *deficit*.

For a fixed deficit, reducing the number of variables makes the required deviation larger relative to the standard deviation. This explains why shorter constraints can prune more. If instead the deficit grows proportionally to $\sqrt{k}$, the required number of standard deviations stays constant, and constraints of different lengths have comparable pruning power. For example,

$$\sum_{i=1}^{100} z_i \leq 40, \qquad \sum_{i=1}^{400} z_i \leq 180$$

both require a deviation of two standard deviations below the mean. The second constraint is four times longer, but the fractions of assignments satisfying them are comparable. If the deficit is a fixed positive fraction of k, however, only an exponentially small fraction of assignments satisfy the constraint. In that case, a longer constraint can be substantially stronger. Thus length alone does not determine pruning power: the deficit must be compared with the fluctuations that the remaining variables can provide.

Now consider a path obtained by choosing each branching value with probability 1/2. About half of its branching steps take the nonpreferred value, so a linear number of branching steps creates a linear number of constraints. Their requirements accumulate. This alone would not suffice: in an adversarial formula, many constraints might ask for essentially the same assignment to their shared variables. One choice of a shared variable could then help satisfy all of them. Random literal signs prevent this alignment from persisting throughout the accumulated constraints.

The role of opposing signs can be seen in a small example. Suppose two constraints contain an unassigned variable y as follows:

$$y + A \le b, \qquad \neg y + B \le c,$$

where A and B are sums of other literal occurrences. Adding the constraints and using $y + \neg y = 1$ gives

$$A + B \le b + c - 1.$$

Whichever value we give y, it helps one constraint and hurts the other. Its contribution to their sum is fixed. If both constraints instead contained y, setting $y = 0$ would help both. A single cancellation need not yield substantial pruning, but many such cancellations limit how much the remaining assignments can help satisfy the combined requirement.

This is the useful effect of overlap in random formulas. Shared variables occur with both polarities, so positive and negative occurrences largely cancel when we add the constraints. The relevant diversity is therefore in the literal signs, not necessarily in the sets of neighboring variables: even constraints with the same variable set can impose different demands. Cancellation preserves the accumulated shortfall from the average while reducing the variation that the remaining variables can provide. The summed constraint can consequently admit very few completions even before it becomes infeasible. The algorithm exploits this by enumerating the completions of that sum, as described below.

For the regular-graph example, the requirements accumulated over a linear number of branching steps have scale $n\sqrt{d}$, whereas the ordinary fluctuations in the summed contributions have scale $\sqrt{nd}$. The proof makes this comparison precise and bounds the exceptional assignments that provide enough compensation by an exponentially small fraction. It does not assume that the learned constraints are independent, nor that every path has the stated behavior. Instead, it bounds the total enumeration work over all paths and proves an exponential saving for most random inputs. The algorithm itself is deterministic.

For an irregular graph, the degree-dependent weights introduced below put these requirements on a comparable scale before adding them. A larger constraint does not necessarily prune a larger fraction of assignments: in the random setting, its larger typical imbalance is accompanied by larger fluctuations. Dividing by approximately the square root of the degree accounts for this change of scale. Multiplying a constraint by a positive weight does not change its satisfying assignments; the weight determines its contribution to the summed constraint. An unusually large imbalance still gives a stronger requirement after normalization. These weights allow the argument to handle arbitrary degree sequences and several independently sampled clauses on a variable pair.

4.2 Notation and enumeration

We use the same residual formulas, literal counts, and pruning constraints as in Section 3. The new operation is to stop branching after a specified number of variables and enumerate completions satisfying the summed constraint C_Σ defined in (7) below.

Let n denote the number of input variables. The algorithm separately optimizes variables occurring in no binary clause and searches on the remaining $r \le n$ variables. List these as

$\mathbf{u}_1, \ldots, \mathbf{u}_r$ in their original input order, inherited from the fixed input indices. This list specifies the branching order; the pruning constraints do not determine or change it. Write x_t for the Boolean value of $\mathbf{u}_t$. A *prefix assignment of length T* specifies $x_1, \ldots, x_T$; a *suffix assignment* specifies the remaining $r - T$ values. These terms refer to this fixed order. As before, α denotes the current partial assignment and $\mathcal{C}$ the collection of pruning constraints. We retain the individual constraints, substituting each assigned value in them as the search proceeds.

To define C_Σ, give each constraint $C : \sum_i \ell_i \le b_C$ a positive rational weight ω_C, fixed when it is created. Multiply both sides of each constraint by its weight and add the resulting inequalities:

$$C_\Sigma : \qquad \sum_{C \in \mathcal{C}} \omega_C \sum_i \ell_i \le \sum_{C \in \mathcal{C}} \omega_C b_C. \tag{7}$$

These weights are assigned to whole constraints; they are distinct from the coefficients obtained by collecting repeated literal occurrences. When every $\omega_C = 1$, this is the ordinary sum used in Lemma 2.5, and each literal's coefficient is simply its total number of occurrences. The random-formula algorithm specifies the weights in (10). Every assignment satisfying all individual constraints satisfies this sum. The converse need not hold. In an empty collection the sum is $0 \le 0$.

The sum can be simplified just as in Lemma 2.5. If a variable contributes $aw + b\neg w$ with $a, b \ge 0$, replace this by $(a-b)w + b$ when $a \ge b$, or by $(b-a)\neg w + a$ otherwise, and move the constant to the right-hand side. After collecting all occurrences, we obtain

$$\sum_{i=1}^{k} w_i z_i \le B, \qquad w_i \ge 0, \quad z_i \in \{0, 1\}, \tag{8}$$

where each z_i is one remaining variable or its negation, and each remaining variable appears once. Variables with zero coefficient are included as well, so an enumerated assignment is complete.

To enumerate this inequality, assign the z_i one at a time and reject a partial assignment as soon as its assigned contribution exceeds B. Every retained partial assignment has a completion: set all remaining z_i to zero. Lemma 4.4, proved after the main running-time argument, uses this observation to bound enumeration time by the number of satisfying completions, up to polynomial factors.

4.3 The algorithm

The search has two stages. First, it branches on a fixed prefix of the variable order, using the preferred literal and pruning constraint from the literal-gap algorithm. For each prefix assignment that survives, it then enumerates assignments to the remaining variables satisfying C_Σ. These completions may violate an individual constraint; we evaluate their actual scores, so considering them does not affect correctness. Enumeration under C_Σ, in the simplified form (8), has the running-time guarantee just described.

The algorithm receives only F. It constructs the graph $G_F = (V_F, E_F)$ directly from the binary clauses of F, before applying Preprocess:

$$E_F = \{\{u, v\} : F \text{ contains a binary clause on } u, v\}. \tag{9}$$

The vertices are exactly the endpoints, $V_F = \bigcup_{e \in E_F} e$, so G_F has no isolated vertices. Literal signs do not affect whether a pair is an edge, and unit clauses create no edges. This graph remains fixed throughout the search. The graph used to generate F belongs to the distributional assumptions below and is not supplied to the algorithm.

A variable outside V_F occurs only in unit clauses. Set it to the value satisfying more of its units, breaking ties by choosing zero. Let α_0 collect these assignments, let s_0 be their total

score, and let F_{core} contain the clauses on the remaining variables. An optimum for F_{core}, extended by α_0, is an optimum for F, with s_0 added to its score. If V_F is empty, this already solves the instance. Otherwise write $r = |V_F|$ and search on these r variables.

Before preprocessing, record the weights

$$\omega_t = \frac{1}{\lceil \sqrt{\deg_{G_F}(\mathbf{u}_t)} \rceil}. \tag{10}$$

A constraint created when branching on variable $\mathbf{u}_t$ receives weight ω_t. As explained in Section 4.1, this normalizes the degree-dependent scale of the random requirements and their fluctuations. The actual imbalance remains in the bound of each constraint. If G_F is regular, all weights are equal, and the weighted sum is equivalent to the ordinary sum. The weights remain fixed throughout the search.

Set $T = \lceil r/2 \rceil$. This is the number of variables assigned by branching before switching to enumeration; it is not the number of constraints created. At a node where $t < T$ variables have been assigned, the next branching variable is $\mathbf{u}_{t+1}$. Once $t = T$, the search creates no further constraints on that path and enumerates assignments to $\mathbf{u}_{T+1}, \ldots, \mathbf{u}_r$ satisfying the current C_Σ. For example, with eight remaining variables we branch on $\mathbf{u}_1, \ldots, \mathbf{u}_4$, then complete each surviving prefix by enumerating the last four variables. The choice of half the variables gives at most 2^T prefixes while allowing a linear number of branching steps to create constraints; the proof below bounds the total work of completing those prefixes.

The two procedures use the following data. Shared variables are initialized once by RandomFormula and are available to every call of RandomSearch.

Fixed shared data. $U = (\mathbf{u}_1, \ldots, \mathbf{u}_r)$ is the variable order, and $(\omega_1, \ldots, \omega_r)$ contains the constraint weights. Neither changes during the search. All variables remain in U, including variables that disappear from the residual formula.

Shared best solution. M is the largest score in F_{core} found so far, and $\alpha^\star$ assigns all variables of V_F and attains that score. The outer procedure adds s_0 and α_0 when returning the solution for the original formula. Initially $M = -\infty$ and $\alpha^\star = \emptyset$, indicating that no complete assignment has yet been evaluated. Every recursive call may update this same pair. An improvement is retained when the call returns and when the search backtracks to explore another branch.

Parameters of a recursive call. F is the residual soft formula, $\mathcal{C}$ is the collection of pruning constraints with their weights, and α is the current partial assignment. The value s is the score offset accumulated through preprocessing and assignments: for any completion β, its score in the formula F_{core} is s plus the number of clauses it satisfies in F. These four parameters belong to the current node and are copied separately for its children. The parameter T is the stopping depth; it is passed unchanged to every recursive call.

At each call, α assigns exactly the first $|\alpha|$ variables of U. In particular, M and $\alpha^\star$ are shared across branches, whereas s and α describe only the current branch. The identities used by Preprocess apply to copies of clauses as well, so the procedure also preserves scores for the bounded multisets allowed by the random model.

Algorithm 5 RandomFormula(F)

Construct $G_F = (V_F, E_F)$ from F by (9), before preprocessing.
Optimize variables outside V_F, obtaining $F_{\text{core}}, s_0, \alpha_0$ as above.
if $V_F = \emptyset$ **then**
  return (s_0, α_0).
end if
Initialize shared $U = (\mathbf{u}_1, \ldots, \mathbf{u}_r)$ in the original input order, where $r = |V_F|$.
$T \leftarrow \lceil r/2 \rceil$. *// Branch on the first T variables, then enumerate completions.*
Initialize the shared weights ω_t from G_F by (10).
Initialize shared $M \leftarrow -\infty$ and $\alpha^\star \leftarrow \emptyset$.
$(F, s) \leftarrow$ Preprocess(F_{core}).
RandomSearch($F, \emptyset, \emptyset, s, T$).
return $(M + s_0, \alpha^\star \cup \alpha_0)$.

Algorithm 6 RandomSearch($F, \mathcal{C}, \alpha, s, T$)

Require: Fixed shared data $U, (\omega_t)_{t=1}^r$ and shared best solution $M, \alpha^\star$, initialized above.
$t \leftarrow |\alpha|$. *// The variables $\mathbf{u}_1, \ldots, \mathbf{u}_t$ have been assigned.*
Form C_Σ by (7) and simplify it to (8).
if its bound B is negative **then**
  return
end if
if $t = T$ **then**
  for each suffix assignment β satisfying C_Σ, by Lemma 4.4 **do**
    Let q be s plus the number of clauses of F satisfied by β.
    If $q > M$, update shared $M \leftarrow q$ and $\alpha^\star \leftarrow \alpha \cup \beta$.
  end for
  return
end if
Let $\mathbf{u} \leftarrow \mathbf{u}_{t+1}$. *// The next variable in the fixed list U.*
Choose its preferred literal v by (3).
Construct C from F and v by Lemma 2.4; give it weight ω_{t+1}.
for $p \in [1, 0]$ **do**
  Make separate copies $(F_p, \mathcal{C}_p, \alpha_p, s_p)$ of the current state.
  if $p = 0$ **then**
    $\mathcal{C}_p \leftarrow \mathcal{C}_p \cup \{C\}$.
  end if
  Assign $v = p$, updating the copied formula, constraints, assignment, and score.
  RandomSearch($F_p, \mathcal{C}_p, \alpha_p, s_p, T$).
end for

This algorithm uses substitution, cancellation, and rejection during the prefix search, but does not call the full Simplify procedure. In particular, it makes no forced assignments before reaching depth T. This keeps the first T variables fixed, as required by the analysis below. Created constraints remain in $\mathcal{C}$ with their updated bounds, even if they have become automatically satisfied.

4.4 Correctness

The comparisons are the same as in the literal-gap search. The only new point is that suffix enumeration enforces C_Σ from (7), which is weaker than enforcing every constraint separately. The next lemma shows that this still retains an optimum; its proof does not use randomness.

Lemma 4.1 (Correctness of weighted enumeration)**.** *For every formula F, including formulas with repeated clauses,* RandomFormula *returns an optimum score and an assignment attaining it.*

Proof. Variables outside V_F can be optimized independently from their units. It therefore suffices to prove that the recursive search finds an optimum for F_{core}. Preprocessing changes the score of every assignment by the same recorded constant. It therefore preserves both optimal assignments and all comparisons between the two values of a variable.

Consider the full decision tree in the fixed variable order, restoring any prefixes rejected by the algorithm. At each of the first T levels, place the preferred child first. Complete the rest of the tree in any fixed order, and choose the earliest optimal complete assignment x^*. If its path takes a nonpreferred child at variable t, changing only that value to the preferred one produces an earlier assignment. Its score must be strictly smaller: a larger score contradicts optimality, and an equal score contradicts the choice of x^*. Lemma 2.4 therefore implies that x^* satisfies the constraint added at that step.

Thus x^* satisfies every constraint created on its prefix, and hence their positively weighted sum at every earlier depth. The sum test never rejects that prefix, and enumeration at depth T visits its suffix. Every enumerated candidate is scored in F_{core}, including the preprocessing and assignment offsets. Adding the separately optimized score s_0 and extending by α_0 therefore returns an optimum for F. □

4.5 Random inputs and the running-time guarantee

Two familiar models motivate the assumptions below. Fix any set E of pairs of distinct input variables, of any cardinality. There is no density, regularity, or connectivity requirement. The generating graph is $G = (V, E)$, where $V = \bigcup_{e \in E} e$: an edge means that its pair is selected for clause sampling. The graph is fixed before sampling and is not an input to the optimization algorithm. On each selected pair we may choose exactly one of the four possible clauses, uniformly at random and independently across pairs.

Alternatively, include each of the four possible clauses independently with probability $1/4$, with all choices independent across pairs as well. The number X_e of clauses on a selected pair $e \in E$ then satisfies

$$X_e \sim \text{Binomial}(4, 1/4), \qquad \mathbb{E}X_e = 1.$$

Thus a selected pair can contribute zero, one, two, three, or four clauses. There is one clause per selected pair in expectation, and $|E|$ binary clauses in total in expectation. Pairs outside E contribute no binary clauses. More generally, any fixed inclusion probability $q \in (0, 1)$ is allowed: then $X_e \sim \text{Binomial}(4, q)$ and $\mathbb{E}X_e = 4q$. The observed graph G_F contains only pairs that actually receive a clause. Thus $E_F \subseteq E$, with equality in the one-clause-per-pair model but not necessarily under independent inclusion.

Both models treat the two polarities of each endpoint symmetrically. The next definition allows both models, while excluding pair distributions whose endpoint literal counts are always exactly balanced.

Definition 4.2 (Complementation-invariant random signs)**.** Fix constants $m_0 \geq 1$ and $0 < \kappa \leq m_0^2$, a set of n input variables, and any set E of unordered pairs of distinct input variables. Let $V = \bigcup_{e \in E} e$ and $G = (V, E)$. For each edge $e = \{u, v\}$ draw a multiset M_e of

binary clauses on u, v, independently over the edges. The binary part of F is the multiset union of these M_e; pairs outside E contribute no binary clauses. A sampled M_e may be empty, so the observed graph G_F in (9) may have fewer edges than G. Assume that

1. $|M_e| \le m_0$;
2. complementing the literal of either endpoint in every clause of M_e leaves its distribution unchanged;
3. for each endpoint u,
$$\mathbb{E}\left[\left(\sum_{c \in M_e} s_u(c)\right)^2\right] \ge \kappa,$$
where $s_u(c) = 1$ if the literal of u in c is positive and $s_u(c) = -1$ otherwise.

Add any multiset of unit clauses on the input variables, fixed independently of these draws.

The last condition requires some variation in the difference between positive and negative occurrences at an endpoint. One uniformly signed clause per edge has $m_0 = \kappa = 1$. Including each of the four clauses independently with probability $q \in (0, 1)$ has $m_0 = 4$ and $\kappa = 4q(1-q)$. In particular, $q = 1/4$ gives the independent-clause model discussed in the introduction.

The theorem bounds the complexity of solving these random inputs. Its constants depend on the pair distribution but not on the density, degree sequence, or choice of the fixed support graph.

Theorem 4.3 (Random signs on every graph)**.** *Fix $m_0 \ge 1$ and $0 < \kappa \le m_0^2$. For a formula F on n variables drawn according to Definition 4.2 with any fixed generating graph G, an optimal assignment can be computed from F alone by a deterministic polynomial-space algorithm with expected running time*

$$O^*(2^{n/2} + 2^n e^{-c_0 n}), \qquad c_0 = \frac{c_1^2}{10816 m_0^2}, \qquad c_1 = \frac{\sqrt{\kappa}}{2\sqrt{1 + 4m_0^4/\kappa^2}}.$$

With probability at least $1 - 3e^{-c_0 n/2}$, its running time is $O^*(2^n e^{-c_0 n/2})$. The algorithm returns an optimum on every input; the expectations and probabilities refer only to the random input formula.

For one uniformly signed clause per edge, $c_0 = 1/216320$. Appendix A gives a sharper constant on almost-regular graphs without units. The theorem above also covers irregular supports and independently fixed unit clauses. The bound holds for any variable order fixed independently of the sampled signs.

4.6 Proof of the running-time bound

We first condition on the observed graph G_F. Lemma 4.5 shows that the surviving pairs still satisfy the random-model assumptions, with the same constants. Thus the graph used by the algorithm can be treated as fixed for the conditional analysis, even though it was constructed from F. Write $r = |V_F|$ for the number of variables searched. We count assignments to these variables that satisfy the summed constraint at their own prefix. The enumeration lemma then turns that count into a time bound. The proof has three ingredients.

1. Along a uniform path, each nonpreferred choice adds a positive requirement measured relative to the average contribution of the remaining literals. Their weighted total is at least of order r on most paths. Lemma 4.9 bounds the probability that it is much smaller.

2. The actual remaining values can help satisfy these requirements, but their effects have random signs. Lemma 4.10 bounds the chance of an unusually favorable total contribution. Its proof uses the randomness of unrevealed endpoint signs, not independence of the learned constraints.

3. Combining these estimates shows that an exponentially small fraction of conditional formula–assignment pairs pass the sum. Averaging over assignments and then over the observed graph bounds the expected enumeration work; a final probability bound gives the running time for most formulas.

The supporting lemmas and their elementary probability estimates follow the main proof.

Proof of Theorem 4.3. Use RandomFormula on F alone. Correctness is Lemma 4.1. Depth-first search and suffix enumeration use polynomial space. Constructing G_F and optimizing variables outside V_F take polynomial time.

Condition on $G_F = H$, where H is any possible nonempty observed graph, and write $r = |V(H)|$. By Lemma 4.5, its pair laws satisfy the assumptions of Definition 4.2 with the same m_0, κ. Its vertex order, inherited from the input indices, and its degrees are now fixed. All probabilities in Steps 1 and 2 refer to this conditional model. The weights are rational with polynomial encoding length, and satisfy

$$\frac{1}{2\sqrt{\deg_H(\mathbf{u}_t)}} \le \omega_t \le \frac{1}{\sqrt{\deg_H(\mathbf{u}_t)}}. \tag{11}$$

Thus all arithmetic per search state has polynomial cost.

1. Express the summed requirement. Choose a uniform assignment $x \in \{0,1\}^r$, independently of the remaining random choices in F. For the analysis only, construct its first $T = \lceil r/2 \rceil$ comparisons even if an earlier sum test would have rejected its prefix. Let $\mathrm{gain}_t(x)$ be the score with $x_t = 1$ minus the score with $x_t = 0$, keeping all other values fixed. It does not depend on x_t. Let Δ_t be its average over uniform later values, with $x_1, \dots, x_{t-1}$ fixed. The preferred value of the original variable is $p_t = 1$ when $\Delta_t \ge 0$ and $p_t = 0$ otherwise. Equivalently, the algorithm's preferred literal is x_t or $\neg x_t$, respectively, by (3).

Put $I = \{t \le T : x_t \ne p_t\}$, the indices at which a constraint is created. Lemma 2.4 expresses each such constraint as $(2p_t - 1)\mathrm{gain}_t(x) + 1 \le 0$. Hence its weighted sum is

$$A_T + Q_T \le 0, \tag{12}$$

where

$$A_T = \sum_{t \in I} \omega_t(|\Delta_t| + 1),$$
$$Q_T = \sum_{t \in I} \omega_t(2p_t - 1)\big(\mathrm{gain}_t(x) - \Delta_t\big).$$

Here A_T records the requirements at creation, while Q_T records the deviation supplied by the subsequent values. The identity follows from $(2p_t - 1)\Delta_t = |\Delta_t|$. Preprocessing preserves each score difference, so these expressions computed from the original clauses are exactly the ones used by the search.

2. Bound the fraction that satisfies the sum. Lemma 4.9 gives, for any $a, b > 0$,

$$\mathbb{P}\left(A_T < \frac{c_1 T}{26} - a - b\right) \le e^{-2a^2/(Tm_0^2)} + e^{-169b^2/(m_0^2 r)}.$$

Lemma 4.10 gives, for any $c > 0$,

$$\mathbb{P}(Q_T < -c) \le e^{-2c^2/(m_0^2 T)}.$$

Take $a = b = c = c_1T/104$. Outside these three exceptional events,

$$A_T + Q_T \geq \frac{c_1T}{26} - 3\frac{c_1T}{104} = \frac{c_1T}{104} > 0.$$

Thus the summed constraint fails. The first and third exponents equal $c_1^2T/(5408m_0^2) \geq c_0r$. The second is $169c_1^2T^2/(10816m_0^2r) \geq c_0r$, since $T \geq r/2$. Consequently

$$\mathbb{P}_{F,x}(A_T + Q_T \leq 0 \mid G_F = H) \leq 3e^{-c_0r}.$$

3. Convert the count to running time. For a prefix α, let $N(\alpha)$ count suffix assignments satisfying its sum, and put $W(F) = \sum_{\alpha\in\{0,1\}^T} N(\alpha)$. This definition includes prefixes rejected early, and so bounds the number of completions actually enumerated. Since x is uniform,

$$\mathbb{E}[W(F) \mid G_F = H] = 2^r\mathbb{P}_{F,x}(A_T + Q_T \leq 0 \mid G_F = H) \leq 3 \cdot 2^r e^{-c_0r}.$$

If H is empty, the algorithm finishes in polynomial time; define $W(F) = 1$ in this case, so the same bound holds with $r = 0$. Since $\kappa \leq m_0^2$, we have $c_0 \leq 1/216320 < \ln 2$. The quantity $2^re^{-c_0r}$ is therefore increasing in r. Using $r \leq n$ and averaging over G_F gives

$$\mathbb{E}W(F) \leq 3 \cdot 2^n e^{-c_0n}.$$

Lemma 4.4 bounds the search work by $O^*(2^T+W(F))$. The deterministic inequality $2^T \leq 2^{\lceil n/2\rceil}$ now proves the expected-time assertion. Markov's inequality applied to the unconditional expectation also gives

$$\mathbb{P}_F\left(W(F) > 2^n e^{-c_0n/2}\right) \leq 3e^{-c_0n/2}.$$

The prefix cost 2^T is smaller than the claimed high-probability time bound, up to a constant factor, because $c_0 < \ln 2$. This completes the proof. □

4.7 Supporting lemmas

We first justify the enumeration cost used in Step 3. A small number of satisfying completions does not by itself bound a search over arbitrary constraints. For the single inequality (8), however, every retained partial assignment has a completion obtained by setting its remaining oriented variables to zero.

Lemma 4.4 (Enumerating the summed constraint)**.** *Suppose an inequality $\sum_{i=1}^k w_iz_i \leq B$ with nonnegative rational weights has N satisfying Boolean assignments. Depth-first enumeration that rejects a partial assignment as soon as its assigned contribution exceeds B lists all these assignments in polynomial space and* $\mathrm{poly}(n)(N+1)$ *time, including polynomial arithmetic costs. Consequently, if learning is stopped after the first T variables in a fixed order, and $N(\alpha)$ counts the suffix assignments satisfying the sum at prefix assignment α, the total enumeration time is*

$$O^*(2^T + W(F)), \qquad W(F) = \sum_{\alpha\in\{0,1\}^T} N(\alpha). \tag{13}$$

Proof. If $B < 0$, even the minimum left-hand side, zero, exceeds B, so the root is rejected. Otherwise, every retained partial assignment extends to a solution by setting all its unassigned z_i to zero. At a fixed depth, distinct retained partial assignments have disjoint sets of completions. There are therefore at most N retained nodes at that depth. Each has at most two children, so rejected children increase the count by at most a constant factor. Summing over the $k+1$ depths gives the first bound. The depth-first stack has polynomial size.

There are 2^T prefix assignments. Constructing their sums costs polynomial time per prefix; applying the first bound to each and adding the costs gives the second bound. Prefixes rejected earlier only reduce the work. □

The algorithm does not know which pairs were available for sampling. The next lemma justifies analyzing it after fixing the graph it actually observes. Crucially, whether a pair is empty does not distinguish between the literal polarities.

Lemma 4.5 (Conditioning on the observed pairs)**.** *Suppose F is generated according to Definition 4.2. Conditional on any observed graph $G_F = H$ of positive probability, the clause multisets on its edges are independent. Each has its original distribution conditioned on being nonempty, and satisfies all three conditions of that definition with the same m_0 and κ.*

Proof. The event $G_F = H$ requires $M_e \neq \emptyset$ for each edge of H and $M_e = \emptyset$ for every other edge of the generating graph. These events concern separate independent draws. Conditioning on them therefore preserves independence, and conditions each retained pair only on its own nonemptiness. The clause-count bound remains m_0.

Complementing either endpoint does not change whether a multiset is empty. Its conditional distribution is consequently still invariant under either complementation. Finally, put $Z_{u,e} = \sum_{c \in M_e} s_u(c)$ and $p_e = \mathbb{P}(M_e \neq \emptyset)$. An empty multiset contributes zero to $Z_{u,e}^2$, so for every retained edge

$$\mathbb{E}[Z_{u,e}^2 \mid M_e \neq \emptyset] = \frac{\mathbb{E} Z_{u,e}^2}{p_e} \geq \kappa.$$

The unit clauses were fixed independently of the pair draws and remain fixed under this conditioning. This proves the claim. □

For the probability estimates below, fix a nonempty observed graph $G_F = H = (V_H, E_H)$ and put $r = |V_H|$. All probabilities and expectations refer to the conditional model in Lemma 4.5. We identify the variables with their positions $1, \dots, r$ in U when writing edges and neighbors.

The next lemma identifies the randomness available when a constraint is created and when a later variable is assigned. It is the reason that overlapping constraints can be analyzed without treating them as independent. We use the original sampled clauses in this analysis, since preprocessing preserves all gains.

Let b_t be the number of original units (x_t) minus the number of units ($\neg x_t$). For an edge $\{t, v\}$, define its contribution to the gain at t by

$$G_{tv}(x_v) = \sum_{\substack{c \in M_{\{t,v\}} \\ \text{the literal on } v \text{ is false}}} s_t(c).$$

A clause contributes its sign at t exactly when its partner literal is false. Thus

$$\begin{aligned} \operatorname{gain}_t(x) &= b_t + \sum_{v \sim t} G_{tv}(x_v), \\ \Delta_t &= b_t + \sum_{\substack{v < t \\ v \sim t}} G_{tv}(x_v) + \frac{1}{2} \sum_{\substack{v > t \\ v \sim t}} \sum_{c \in M_{\{t,v\}}} s_t(c). \end{aligned}$$

Write $d_t = \deg_H(\mathbf{u}_t)$.

Lemma 4.6 (Revealing the endpoint signs)**.** *The random formula and the independent uniform assignment can be exposed in variable order so that, just before the signs at variable t are revealed,*

$$\Delta_t = b_t + \sum_{v \sim t} \rho_{t,\{t,v\}} a_{tv}, \qquad Y_t = \sum_{v \sim t} a_{tv}^2 \leq m_0^2 d_t.$$

The coefficients a_{tv} are already known, and the signs $\rho_{t,\{t,v\}}$ are independent fair signs. After they are revealed, x_t is still a fair bit independent of Δ_t. For each edge $\{t,v\}$ with $t<v$, the later deviation has the form

$$G_{tv}(x_v)-\frac{1}{2}\sum_{c\in M_{\{t,v\}}} s_t(c)=\rho_{t,\{t,v\}}\eta_{tv}\theta_{v,\{t,v\}},\qquad |\eta_{tv}|\le m_0/2.$$

Here η_{tv} is known from the start. The signs $\theta_{v,\{t,v\}}$ are fresh independent fair signs when endpoint v is reached, conditional on all earlier exposed information.

Proof. For each edge e, first draw an independent base multiset K_e with the prescribed law of M_e. Write $\sigma_u(c)$ for an endpoint sign in this base multiset. Independently draw a fair sign $\rho_{u,e}$ for each endpoint of each edge. Obtain M_e by replacing its endpoint signs by $\rho_{u,e}\sigma_u(c)$. Complementation invariance implies that the result has the required distribution, independently on different edges.

Draw the uniform assignment independently and put $\theta_{u,e}=\rho_{u,e}(2x_u-1)$. Conditional on any assignment x, the transformation from the ρ signs to the θ signs preserves their independent fair distribution. Consequently all the θ signs are mutually independent and independent of x and of the base multisets. A partner literal on v is false precisely when $\sigma_v(c)\theta_{v,e}=-1$.

Reveal all K_e first. At step t, reveal the signs $\rho_{t,e}$ on edges incident to t, followed by x_t. Let $\mathcal{H}_t$ be the information after this step. For $e=\{t,v\}$ set

$$a_{tv}=\begin{cases}\displaystyle\sum_{\substack{c\in K_e\\ \sigma_v(c)\theta_{v,e}=-1}}\sigma_t(c), & v<t,\\[2ex] \displaystyle\frac{1}{2}\sum_{c\in K_e}\sigma_t(c), & v>t.\end{cases}$$

These coefficients are determined by $\mathcal{H}_{t-1}$. Multiplying them by $\rho_{t,e}$ gives precisely the earlier-pair contributions or the average later-pair contributions to Δ_t. This proves its representation. The conditional variance of the signed sum is Y_t, because the new signs are independent and have mean zero. Also $|a_{tv}|\le m_0$, proving the upper bound on Y_t.

For $t<v$ define

$$\eta_{tv}=\sum_{\substack{c\in K_e\\ \sigma_v(c)=-1}}\sigma_t(c)-\frac{1}{2}\sum_{c\in K_e}\sigma_t(c).$$

This is half the difference of the signed sums over the two partner polarities, so $|\eta_{tv}|\le |K_e|/2\le m_0/2$. If $\theta_{v,e}=1$, the first sum selects the false partner literals; if $\theta_{v,e}=-1$, their complementary set is selected. Subtracting the average gives the asserted deviation with opposite signs in these two cases. The θ signs at v are jointly independent and fair conditional on $\mathcal{H}_{v-1}$, as required. □

We will apply the following probability estimate both to sequential choices and to functions of the independent edge data. Its interval form gives the constants used in Step 2 of the main proof.

Lemma 4.7 (Conditional ranges and bounded differences)**.** *Let $X_1, \dots, X_k$ be exposed in order, and let $\mathcal{H}_j$ denote the information after step j. Conditional on the earlier information, suppose X_j lies in an interval of length ℓ_j, where the interval is known before X_j is revealed. If $\sum_j \ell_j^2 \le B$ on every path for a deterministic $B > 0$, then, for $a > 0$,*

$$\mathbb{P}\left(\sum_{j=1}^{k}(X_j - \mathbb{E}[X_j \mid \mathcal{H}_{j-1}]) < -a\right) \le e^{-2a^2/B}.$$

In particular, if f is a function of independent inputs and changing input j changes f by at most c_j, with $\sum_j c_j^2 > 0$, then

$$\mathbb{P}(f < \mathbb{E}f - a) \le e^{-2a^2/\sum_j c_j^2}.$$

If all c_j are zero, f is constant and this probability is zero.

Proof. A centered random variable Z supported on an interval of length ℓ satisfies $\mathbb{E}e^{\lambda Z} \le e^{\lambda^2\ell^2/8}$. To see this, let $\psi(\lambda) = \log \mathbb{E}e^{\lambda Z}$. Then $\psi(0) = \psi'(0) = 0$, and $\psi''(\lambda)$ is the variance under the distribution that reweights each outcome by $e^{\lambda Z}$. This distribution has the same support. Every random variable in an interval of length ℓ has variance at most $\ell^2/4$: subtract the midpoint, bound the second moment by $\ell^2/4$, and use that variance is no larger than this second moment. Integrating the bound on ψ'' twice proves the exponential estimate.

Apply it conditionally at each step, with either sign of λ. Successive conditional expectations and the bound on the sum of squared interval lengths give

$$\mathbb{P}\left(\sum_{j}(X_j - \mathbb{E}[X_j \mid \mathcal{H}_{j-1}]) < -a\right) \le \exp(-\lambda a + \lambda^2 B/8).$$

Taking $\lambda = 4a/B$ proves the first assertion. For the second assertion, reveal the independent inputs one at a time and consider the successive conditional expectations of f. Their increments have conditional mean zero and conditional range at most c_j. Indeed, two possible values of input j change f by at most c_j for every fixed choice of the remaining inputs, and hence also after averaging those inputs. The first assertion applies. □

An average gain may include an arbitrary unit-clause contribution. The next lemma shows that such a shift cannot remove the useful absolute fluctuations of a symmetric signed sum. Its second assertion converts a variance estimate into the average gap needed below.

Lemma 4.8 (Absolute gains and second moments)**.** *Let $Z = \sum_i a_i\rho_i$ for independent fair signs ρ_i, and put $Y = \sum_i a_i^2$. For every real b and every $K > 0$,*

$$\mathbb{E}\min\{|b + Z|, K\} \ge \frac{1}{2}\left(\sqrt{Y/3} - \frac{Y}{4K}\right).$$

For any nonnegative random variable U with $\mathbb{E}U > 0$ and finite second moment,

$$\mathbb{E}\sqrt{U} \ge \frac{(\mathbb{E}U)^{3/2}}{(\mathbb{E}U^2)^{1/2}}.$$

Proof. Independence and zero means give $\mathbb{E}Z^2 = Y$ and

$$\mathbb{E}Z^4 = \sum_i a_i^4 + 6\sum_{i<j} a_i^2 a_j^2 \le 3Y^2.$$

The Hölder inequality, with exponents 3/2 and 3, gives

$$\mathbb{E}|Z|^2 = \mathbb{E}(|Z|^{2/3}|Z|^{4/3}) \le (\mathbb{E}|Z|)^{2/3}(\mathbb{E}|Z|^4)^{1/3}.$$

Rearranging yields $\mathbb{E}|Z| \ge \sqrt{Y/3}$ when $Y > 0$; the claim is immediate when $Y = 0$. Write $r_+ = \max\{r, 0\}$. Then $(z - K)_+ \le z^2/(4K)$ for $z \ge 0$, because for $z > K$ the inequality is equivalent to $(z - 2K)^2 \ge 0$. It follows that

$$\mathbb{E}\min\{|Z|, K\} \ge \sqrt{Y/3} - Y/(4K).$$

The triangle inequality gives $|b + z| + |b - z| \ge 2|z|$. After truncation this implies $\min\{|b + z|, K\} + \min\{|b - z|, K\} \ge \min\{|z|, K\}$. Symmetry of Z makes the two expectations on the left equal, proving the first assertion.

For the second assertion, apply the same Hölder inequality to $(\sqrt{U})^{2/3}U^{2/3}$:

$$\mathbb{E}U \le (\mathbb{E}\sqrt{U})^{2/3}(\mathbb{E}U^2)^{1/3}.$$

Rearranging proves the claim. □

We can now bound the accumulated requirements A_T in the main proof. There are two sources of variation: the requirement at an individual branching step, and the amount of conditional variance available across the selected variables. The two terms in the following bound account for these separately.

Lemma 4.9 (Accumulated requirements)**.** *With A_T, c_1, T, r as in the proof of Theorem 4.3, for all $a, b > 0$,*

$$\mathbb{P}\left(A_T < \frac{c_1 T}{26} - a - b\right) \le e^{-2a^2/(Tm_0^2)} + e^{-169b^2/(m_0^2 r)}.$$

Proof. Use the exposure of Lemma 4.6.

1. A lower bound at each step. Put

$$\widetilde{X}_t = [t \in I]\, \omega_t \min\{|\Delta_t|, m_0\sqrt{d_t}\}.$$

Here $[t \in I]$ is one when a constraint is created at step t, and zero otherwise. Thus $0 \le \widetilde{X}_t \le m_0$ and $A_T \ge \sum_{t \le T} \widetilde{X}_t$. Conditional on $\mathcal{H}_{t-1}$, the gain is the fixed shift b_t plus a symmetric signed sum of variance Y_t. The bit x_t is independent of this gain and chooses the nonpreferred value with probability one half. Lemma 4.8 gives

$$\begin{aligned}\mathbb{E}[\widetilde{X}_t \mid \mathcal{H}_{t-1}] &\ge \frac{\omega_t}{4}\left(\sqrt{Y_t/3} - \frac{Y_t}{4m_0\sqrt{d_t}}\right) \\ &\ge \frac{\omega_t\sqrt{Y_t}}{4}\left(\frac{1}{\sqrt{3}} - \frac{1}{4}\right) \ge \mu_t := \frac{\omega_t\sqrt{Y_t}}{13}.\end{aligned}$$

The middle inequality uses $Y_t \le m_0^2 d_t$. Lemma 4.7, with interval lengths m_0, therefore gives

$$\mathbb{P}\left(A_T < \sum_{t \le T} \mu_t - a\right) \le e^{-2a^2/(Tm_0^2)}.$$

2. The average available variance. For an earlier neighbor $v < t$, condition on the base multiset and let q_+ and q_- be its signed sums at t over the two polarities at v. The coefficient a_{tv} takes these two values with equal probabilities. Thus

$$\mathbb{E}[a_{tv}^2 \mid K_{\{t,v\}}] = \frac{q_+^2 + q_-^2}{2} \ge \frac{(q_+ + q_-)^2}{4}.$$

Averaging and applying the model's third condition gives $\mathbb{E}a_{tv}^2 \geq \kappa/4$. For a later neighbor, a_{tv} is half the endpoint signed sum, giving the same lower bound directly. The terms of $Y_t = \sum_{v\sim t} a_{tv}^2$ depend on independent edge data, and each lies in $[0, m_0^2]$. Consequently

$$\mathbb{E}Y_t \geq \frac{\kappa d_t}{4}, \qquad \mathrm{Var}(Y_t) \leq \frac{m_0^4 d_t}{4}.$$

By Lemma 4.8,

$$\begin{aligned}\mathbb{E}\sqrt{Y_t} &\geq \frac{\sqrt{\mathbb{E}Y_t}}{\sqrt{1+\mathrm{Var}(Y_t)/(\mathbb{E}Y_t)^2}} \\ &\geq \frac{\sqrt{\kappa d_t}/2}{\sqrt{1+4m_0^4/(\kappa^2 d_t)}} \geq c_1\sqrt{d_t}.\end{aligned}$$

We used $d_t \geq 1$ in the last step. The lower weight bound (11) now implies $\mathbb{E}\sum_{t\leq T}\mu_t \geq c_1 T/26$.

3. Concentration over the edge data. Regard each base multiset and its two θ signs as one independent edge input. Changing the input on $\{u, v\}$ affects only the terms μ_u, μ_v that belong to the prefix. In each affected square root, only one squared coefficient changes, with absolute coefficient in $[0, m_0]$. For $R \geq 0$ and $a, b \in [0, m_0]$, $|\sqrt{R+a^2} - \sqrt{R+b^2}| \leq |a-b| \leq m_0$. Thus an edge changes the sum by at most $m_0(\omega_u + \omega_v)/13$. The sum of squared change bounds is at most

$$\frac{m_0^2}{169}\sum_{\{u,v\}\in E_H}(\omega_u+\omega_v)^2 \leq \frac{2m_0^2}{169}\sum_v \omega_v^2 d_v \leq \frac{2m_0^2 r}{169}.$$

Lemma 4.7 gives

$$\mathbb{P}\left(\sum_{t\leq T}\mu_t < c_1 T/26 - b\right) \leq e^{-169b^2/(m_0^2 r)}.$$

Combining this with the bound from the first step proves the lemma. □

The last estimate controls how much the later values can help satisfy the accumulated constraints. A shared variable may affect many constraints, but the relevant endpoint signs on its different edges are still independent when that variable is reached. This is the estimate for Q_T used in Step 2 of the main proof.

Lemma 4.10 (Favorable contributions from later values)**.** *For every $c > 0$,*

$$\mathbb{P}(Q_T < -c) \leq e^{-2c^2/(m_0^2 T)}.$$

Proof. By Lemma 4.6, the contribution from $\{t, v\}$, where $t \leq T$ and $t < v$, is

$$[t \in I]\,\omega_t(2p_t - 1)\rho_{t,\{t,v\}}\eta_{tv}\theta_{v,\{t,v\}}.$$

Order these terms by their later endpoint v. Their coefficients are known before step v, while the θ signs at v are conditionally independent fair signs. We may reveal these signs one at a time and then reveal x_v; the corresponding $\rho_{v,e}$ are then determined. This is equivalent to the earlier exposure and gives a conditional mean-zero increment for each displayed term.

Its conditional range has length at most $\omega_t m_0$. The sum of squared range lengths is bounded by

$$m_0^2\sum_{t\leq T}\omega_t^2|\{v > t : \{t, v\} \in E_H\}| \leq m_0^2\sum_{t\leq T}\omega_t^2 d_t \leq m_0^2 T.$$

Lemma 4.7 proves the claim. This argument does not require independence between A_T and Q_T. □

4.8 Independent clauses on an arbitrary support

The general pair-distribution condition has a particularly simple consequence for independent clauses. It allows any prescribed set of variable pairs; neither a density condition nor regularity is needed.

Corollary 4.11 (Arbitrary supports)**.** *Fix $q \in (0,1)$ independently of n. Fix any set of variable pairs and include each of the four clauses on each pair independently with probability q. Add unit clauses fixed independently of the binary draw. An optimal assignment can be computed by a deterministic polynomial-space algorithm with expected running time $O^*(2^n e^{-c_0(q)n})$, where $c_0(q) > 0$ depends only on q. The algorithm is exact on every input.*

Proof. Apply Theorem 4.3 with $m_0 = 4$ and $\kappa = 4q(1-q)$. The theorem already gives an algorithm whose only input is F. Its prefix term is absorbed into $O^*(2^n e^{-c_0(q)n})$ because $c_0(q) < (\ln 2)/2$. The same bound holds if the selected pair set is itself random and chosen independently of the clause signs: condition on that set, apply the uniform bound of the theorem, and then average. □

The saving is uniform over supports and their degree sequences, but not over q tending to zero or one. This fixed-q result needs no logarithmic degree threshold. Unit clauses chosen after observing the binary draw are not covered by its exposure argument. The extension below permits that stronger choice under a density assumption. Complementary pairs alone have endpoint signed sum zero and fail the third condition of Definition 4.2; the Max-Cut results in Section 5 treat such inputs by a different analysis.

4.9 An extension using assignment counts

The preceding argument relies on fresh symmetric endpoint signs. We next allow a systematic preference for equal or opposite literal polarities on a pair, and unit clauses chosen after the binary formula is known. To handle the systematic part, we guess one integer: the total number of variables set to one. There are only $n+1$ possibilities. Within each guess we again form comparison constraints and enumerate completions satisfying their sum.

We introduce the additional notation here because the weighted search above does not need it. Expanding the truth indicator of each clause expresses the number of satisfied clauses as

$$F(x) = c + \sum_{v} a_v x_v + \sum_{u<v} J_{uv} x_u x_v, \qquad x \in \{0,1\}^n. \tag{14}$$

The matrix J is symmetric with zero diagonal. A binary clause with equal literal polarities contributes -1 to J_{uv}; one with opposite polarities contributes $+1$. Unit clauses change only the constant and linear terms. The gain from changing x_t from zero to one is therefore

$$\mathrm{gain}_t(x) = a_t + \sum_{w \neq t} J_{tw} x_w.$$

These quantities are computed by scanning clauses.

Fix a rational number β representing the common part of the pair coefficients. For a guess $h \in \{0, \dots, n\}$ define

$$H_t^{h,\beta}(x) = \mathrm{gain}_t(x) + \beta \left(h - \sum_{w \neq t} x_w \right) = a_t + \beta h + \sum_{w \neq t} (J_{tw} - \beta) x_w. \tag{15}$$

The second expression explains the adjustment: once h is fixed, the common pair contribution β is replaced by the known constant βh. The remaining coefficients are the deviations $J_{tw} - \beta$. When $h = \sum_w x_w$ is the correct guess, $H_t^{h,\beta}(x) = \mathrm{gain}_t(x) + \beta x_t$. We account for this last term by allowing a slack of $|\beta|$ in a comparison.

The algorithm. Choose a prefix of T variables in a fixed order and partition it into ordered, nonempty rounds $B_1, \dots, B_\ell$. For the random formulas below each round contains one variable. We allow larger rounds because Section 5 uses the same procedure with two rounds. Within round B_j, choose each preference using only the assigned values in earlier rounds, whose union is denoted by $B_{<j}$. The values of all other variables, including those in the current round, are averaged uniformly. This average is computed exactly by substituting $1/2$ in (15); no sampling is involved.

If this average is nonnegative, the preferred value p_t is one; otherwise it is zero. When the prefix takes the nonpreferred value, add the comparison

$$(2p_t - 1)H_t^{h,\beta}(x) \le |\beta|.$$

The code stores the sum of left-hand sides minus right-hand sides in $L(x)$, so the single inequality for enumeration is $L(x) \le 0$. The global best score and assignment again are $M, \alpha^\star$.

Algorithm 7 CountComparisons$(F, \beta, B_1, \dots, B_\ell)$

$M \leftarrow -\infty$; $\alpha^\star \leftarrow \emptyset$.
$(F, s_0) \leftarrow$ Preprocess(F).
Set $\mathcal{H} \leftarrow \{0, \dots, n\}$ if $\beta \neq 0$, and $\mathcal{H} \leftarrow \{0\}$ otherwise. *// One run suffices when the comparison does not depend on h.*
for each $h \in \mathcal{H}$ **do**
  for each prefix assignment α on $B_1 \cup \cdots \cup B_\ell$ **do**
    $L(x) \leftarrow 0$.
    for each round B_j, and each variable $t \in B_j$ **do**
      Compute the average Δ_t of $H_t^{h,\beta}$ with $x_{B_{<j}} = \alpha_{B_{<j}}$.
      Set $p_t \leftarrow 1$ if $\Delta_t \ge 0$, and $p_t \leftarrow 0$ otherwise.
      if $\alpha_t \neq p_t$ **then**
        $L(x) \leftarrow L(x) + (2p_t - 1)H_t^{h,\beta}(x) - |\beta|$.
      end if
    end for
    for each suffix y with $L(\alpha, y) \le 0$, by Lemma 4.4 **do**
      Let $q \leftarrow s_0 + F(\alpha, y)$.
      If $q > M$, set $M \leftarrow q$ and $\alpha^\star \leftarrow \alpha \cup y$.
    end for
  end for
end for
return $(M, \alpha^\star)$.

Each guess uses a separate set of comparisons, and guesses are processed consecutively. Enumeration need not enforce that a candidate has exactly h ones: every candidate is evaluated in the actual formula. For the biased model below, use singleton rounds through $T = \lfloor n/2 \rfloor$ and

$$\beta = \binom{n}{2}^{-1} \sum_{u<v} J_{uv}.$$

This is the empirical mean of the signed pair coefficients. For the zero-mean models we may instead set $\beta = 0$. Neither choice requires an eigenvalue or eigenvector calculation.

Correctness. The count guess changes the comparisons, so we must check that a correct guess preserves an optimum. The following lemma accounts for the added slack. This argument is independent of the distribution of F and also justifies the round version used later.

Lemma 4.12 (Correctness of count comparisons)**.** *For every F, every rational β, and every ordered prefix partition,* CountComparisons *returns an optimum score and an assignment attaining it.*

Proof. Fix an optimal complete assignment x^* and its correct count $h = \sum_w x^*_w$. If $x^*_t \neq p_t$, changing just this variable to p_t cannot increase the score, so $(2p_t - 1)\mathrm{gain}_t(x^*) \leq 0$. For the correct count, (15) therefore gives

$$(2p_t - 1)H_t^{h,\beta}(x^*) = (2p_t - 1)\mathrm{gain}_t(x^*) + (2p_t - 1)\beta x^*_t \leq |\beta|.$$

Every comparison on its prefix holds, and hence so does their sum. The corresponding suffix is enumerated in this run. When $\beta = 0$, the comparisons do not depend on h, so the single run also retains x^*. Preprocessing preserves all gains and records its constant score change. Scoring candidates with that constant included proves the claim. □

4.10 Random clauses with weaker symmetry assumptions

The next result applies the count adjustment to two explicit distributions. It permits probabilities tending to zero and unit clauses chosen after seeing the binary formula. Its density threshold supplies enough random variation for the accumulated comparisons to remain effective. These additional freedoms distinguish it from the fixed-q result above.

Corollary 4.13 (Random formulas with several clauses per pair)**.** *There are absolute constants $C, \gamma_0 > 0$ such that, for sufficiently large n, each of the following models with $d \geq C \log n$ admits a deterministic polynomial-space exact algorithm taking time $O^*(2^{(1-\gamma_0)n})$ with probability at least $1 - 1/n$.*

1. ***Independent clauses.*** *On every variable pair, include each equal-polarity clause with probability $q(1+\epsilon)$ and each mixed-polarity clause with probability $q(1-\epsilon)$, independently across all clauses. Here $\epsilon \in [-1, 1]$, $q(1 + |\epsilon|) \leq 1/2$, and $d = 4q(n-1)$.*
2. ***A random set of pairs.*** *Select each pair for $S \subseteq \binom{[n]}{2}$ independently with probability p. On each selected pair, include each of the four clauses independently with probability $1/4$; include none on the other pairs. Here $d = p(n-1)$.*

Both models allow arbitrary additional unit clauses, even chosen after the binary clauses are known. The algorithm is exact on every input; the probability qualifies its running time.

Here $d = 2\mathbb{E}m_2/n$, so the expected number of binary clauses per variable is $d/2$, and log is natural. In the first model with $\epsilon = 0$, and in the second model, the mean pair coefficient is zero and one run with $\beta = 0$ suffices. The original strict pruning constraints may also be used in these zero-mean cases: they retain an optimum by the earlier correctness argument and are stronger than the comparisons with zero slack.

How the running-time argument applies. Appendix B provides an independent-pair criterion (Proposition B.1) for the accumulation theorem in Appendix C. It reduces the task here to checking the mean, variance, and boundedness of the pair coefficients J_{uv}. We perform those checks explicitly below. The criterion's probability event depends only on J, while its running-time conclusion is uniform over the linear coefficients. This is why even unit clauses chosen after the random draw are allowed.

Proof of Corollary 4.13. Use CountComparisons with singleton rounds through $T = \lfloor n/2 \rfloor$. Its correctness follows from Lemma 4.12.

1. State the condition to be checked. Set $\tau = 1/2$ and $C_0 = 10^9/\tau$, as in Proposition B.1. All pairs have the same distribution, so its variance-ratio parameter is $\chi = 1$. Write $\mu = \mathbb{E}J_{uv}$

and $\nu = \mathrm{Var}(J_{uv})$. Let B bound $|J_{uv} - \mu|$, and put $\sigma_0 = |\mu|$. It suffices to verify independence across pairs and

$$n\nu \geq C_0(B^2 \log n + \sigma_0^2).$$

For sufficiently large n, the proposition then supplies all the conditions needed by the sequential case of Theorem C.2.

2. Independent clause inclusions. Let $B_{++}, B_{--}, B_{+-}, B_{-+}$ be the four independent inclusion indicators on one pair. By (14),

$$J = B_{+-} + B_{-+} - B_{++} - B_{--}.$$

Taking expectations and summing the variances of independent indicators gives

$$\mu = -4q\epsilon, \qquad \nu = 4q - 4q^2(1 + \epsilon^2) \geq 2q.$$

The last inequality follows because every inclusion probability is at most $1/2$, so each indicator has variance at least half its mean. Since $|J| \leq 2$ and $|\mu| \leq 2$, take $B = 4$ and $\sigma_0 \leq 2$. Also $n\nu \geq d/2$. For sufficiently large n, the bound $d \geq 34C_0 \log n$ implies

$$n\nu \geq 17C_0 \log n \geq C_0(16 \log n + 4),$$

as required. Use the empirical mean β; when $\epsilon = 0$, the known zero mean $\beta = 0$ is also allowed.

3. A random support. On one pair write $J = EZ$, where E is its Bernoulli-p selection indicator, and Z is the difference between the sums of two positive and two negative independent Bernoulli-$1/4$ indicators. The variables E and Z are independent. Therefore

$$\mathbb{E}J = 0, \qquad \nu = \mathbb{E}J^2 = p\,\mathbb{E}Z^2 = \frac{3p}{4}, \qquad |J| \leq 2.$$

The pair coefficients are independent, even though the four clause indicators on a pair share its selection indicator. We may take $B = 2$, $\sigma_0 = 0$, and $\beta = 0$. Since $n\nu \geq 3d/4$, the condition $d \geq 6C_0 \log n$ is sufficient.

4. Obtain a fixed exponential saving. In both cases Proposition B.1 gives its event with probability at least $1 - 1/n$, with $K = 9$ and $g = \tau^{3/2}/12$. In the notation of Theorem C.2, the sequential case then has

$$\alpha = \frac{81}{\sqrt{\tau}}, \qquad c_* = \frac{\tau^3}{144 \cdot 64 \cdot 81^2 \cdot 9}.$$

That theorem bounds the time by $O^*((n+1)(2^{\tau n} + 2^n e^{-c_* n}))$. The factor $n + 1$ accounts for count guesses and is polynomial; it is absent when $\beta = 0$. At $\tau = 1/2$, $\gamma_0 = 10^{-10}$ is smaller than both $1 - \tau$ and $c_*/\ln 2$, so it gives the claimed binary-exponent saving. The common threshold $C = 34C_0$ suffices for the two models. The event concerns only the binary pair coefficients, so the conclusion holds for every choice of additional units on that event. □

5 Max-Cut families

Two familiar graph families illustrate accumulation without an initial literal gap: regular graphs with small nontrivial eigenvalues, and graphs drawn from the stochastic block model. Encode an edge $\{u, v\}$ by $(u \vee v)$ and $(\neg u \vee \neg v)$. It contributes $1 + [x_u \neq x_v]$, so maximizing the resulting formula is exactly Max-Cut up to the constant number of edges. Every variable has equal positive and negative binary literal degrees. Additional units are arbitrary linear fields, and both results below allow them even when they depend on the graph.

The regular-graph algorithm uses only gains, degrees and a guessed count; its spectral hypothesis is used in the analysis. For hidden communities, the algorithm also computes approximate spectral directions. This is an explicit preprocessing step, described after the class statement.

5.1 Max-Cut on graphs with small nontrivial eigenvalues

For a d-regular simple graph G, let $\lambda(G) = \max_{i \geq 2} |\lambda_i(A_G)|$, retaining multiplicities.

Corollary 5.1 (Max-Cut with external fields)**.** *Fix $\lambda_0 \geq 1$ and $0 < \varepsilon \leq 1/2$. There is $c = c(\lambda_0, \varepsilon) < 2$ such that every d-regular graph satisfying*

$$d \leq (1-\varepsilon)n, \qquad \lambda(G) \leq \lambda_0 \sqrt{d}$$

admits a deterministic polynomial-space Max-Cut algorithm of time $O^*(c^n)$. The same bound holds after adding any multiset of unit clauses to the Max 2-SAT encoding with $(u \vee v)$ and $(\neg u \vee \neg v)$ on each edge.

For $d \geq 64\lambda_0^2$, use Algorithm 7 with $\beta = -2d/(n-1)$, $T = \lfloor n/2 \rfloor$, and two prefix rounds of sizes $\lfloor n/4 \rfloor$ and $T - \lfloor n/4 \rfloor$. For smaller d, branch outside a greedy independent set and optimize its vertices separately. The solver computes no spectral quantities; the eigenvalue condition is used only in the following analysis.

Proof. If $d < 64\lambda_0^2$, greedily find an independent set of size at least $n/(d+1)$, branch on its complement, and optimize its variables independently. This already gives a uniform fixed saving in this range. Assume henceforth $d \geq 64\lambda_0^2$ and n is sufficiently large in terms of the fixed parameters; finitely many smaller sizes can be absorbed in the constant.

Apply the round version of Theorem C.2 with $\tau = 1/2$, two prefix rounds of sizes $\lfloor n/4 \rfloor$ and $\lfloor n/2 \rfloor - \lfloor n/4 \rfloor$ in any order, and

$$J = -2A_G, \quad P = \beta \mathbf{1}\mathbf{1}^{\mathsf{T}}, \quad \beta = -2d/(n-1), \quad N = -2A_G - \beta(\mathbf{1}\mathbf{1}^{\mathsf{T}} - I).$$

Use exact counts $0, \ldots, n$, so $s_t = |\beta|$. The matrix N annihilates $\mathbf{1}$ and has other eigenvalues $-2\lambda_i(A_G) + \beta$. Since $|\beta| \leq 2$,

$$\|N\| \leq 3\lambda_0 \sqrt{d}, \qquad D_N = 4d(1 - d/(n-1)) \geq 2\varepsilon d$$

for $n \geq 2/\varepsilon$. Thus $K = 9\lambda_0^2/(2\varepsilon)$ suffices.

An edge between the two prefix rounds contributes at least ε^2 to W_t. The expander mixing inequality gives, for $n \geq 80$,

$$e(B_1, B_2) \geq \frac{d|B_1||B_2|}{n} - \lambda(G)\sqrt{|B_1||B_2|} \geq \frac{dn}{32} - \frac{d}{2} \geq \frac{dn}{40}.$$

Indeed $|B_1||B_2| \geq (n/4-1)^2$, $\sqrt{|B_1||B_2|} \leq n/4$, and $\lambda(G) \leq d/8$. The mixing inequality itself follows by writing each indicator vector as its constant component plus a vector perpendicular to $\mathbf{1}$ and applying the spectral bound to the latter components. Consequently

$$\sum_{t \in B_2} \sqrt{W_t} \geq \varepsilon \sum_{t \in B_2} \sqrt{\deg_{B_1}(t)} \geq \frac{\varepsilon e(B_1, B_2)}{\sqrt{d}} \geq \frac{\varepsilon n \sqrt{d}}{40}.$$

Since $\sqrt{D_N} \leq 2\sqrt{d}$, choose $g = \varepsilon/(320\sqrt{3})$. Finally $\sum_{t \leq T} s_t \leq 2d = o(n\sqrt{d})$ uniformly for $d \leq n$, so the slack condition holds for all sufficiently large n. There are $n+1$ runs and all theorem parameters are fixed. □

Concrete graph families. Paley graphs on $q \equiv 1 \pmod 4$ vertices, for prime powers q, have degree $(q-1)/2$ and nontrivial eigenvalues $(-1 \pm \sqrt{q})/2$ [13, Section 9.1.2]. Thus $\lambda_0 = 1$, $\varepsilon = 1/2$ suffice for all sufficiently large q. Connected nonbipartite Ramanujan graphs satisfy $\lambda(G) \le 2\sqrt{d-1}$, so they give another family when the density restriction holds. Bipartite Ramanujan graphs have the additional eigenvalue $-d$ and do not satisfy this hypothesis at growing degree; ordinary Max-Cut on bipartite graphs is of course immediate. For each fixed $a > 0$, uniform random d-regular graphs with $n^a \le d \le n/2$ satisfy $\lambda(G) \le C_a\sqrt{d}$ with probability at least $1 - 1/n$ by Tikhomirov and Youssef [14]. The corollary therefore supplies a fixed saving on that random graph family as well. These are applications of one deterministic class theorem, not separate algorithm designs. No priority claim is made here.

5.2 The stochastic block model

Fix a partition $V = V_1 \cup \cdots \cup V_k$ and a symmetric probability matrix $\Pi = (\pi_{ab})_{a,b \le k}$. Include each unordered pair of distinct vertices independently, with probability π_{ab} when its endpoints belong to V_a, V_b. This is the stochastic block model [15]. The partition can be arbitrary and need not be balanced. Alternatively, it can be drawn first; the result below holds uniformly conditional on it. The number k is fixed, while the probabilities may depend on n.

Define the smallest and largest edge variances by

$$v_{\min} = \min_{a,b} \pi_{ab}(1 - \pi_{ab}), \qquad v_{\max} = \max_{a,b} \pi_{ab}(1 - \pi_{ab}).$$

These conditions measure the amount of random variation, not how well the communities can be distinguished.

Theorem 5.2 (Exact Max-Cut on the stochastic block model)**.** *Fix $k \ge 1$ and $\chi \ge 1$. There are constants $C = C(k, \chi)$ and $\gamma = \gamma(k, \chi) > 0$ such that, whenever*

$$0 < v_{\min}, \qquad v_{\max} \le \chi v_{\min}, \qquad n v_{\min} \ge C \log n, \tag{16}$$

Max-Cut on a graph from this model can be solved by a deterministic polynomial-space exact algorithm taking $O^*(2^{(1-\gamma)n})$ time with probability $1 - o(1)$. The algorithm is given the fixed class parameters k, χ, but neither the partition nor the probabilities π_{ab}. The same statement holds with arbitrary additional unit clauses, even chosen after the graph is drawn. Exactness holds on every input; only the time bound is probabilistic.

Proof. We verify the auxiliary discovery lemma in Appendix D. In the Max 2-SAT encoding, $J_{uv} = -2A_{uv}$ for the adjacency indicators A_{uv}. Conditional on the partition, these coefficients are independent, with

$$\mathbb{E} J_{uv} = -2\pi_{ab}, \qquad \operatorname{Var} J_{uv} = 4\pi_{ab}(1 - \pi_{ab}), \qquad |J_{uv} - \mathbb{E} J_{uv}| \le 2.$$

Let $Z_{ua} = [u \in V_a]$. The matrix $\bar{A} = -2Z\Pi Z^{\mathsf{T}}$ agrees with the means off the diagonal, has rank at most k, has diagonal entries of absolute value at most 2, and satisfies $\|\bar{A}\| \le 2n$ by its maximum absolute row sum. In Lemma D.1, take

$$r = k, \quad \nu = 4v_{\min}, \quad B = B_0 = \mu = 2.$$

Its threshold becomes $4nv_{\min} \ge C_0(4 \log n + 8)$, which follows from (16) after increasing C. The lemma's algorithm does not use $\bar{A}$ or its probabilistic parameters. Its good event depends only on the pair coefficients and its bound is uniform over linear fields, proving the assertion about units. □

The usual within/between model. When $\pi_{aa} = p$ and $\pi_{ab} = q$ for $a \neq b$, the assumptions reduce to

$$\max\{p(1-p), q(1-q)\} \leq \chi \min\{p(1-p), q(1-q)\}, \qquad n \min\{p(1-p), q(1-q)\} \geq C \log n.$$

No separation between p and q is required; in particular $p = q$ is allowed. Concrete regimes are:

- *Dense blocks:* all $\pi_{ab} \in [\delta, 1-\delta]$ for a fixed $0 < \delta \leq 1/2$. The hypotheses hold for sufficiently large n, with constants depending only on k, δ.
- *Sparse blocks:* $\pi_{ab} = c_{ab}\rho$ with fixed $0 < a \leq c_{ab} \leq b$, $\rho = o(1)$, and $n\rho \geq C' \log n$.
- *Almost complete blocks:* $\pi_{ab} = 1 - c_{ab}\rho$ with the same bounds on c_{ab} and ρ.

For the last two regimes, eventually $a\rho/2 \leq \pi_{ab}(1-\pi_{ab}) \leq b\rho$; one can take $\chi = 2b/a$ and increase C'. The saving is fixed throughout each regime.

How the hidden-community algorithm uses the graph. The group probabilities create a systematic part of the comparisons. The algorithm accounts for it without recovering the community labels:

1. Compute a rational approximate diagonalization of $J = -2A_G$ and retain at most k directions with the largest absolute diagonal coefficients. This takes polynomial time; the precise construction is in Appendix D.
2. Enumerate polynomially many guesses for the projections of an assignment onto those directions. Each guess defines adjusted gains and an explicit error allowance.
3. In each run use two prefix rounds, add hard comparisons in the nonpreferred branches, and enumerate the suffixes satisfying their sum. Evaluate the original objective and retain the best candidate over all runs.

At least one run preserves an optimum. Constraints are local to their run. Thus the spectral step changes how gains are adjusted; the search is still pruned by accumulated hard comparisons. Lemma D.1 and the generator in Appendix C specify the arithmetic, grid sizes and all-input fallback exactly.

When the communities are supplied. If the partition and rational probabilities of polynomial encoding length are given, their role has an elementary description. Guess the number h_a of vertices assigned 1 in each V_a: there are at most $(n+1)^k$ tuples. For $t \in V_a$, use the adjusted gain

$$H_t^{\boldsymbol{h}}(x) = G_t(x) - 2\sum_{b=1}^{k} \pi_{ab} \left(h_b - \sum_{w \in V_b \setminus \{t\}} x_w \right).$$

On an assignment with these counts it equals $G_t(x) - 2\pi_{aa}x_t$, so the comparison allowance is $2\pi_{aa}$. The same two-round accumulation proof applies, with total allowance at most $2n$. The centered independent entries satisfy the required norm and cross-round mass bounds under (16); this is the proof of Lemma D.1 with the mean supplied and no approximation error. Consequently the same type of fixed exponential saving follows using only group counts and gains. The spectral preprocessing in the theorem removes the need to supply this information.

Deterministic endpoints and the role of the assumptions. The variance hypotheses are sufficient, not necessary. If $p = q = 1$, the graph is a clique and ordinary Max-Cut has value $\lfloor n^2/4 \rfloor$. With units, sort the linear fields: for each count h, take the h largest fields and

maximize $h(n-h)$ plus their sum. This is polynomial time. A clique has $\Theta(2^n/\sqrt{n})$ optimal Boolean assignments, so enumerating every optimum would give no fixed exponential saving; count compression avoids that enumeration. More generally, when a supplied fixed-block model has only probabilities 0 or 1, the objective depends on the group counts, and sorting fields within each group gives an $n^{O(k)}$ algorithm. These endpoints use repetition, whereas Theorem 5.2 uses random variation. The theorem does not assert that arbitrary intermediate or highly unequal variances are covered.

6 Repeated interactions across a cut

6.1 Several repeated interaction profiles

An interaction profile toward B is the vector $(J_{uv} : v \in B)$. Suppose $A = A_1 \dot\cup \cdots \dot\cup A_t$ and all vertices of A_i have the same profile toward B. Put

$$P = \prod_{i=1}^{t} (|A_i| + 1), \qquad \Gamma = |A| - \log_2 P.$$

Theorem 6.1 (Compression by group counts)**.** *Given A, B and their profile classes, the instance can be solved in polynomial space and time*

$$O^*\left(2^{n-|B|} + 2^{n-\Gamma}\right).$$

Proof. Fix the outside variables. The effect on B depends on an assignment to A only through its vector of group counts. For each such vector, find the best assignment to A with those counts, then enumerate B. Count-restricted generation over all vectors visits exactly $2^{|A|}$ assignments to A; enumeration on B costs $P2^{|B|}$. Stream the vectors and keep one representative at a time, so space stays polynomial even when P is exponential. Multiply by the outside cost. □

For example, $t = o(n)$ and $|A|, |B| = \Omega(n)$ imply $\log_2 P \le t \log_2(1 + |A|/t) = o(n)$ and hence a fixed saving. This is a statement about repeated profiles across a cut; it does not require literal cliques within either side.

6.2 Literal bicliques and a linear number of defects

Let A, B be disjoint sets, with one prescribed literal for each variable. A good cross pair contains exactly the prescribed clause. A missing clause, a different clause, or an additional binary clause makes the pair defective. All clauses outside the cross pairs and all units are arbitrary.

Corollary 6.2 (An exact literal biclique)**.** *Given an exact literal biclique with $|A| = a$ and $|B| = b$, Max 2-SAT can be solved in polynomial space and time $O^*(2^{n-\min(a,b)})$.*

Proof. Complement variables so that the prescribed literals are positive. All rows of $J_{A,B}$ then agree, so there is one profile class. Theorem 6.1 gives $O^*(2^{n-b} + (a+1)2^{n-a}) = O^*(2^{n-\min(a,b)})$. □

Lemma 6.3 (Extracting a core from average defect)**.** *If the A–B cross pairs contain q defects, there is an exact literal biclique with both sides of size at least*

$$k = \left\lfloor \min\left(a, \frac{ab}{a+q}\right) \right\rfloor.$$

When A, B and their literals are supplied, it can be found in polynomial time.

Proof. Take the k vertices of A with smallest defect counts. Their total number of defects is at most kq/a. Deleting their defective partners leaves at least $b - kq/a \geq k$ vertices of B. □

Thus $a, b = \Omega(n)$ and $q = O(n)$ suffice. The condition allows a linear number of defective pairs, not an arbitrary fixed fraction of the $\Theta(n^2)$ cross pairs.

Proposition 6.4 (Detection can be paid for)**.** *Fix $0 < \delta \leq 1/8$ and put $k = \lfloor \delta n \rfloor$. The existence of an exact literal biclique with both sides at least k can be detected in time $O^*(\binom{n}{k}2^k)$ and polynomial space. Upon detection, the instance can be solved in time*

$$O^*\left(2^{(H_2(\delta)+\delta)n} + 2^{n-k}\right).$$

Both exponents have a fixed saving for fixed $\delta > 0$.

Proof. Enumerate a set of k variables and its literal choices. Intersect their good literal neighborhoods and check whether it contains literals on at least k distinct outside variables. A promised biclique is eventually found. Run Corollary 6.2 once, after detection. The usual binomial entropy bound gives the displayed cost; $H_2(1/8) + 1/8 < 0.669$. □

6.3 Why bias and cross structure differ

Let $n = 2s$, with equal sides A, B. Put positive-positive clauses on every cross pair and negative-negative clauses on every within-side pair. Every variable has positive degree s and negative degree $s - 1$, so its gap is only one, but Corollary 6.2 applies.

Conversely, on the complete graph, choose each incidence to be positive independently with probability 3/4. With high probability, every variable has a positive polarity gap linear in n. Yet there is no prescribed literal biclique with linear sides and only $O(n)$ defects, with high probability. For any fixed sides and literal choices, each cross pair is good with probability at most 9/16, independently across pairs. The probability that all but $O(n)$ of $\Theta(n^2)$ pairs are good is $e^{-\Omega(n^2)}$. There are only $e^{O(n)}$ choices of sides and literals, so a union bound proves the claim.

7 Collective redundancy in bounded blocks

For a block B and two different assignments $a, b \in \{0,1\}^B$, the gain from replacing a by b is

$$F(b, y) - F(a, y) = \kappa_B(a, b) + \sum_{v \notin B} d_v y_v, \qquad d_v = \sum_{u \in B} J_{uv}(b_u - a_u). \tag{17}$$

Thus equality of outside effects concerns sums of interaction rows. Pairwise similar neighborhoods are only one source of such equality.

Proposition 7.1 (An unconditional block clause)**.** *If $\kappa_B(a, b) + \sum_{v \notin B} \min(0, d_v) > 0$, pattern a can be forbidden. The weak inequality also suffices if the common tie convention prefers b to a. If all $d_v = 0$, one of a, b can always be forbidden.*

Proof. The expression is the minimum gain over all assignments outside the block. Fix one order on complete assignments to break objective ties, for example lexicographic order in a fixed variable order. Among the optima choose the first assignment in this order. If it contains pattern a and replacing a by b always increases the score, this contradicts optimality. If the gain is only nonnegative but the fixed order prefers the replacement for every outside assignment, either the score increases or an earlier optimum is obtained. Both are contradictions. Thus the selected optimum avoids a. The same reasoning applies simultaneously to all such exclusions using that one order. If all $d_v = 0$, the comparison is independent of the outside assignment, so the smaller score, or the later tied pattern, can be excluded. □

For instance, if four vertices satisfy $J_{u_1v} + J_{u_2v} = J_{u_3v} + J_{u_4v}$ for every outside v, the patterns 1100 and 0011 have identical outside effects. Their worse pattern is excluded by a clause of length four. No two of the four rows need agree.

Theorem 7.2 (Blocks with exceptional outside variables)**.** *Suppose $B_1, \dots, B_q$ are disjoint nonempty blocks of size at most b. For each block there are two patterns for which the discrepancy set $D_i = \{v \notin B_i : d_v \neq 0\}$ has size at most r. Given these certificates, the instance can be solved in polynomial space and time*

$$O^*\left(2^n(1-2^{-b})^{q/(2r+1)}\right).$$

For constant b, r and $q = \Omega(n)$ this is a fixed saving, at any density.

Proof. Make a graph on the blocks, joining i, j if D_i meets B_j or D_j meets B_i. Orient each such witness from its discrepancy source; there are at most r outgoing witnesses per block. Every induced graph has at most r times its number of vertices in edges, so it is $2r$-degenerate. A greedy coloring with $2r+1$ colors gives an independent set I of size at least $q/(2r+1)$.

Assign all variables outside $\bigcup_{i\in I} B_i$. Every discrepancy variable of a retained block is now fixed. Its replacement gain is a constant independent of the other retained blocks, so one of its two patterns can be excluded. Choose ties consistently. The surviving states number

$$2^{n-\sum_{i\in I}|B_i|}\prod_{i\in I}(2^{|B_i|}-1) \le 2^n(1-2^{-b})^{|I|}.$$

Direct enumeration over these block states has this cost and polynomial space, and preserves an optimum for each outside assignment. □

For fixed b, all candidate blocks and pairs of patterns can be checked in polynomial time. If a disjoint family of size q exists, a maximal greedy disjoint selection has size at least q/b: each chosen block intersects at most b blocks of the promised disjoint family. Thus the existential linear-family version also has a polynomial-time certificate search, with a constant-factor loss in the saving.

This gives a robust structural exit. It does not justify excluding blocks with an unbounded number of exceptions at the same cost: the displayed exponent improvement can tend to zero when r grows.

8 Comparing regions and retaining information

A region of complete assignments is *covered* when the best score already recorded is at least the score of every assignment in that region. Exhaustive exploration establishes coverage; a sound domination certificate can establish it as well. In the following rule, coverage is an explicit hypothesis about the entire stated region.

The preceding arguments motivate a broader learning rule. Fix the same assigned variable set in two prefixes α, β, and let their residual objectives be $F_\alpha(y), F_\beta(y)$. Suppose the region $R_\alpha(y)$ under α has already been covered.

Proposition 8.1 (Conditional comparison)**.** *Within the β-branch it is sound to impose*

$$\neg R_\alpha(y) \ \vee \ [F_\beta(y) > F_\alpha(y)]. \tag{18}$$

The difference $F_\alpha - F_\beta$ is affine, because their remaining quadratic terms agree.

Proof. Every excluded completion lies in a covered region under α and has score at most its counterpart there. It cannot improve the incumbent. Restricting (14) proves the affine-difference claim. □

For equal residual objectives up to a favorable additive constant, this retains $R_\beta \wedge \neg R_\alpha$. If $R_\alpha = C \wedge P$ and $R_\beta = C \wedge Q$, it simplifies to $C \wedge Q \wedge \neg P$. For example, if P is the clause $y_1 \vee y_2 \vee y_3$, the second branch needs only $y_1 = y_2 = y_3 = 0$. If P is a conjunction of s ordinary clauses each on at least ℓ distinct remaining variables, its complement contains at most $s2^{|U|-\ell}$ assignments. Enumerating the subcubes that falsify the clauses gives the same upper bound, allowing duplication. The cost of covering α must still be included.

The negation of a conjunction is not generally an ordinary clause. To learn a clause, find a partial suffix assignment τ for which the current constraints certify R_α and $F_\alpha - F_\beta \geq 0$. Then forbid τ within β. If retained across backtracking, the clause must include the guard specifying the β-prefix. Sufficient certificates can use syntactic clause witnesses and the exact minimum of an affine function over a Boolean subcube; a general implication oracle is not assumed.

8.1 A bounded cache for exact repeated effects

Normalize a residual objective as $c_\alpha + G_\alpha(y)$ with $G_\alpha(0) = 0$. A completed node can store its G_α and constant c_α. If the whole subtree is covered, a later node with the same G and a no-larger constant can be discarded regardless of its inherited auxiliary constraints. Equality of soft objectives is not an assertion that the constrained subproblems are identical; whole-region coverage is what makes this rule sound.

For a fixed variable order, the quadratic terms on unassigned variables at a given depth are fixed, so a cache key needs only the residual linear coefficients. Retain at most $n + 1$ records at each depth, evicting the least recently accessed. Only completed or already certified nodes are inserted. Eviction loses a possible cut but cannot invalidate one.

Proposition 8.2 (Bounded caching on a repeated-profile cut)**.** *Suppose all rows of $J_{A,B}$ agree up to sign, and use the order $V \setminus (A \cup B), A, B$. The preceding cache, together with any coverage-preserving cuts in this fixed search, gives running time $O^*(2^{n-\min(|A|,|B|)})$ and polynomial space.*

Proof. For each fixed outside assignment, the residual field on B depends on A only through $\sum_{u \in A} s_u x_u$, with $s_u \in \{-1, 1\}$. There are at most $|A| + 1$ field vectors. At depth $|V \setminus B|$, all records touched during this phase fit in the cache, so a record touched in the phase cannot be evicted during it. A field vector is explored again only when its constant strictly improves. Since $c_\alpha = F(\alpha, 0) \in \{0, \ldots, m\}$, there are at most $m + 1$ such explorations per vector. The phase costs $O^*(2^{|A|} + (|A| + 1)(m + 1)2^{|B|})$. Multiply by the outside cost. □

The explicit unweighted input is relevant to the polynomial bound on record improvements. Binary-encoded arbitrary weights require a different argument. This proposition is a structured application of caching, not a claim that a polynomial cache always finds all useful similar states.

9 Toward one branching algorithm

A concrete candidate uses the same search operations throughout:

1. Maintain the residual objective, the scoped comparison constraints, and a bounded cache of covered residuals.

2. Reject a node if its constraints have an infeasible linear relaxation, or a cached covered residual dominates it.

3. Select the next variable in a prescribed order. Try first the value preferred when remaining variables are set to 1/2.

4. After that child is covered, impose its affine comparison in the other child and recurse. Record coverage on returning.

The linear relaxation is an explicit polynomial-time test, not an oracle for Boolean satisfiability. Single-variable comparisons and cache cuts are sound together by induction on the order in which regions become covered.

Proposition 9.1 (Why no enumeration switch is needed for the basic bounds)**.** *For a fixed order, this engine inherits the bounds of Theorems 3.2 and 4.3, the sharper Corollary A.1, and the unbiased cases of Corollary 4.13. With the indicated repeated-profile order, it also inherits the preceding caching bound.*

Proof. The bias analysis applies with no forced steps when only binary branching and the sum test are used. Linear infeasibility detects every rejection made by that sum test, so the same time bound holds for this engine. For the random-sign argument, fix any prefix at the analysis depth T, and freeze the sum of constraints already created there. Every later node retained by the relaxation has a fractional completion satisfying this one affine inequality. The minimum of one affine function over a Boolean subcube is attained at a Boolean vertex, so this node also has a Boolean completion satisfying the frozen sum. If that sum has N completions, there are at most N retained descendants at any suffix depth. Thus the cost is bounded by (13), even though the algorithm continues learning and never switches modes. Positive weights cause no difficulty: the linear relaxation implies every positive weighted sum of its comparisons, including that of Theorem 4.3. For the unbiased models of Corollary 4.13, take $P = 0$. The preferred values then agree with those of this engine, and each strict second-child comparison implies the corresponding weak comparison used in (25). The same frozen-sum argument therefore bounds its descendants by the surviving mass in Theorem C.2. The cache only removes nodes of this fixed-order, fixed-preference search. The structured bound was proved directly above. □

Ordering still has a cost. The repeated-profile bound is not asserted for every order. For a fixed $0 < \delta \leq 1/8$, one explicit policy enumerates $k = \lfloor \delta n \rfloor$ element sets B, groups outside rows on B up to sign, and searches for a group A of size k. A successful test supplies an order with the structured bound; if none succeeds, use a fixed default order. The search costs $O^*(\binom{n}{k})$. On every draw from a covered unbiased random model, either the test succeeds and gives its unconditional structured bound, or the default order is used and the exceptional event is contained in the exceptional event for that fixed order. Bias is order-independent. The resulting bound includes the detection cost, the possible successful case $O^*(2^{(1-\delta)n})$, and the default-order bound. Thus it retains a fixed saving, though not necessarily the same numerical saving as the default-order theorem. For a promised biclique family, choose δ below its guaranteed side fraction (or the extracted core fraction). This is a uniform solver for literal bias, the unbiased random models, and the indicated structured cuts, with an explicit ordering search. A purely local adaptive ordering rule with the same guarantees remains open.

The count guesses used for biased random formulas and regular-graph Max-Cut change the gains used to generate comparisons. Hidden communities use the same principle with computed projection directions. The round version also fixes preferences earlier. These operations are handled by the single generator in Theorem C.2, with fresh constraints for every run. One can interleave this generator (using the explicit constructions for the desired classes) and the preceding engine, stopping when an exact solver finishes. A fixed number of polynomial-space solvers costs only a constant factor over the fastest one, up to simulation overhead. This gives a uniform solver for the combined class guarantees. It does not require transferring cache records between different auxiliary objectives. A local adaptive policy achieving all these guarantees without this interleaving remains open. Nor do we claim that arbitrary pre-branching clauses

can be added without reconciling their representative order with its coverage invariant. These are separate integration tasks.

10 Open directions: exploiting redundancy

10.1 A structural explanation of persistent correlation

The most immediate target is to weaken the cancellation assumptions of the balanced class. The bias proof already permits arbitrary correlations. In the balanced argument, coherent comparison coefficients can make the remaining uncertainty large. Some coherent families are resolved by group counts or by the bounded-block theorem. This suggests analyzing the search only after applying those available comparisons.

Open question 10.1 (From poor cancellation to exploitable comparisons)**.** When a large set of generated comparisons fails to accumulate useful pruning, can one efficiently find either a family of bounded block replacements with few exceptions, or a larger group of assignments whose outside effects admit affordable compression?

Absence of an exact short relation gives only a nonzero signed sum of neighborhood vectors. It does not give the quantitative cancellation bound needed in the present proofs. Theorem 7.2 handles a robust version with a constant number of exceptions. Extending this to a growing number of exceptions without losing a fixed exponent saving is a concrete intermediate problem. It may require using the learned conditional comparisons rather than assigning every exception independently.

10.2 Similarity of regions, not just equality of formulas

Open question 10.2 (Affordable regional coverage)**.** Can the algorithm maintain polynomially many conditional comparison certificates that cover most repeated optimization behavior, even when no two residual objectives are identical?

At a fixed depth, residual linear vectors provide a direct representation for comparisons. Differences supported on a small set give local tests; common learned clauses cancel from the regional comparison. The relevant cost is finding and retaining the certificates as well as exploiting them. The affine-difference observation avoids general formula isomorphism, but does not make arbitrary Boolean implication tests polynomial-time.

10.3 An amortized measure for overlapping constraints

Deleting a satisfied clause and shortening an unsatisfied clause have different effects on the surviving completion space. A useful measure should account for both, and should charge only genuinely new restrictions. For a constraint family C on U, the ideal information measure is

$$|U| - \log_2 \#\{y \in \{0,1\}^U : y \models C\}.$$

It is not generally efficiently computable and does not by itself pay for enumeration. A tractable bound based on sums, selected subfamilies, or conditional coverage could support an adaptive branching rule. The measure should also charge the discovery and storage of comparisons.

10.4 Randomized exploration with exact answers

Las Vegas randomization may help choose comparison blocks or enumerate survivors. It must preserve the exact optimum on every run. If a capped attempt costs at most B and succeeds

with probability at least p, independent restarts cost at most B/p in expectation. A small number of satisfying full assignments does not supply this success probability by itself. A useful result must bound the work spent on partial assignments without a surviving completion.

10.5 The long-term target

Desired theorem 10.3 (Growing density with a uniform saving)**.** Find a deterministic polynomial-space exact algorithm and a constant $\gamma > 0$ for all canonical unweighted Max 2-SAT formulas without repeated binary clauses, allowing complementary pairs, running in $O^*(2^{(1-\gamma)n})$ time. An intermediate target is the same guarantee for every formula up to a specified density bound $\rho(n) \to \infty$, with γ independent of n.

This is a desired result, not a consequence or a conjecture established by the present arguments. The current paper's contribution should be assessed through its explicit classes. A structural dichotomy is valuable if both outcomes are algorithmic and their costs combine into a fixed saving; merely showing that some useful relation exists is insufficient.

A A sharper constant for almost-regular random signs

Fix a simple graph G and a vertex order before drawing signs. Each edge contributes one clause, and its two literal signs are independent and uniform; all incidence signs are mutually independent. There are no initial unit clauses. The graph itself is arbitrary. Write $d = \min_v \deg(v)$ and $D = \max_v \deg(v)$.

Corollary A.1 (A fixed saving for random signs)**.** *For $d \geq 6$, there is a deterministic polynomial-space algorithm that is exact on every signing and, with probability at least $1 - 3e^{-nd/(1250D)}$, takes time*

$$O^*\left((2e^{-d/(1250D)})^n\right).$$

Its expected running time over the signs is $O^*((2e^{-d/(625D)})^n)$.

In particular, every regular graph gives the base $2e^{-1/1250} \simeq 1.99840064$, independently of its degree. This includes any sequence of regular graphs of degree $\omega(1)$, and the complete graph. Randomness belongs to the instance distribution; the solver neither guesses an answer nor has a probability of returning an incorrect optimum.

Proof. Learn (6) in the fixed order through depth $T = \lceil n/2 \rceil$, and then enumerate the sum as in Lemma 4.4. Let x be a uniform complete assignment, independent of the signs, and retain even constraints of already cut prefixes for this upper-bound analysis. Define

$$S = \sum_{i \text{ creator},\ i \leq T} (|g_i|/2 + 1), \qquad Q = \sum_{i \text{ creator},\ i \leq T} (L_i(x) - D_i^+/2),$$

where D_i^+ counts later neighbors. An assignment passes its prefix's summed inequality exactly when $S + Q \leq 0$. Therefore

$$\mathbb{E}_F W(F) = 2^n \mathbb{P}_{F,x}(S + Q \leq 0). \tag{19}$$

Reveal the incidence signs of vertices in order, together with their assigned values. Just before vertex i, let r_i count earlier neighbors whose partner literal in their clause with i is false. Conditional on the history, g_i is a sum of independent signs of weights one and two, of variance $Y_i = D_i^+ + 4r_i$. The fourth-moment bound for a Rademacher sum gives $\mathbb{E}|g_i| \geq \sqrt{Y_i/3}$. The

contribution X_i to S thus has conditional mean at least $\mu_i = \sqrt{Y_i}/(4\sqrt{3}) + 1/2$ and second moment at most D, for $D \geq 6$. The same nonnegative-increment estimate as above gives

$$\mathbb{P}\left(S < \sum_{i \leq T} \mu_i - a\right) \leq e^{-a^2/(2TD)}. \tag{20}$$

The earlier endpoint-falsity indicators are independent fair coins, including after averaging over x. Hence r_i has distribution $\mathrm{Bin}(D_i^-, 1/2)$, with $\mathbb{E}Y_i = D_i^+ + 2D_i^- \geq d$ and $\mathrm{Var}(Y_i) \leq 2\mathbb{E}Y_i$. For nonnegative Y of mean μ, the pointwise inequality $\sqrt{z} \geq 1 + (z-1)/2 - (z-1)^2/2$ yields $\mathbb{E}\sqrt{Y} \geq \sqrt{\mu} - \mathrm{Var}(Y)/(2\mu^{3/2})$. It follows that $\mathbb{E}\sum_i \mu_i \geq T\sqrt{d}/(4\sqrt{3})$. Changing one earlier-endpoint coin changes this sum by at most $1/(2\sqrt{3})$. There are at most $TD/2$ relevant internal edges, so bounded differences gives

$$\mathbb{P}\left(\sum_i \mu_i < T\sqrt{d}/(4\sqrt{3}) - b\right) \leq e^{-48b^2/(TD)}. \tag{21}$$

Finally reveal all later endpoints, without creating constraints after T. For each newly revealed vertex, the bad-value indicators from earlier creators use distinct unseen incidence signs. Conditional on the history, they are independent fair bits. There are at most TD relevant edge contributions in total. The conditional cosh bound therefore gives

$$\mathbb{P}(Q < -c) \leq e^{-2c^2/(TD)}. \tag{22}$$

No independence between S and Q is needed.

Take $(a, b, c) = (2/25, 1/100, 1/20)T\sqrt{d}$. Their sum is less than $T\sqrt{d}/(4\sqrt{3})$. Combining the three tails with (19) gives $\mathbb{E}_F W(F) \leq 3 \cdot 2^n e^{-2Td/(625D)}$. Markov's inequality implies

$$\mathbb{P}_F\left(W(F) > 2^n e^{-Td/(625D)}\right) \leq 3e^{-Td/(625D)}.$$

Use $T = \lceil n/2 \rceil$ and (13). The prefix cost 2^T is smaller than the asserted bounds, up to constant factors. □

The argument counts total remaining work, not merely the fraction of prefixes with a conflict. A satisfiable summed constraint can already justify a large saving. Adaptive variable selection based on unrevealed signs is not covered by this exposure proof.

B Independent-pair membership

For this appendix, $P_{uv} = \sum_k \beta_k t_k(u) t_k(v)$ with ternary t_k, $N = \mathrm{offdiag}(J - P)$, $\sigma = \sum_k |\beta_k|$, $T = \lfloor \tau n \rfloor$, $R_{tw} = N_{tw}[w > t]$ for $t \leq T$, and $W_t = \sum_{j<t} N_{tj}^2$. Exact count grids have $s_t \leq \sigma$. Thus the following imply the sequential case of Theorem C.2:

$$\|R\|^2 \leq KD_N, \qquad \sum_{t \leq T} \sqrt{W_t} \geq 4\sqrt{3}\, gn\sqrt{D_N}, \qquad \sigma \leq g\sqrt{D_N}/(8\tau). \tag{23}$$

The following is the independent-pair criterion of Proposition 12 in [16]. We write χ for the variance ratio to distinguish it from clause density. The probability event concerns J alone, which is why the linear coefficients and unit clauses can be arbitrary.

Proposition B.1 (Independent-pair membership)**.** *Fix $0 < \tau < 1$, $\chi \geq 1$, $C_0 = 10^9\chi/\tau$, and $n \geq 10^9\chi^2/\tau^2$. Suppose the coefficients J_{uv}, $u < v$, are independent, with*

$$|J_{uv} - \mathbb{E}J_{uv}| \leq B, \qquad 0 < \nu \leq \mathrm{Var}\, J_{uv} \leq \chi\nu.$$

Choose either of the following decompositions:

1. *All means equal μ, and $A_{uv} = \beta$, where β is the empirical mean over pairs. Put $\sigma_0 = |\mu|$.*
2. *A fixed supplied aligned decomposition has $A_{uv} = \mathbb{E}J_{uv}$. Put $\sigma_0 = \sum_k |\beta_k|$. In particular, if all means vanish, take $A = 0$ and $\sigma_0 = 0$.*

If $n\nu \ge C_0(B^2 \log n + \chi\sigma_0^2)$, then with probability at least $1 - 1/n$, conditions (23) *hold with $K = 9\chi$ and $g = \tau^{3/2}/(12\sqrt{\chi})$. The lower bound on n in Theorem C.2 also holds.*

Proof. Put $X_{uv} = J_{uv} - \mathbb{E}J_{uv}$, and let $\delta = \beta - \mu$ in the first case and $\delta = 0$ in the second. Then $N_{uv} = X_{uv} - \delta$ off the diagonal. Set $\eta = 1/100$.

We use scalar Bernstein concentration [10], Talagrand's convex concentration inequality [11], and the rectangular norm estimate of Bandeira and van Handel [9, Theorem 3.1 and Corollary 3.2]. The last estimate, for independent random signs ξ_{tw} and a fixed coefficient array a, is

$$\mathbb{E}\|(\xi_{tw}a_{tw})\| \le (1+\vartheta)\left(\sigma_1 + \sigma_2 + \frac{5\max_{t,w}|a_{tw}|\sqrt{\log n}}{\sqrt{\log(1+\vartheta)}}\right), \quad 0 < \vartheta \le 1/2,$$

where σ_1, σ_2 are the largest row and column Euclidean norms. For a convex 1-Lipschitz function of independent coordinates in $[-B, B]$, we use the tail $\mathbb{P}(f > \mathbb{E}f + 8\sqrt{\pi}B + r) \le 4e^{-r^2/(16B^2)}$.

Centering and total variance. Bernstein's inequality gives, in the empirical case, $|\delta| \le \delta_0 = 6\sqrt{\chi\nu\log n}/n$ except with probability at most n^{-3}. For example, applying it to $\sum_{u<v} X_{uv}$ gives a two-sided exponent at least $4.5(1 - 1/n)\log n$. The bound on n implies $\delta_0 \le 10^{-3}\sqrt{\nu}$ and $n\delta_0 \le 10^{-3}\sqrt{n\chi\nu}$. Bernstein applied to $\sum X_{uv}^2$ also gives

$$(1-\eta)n(n-1)\nu \le \sum_{u\ne v} X_{uv}^2 \le (1+\eta)n(n-1)\chi\nu.$$

In the empirical case, $D_N = n^{-1}\sum_{u\ne v} X_{uv}^2 - (n-1)\delta^2$; in the expected case, $D_N = n^{-1}\sum_{u\ne v} X_{uv}^2$. Hence, outside another event of probability at most n^{-3},

$$(1-2\eta)(n-1)\nu \le D_N \le (1+\eta)(n-1)\chi\nu. \tag{24}$$

The operator norm. Let R_X have entries X_{tw} for $t \le T$, $w > t$, and zero otherwise. The entries in this array are independent. Symmetrize with an independent copy X' and put $\widehat{X} = X - X'$. Then $\mathbb{E}\|R_X\| \le \mathbb{E}\|R_{\widehat{X}}\|$. Condition on $\widehat{X}$ and introduce independent signs. Since its entries are symmetric, the displayed norm estimate with $\vartheta = \eta$ gives

$$\mathbb{E}\|R_X\| \le (1+\eta)(\mathbb{E}\sigma_1(\widehat{X}) + \mathbb{E}\sigma_2(\widehat{X})) + 102B\sqrt{\log n}.$$

Each squared entry is at most $4B^2$ with expectation at most $2\chi\nu$. Bernstein and a union bound show that every row-square sum is at most $2(1+\eta)n\chi\nu$, and every column-square sum at most $2(\tau+\eta)n\chi\nu$, except with probability $2ne^{-\eta^2 n\chi\nu/(5B^2)} \le n^{-3}$. Bounding each exceptional sum by $4nB^2$ gives

$$\mathbb{E}\sigma_1(\widehat{X})^2 \le 2(1+2\eta)n\chi\nu, \qquad \mathbb{E}\sigma_2(\widehat{X})^2 \le 2(\tau+2\eta)n\chi\nu.$$

Thus $\mathbb{E}\|R_X\| \le 2.89\sqrt{n\chi\nu}$, using $102/\sqrt{C_0} \le 0.004$. Convex concentration with $r = \eta\sqrt{n\chi\nu}$, together with $8\sqrt{\pi}B \le 10^{-3}\sqrt{n\chi\nu}$, now gives

$$\|R\| \le \|R_X\| + n|\delta| \le 2.902\sqrt{n\chi\nu}, \qquad \|R\|^2 \le 8.43n\chi\nu \le 9\chi D_N$$

outside an event of probability at most n^{-3}.

The earlier-neighbor mass. Put $t_0 = \lceil \tau n/10 \rceil$. For $t > t_0$, scalar Bernstein gives $V_t = \sum_{j<t} X_{tj}^2 \geq (1-\eta)(t-1)\nu$ with failure probability at most $e^{-\eta^2 \tau n \nu/(30B^2\chi)} \leq n^{-3}$. Consequently $\sqrt{W_t} \geq \sqrt{V_t} - |\delta|\sqrt{t-1} \geq 0.993\sqrt{(t-1)\nu}$. Summing and using $T \geq \tau n - 1$ yields

$$\sum_{t \leq T} \sqrt{W_t} \geq 0.635\, \tau^{3/2} n^{3/2} \sqrt{\nu} \geq 4\sqrt{3}\, gn\sqrt{D_N};$$

the last expression is at most $0.581\, \tau^{3/2} n^{3/2}\sqrt{\nu}$ by (24).

The allowed comparison slack. Equation (24) gives $g\sqrt{D_N}/(8\tau) \geq 0.0102\sqrt{\tau n \nu/\chi}$. In the expected case, $\sigma = \sigma_0 \leq \sqrt{n\nu/(C_0\chi)} \leq 10^{-4}\sqrt{\tau n\nu/\chi}$. In the empirical case, $\sigma = |\beta| \leq \sigma_0 + \delta_0 \leq 0.002\sqrt{\tau n\nu/\chi}$. This establishes the third condition.

There are at most $n+4$ failure events, each of probability at most n^{-3}, so their union has probability at most $1/n$. Finally $256\pi\tau K/g^2 = 256\pi \cdot 1296\, \chi^2/\tau^2 < 1.1 \cdot 10^6 \chi^2/\tau^2 \leq n$. □

C Analysis of accumulated comparisons

This appendix supplies the common analysis for the random-formula and Max-Cut families in Sections 4 and 5. It also specifies the projection comparison routine used for hidden communities. A few aggregate guesses account for systematic interactions; the remaining comparisons accumulate enough gap while a norm bound controls their combined fluctuation. The matrix formulation is proof machinery for these input families, rather than an additional class definition required of the reader.

C.1 One comparison algorithm

Choose a rational representation $P = \sum_{k=1}^{r} \lambda_k v_k v_k^{\mathsf{T}}$ of polynomial encoding length, and put $N = \operatorname{offdiag}(J - P)$. For each k, a finite rational grid $\mathcal{Y}_k$ must approximate every $\langle v_k, x\rangle$ to within a stated rational error $\delta_k \geq 0$. Define

$$s_t = \sum_k |\lambda_k|\, |v_k(t)|(\delta_k + |v_k(t)|), \qquad \mathcal{Y} = \prod_k \mathcal{Y}_k.$$

The grid is streamed; its description and each point have polynomial size. In a run indexed by $\hat{y} \in \mathcal{Y}$, use

$$a^{\hat{y}} = a + \sum_k \lambda_k \hat{y}_k v_k, \qquad F'_{\hat{y}}(x) = c + \langle a^{\hat{y}}, x\rangle + \sum_{u<v} N_{uv} x_u x_v.$$

Partition a prefix of $T = \lfloor \tau n \rfloor$ variables into consecutive rounds $B_1, \ldots, B_\ell$; the suffix is round $\ell + 1$. Let $\rho(t)$ be the round of t. Throughout round j, prefer $p_t = [\Delta_t \geq 0]$, where

$$\Delta_t = a_t^{\hat{y}} + \sum_{w:\rho(w)<j} N_{tw} x_w + \frac{1}{2} \sum_{w \neq t:\rho(w)\geq j} N_{tw}.$$

The preference uses only earlier rounds. Call t a creator when $x_t \neq p_t$, and let $\gamma_t(x)$ be the gain under $F'_{\hat{y}}$ of flipping it to p_t. At each depth-T prefix enumerate the suffixes satisfying

$$\sum_{t \text{ creator}} (\gamma_t(x) - s_t) \leq 0. \tag{25}$$

This is one affine inequality on the suffix, so Lemma 4.4 applies. Evaluate the original objective on every candidate. Candidates need not have projections near the guessed grid point.

Lemma C.1 (Sound comparisons with approximate projections)**.** *Every optimum is enumerated in a run that approximates its projections. The algorithm is exact for every supplied representation and round partition.*

Proof. Write $y_k = \langle v_k, x\rangle$. The difference between the original $0 \to 1$ gain at t and the auxiliary gain is

$$\sum_k \lambda_k v_k(t)(y_k - \hat{y}_k - v_k(t)x_t).$$

If $|y_k - \hat{y}_k| \le \delta_k$, its absolute value is at most s_t. Every flip at an optimum has original gain at most zero, so each creator has $\gamma_t \le s_t$ in an approximating run. Their sum satisfies (25). Weak inequalities retain ties. □

The accumulation proof below uses only nonnegative comparison allowances, their total size, and the fact that an optimum survives in at least one run. Thus a smaller allowance may replace s_t if its validity is proved directly. The same proof allows a joint finite list of aggregate guesses, rather than a Cartesian product of independent grids. This includes the supplied-community counts and diagonal allowances in Section 5.2; the running bound uses the number of joint guesses.

C.2 A general accumulation theorem

Let $D_N = n^{-1}\sum_{u\neq v} N_{uv}^2 > 0$ and $W_t = \sum_{w:\rho(w)<\rho(t)} N_{tw}^2$. The shared assumptions are

$$\sum_{t\le T}\sqrt{W_t} \ge 4\sqrt{3}\,gn\sqrt{D_N}, \qquad \sum_{t\le T} s_t \le \tfrac{1}{8}gn\sqrt{D_N}. \tag{26}$$

All constants below are independent of n.

Theorem C.2 (Accumulation after projection guessing)**.** *Assume* (26)*, with* $0 < \tau < 1$ *and* $g, K > 0$*. Either of the following conditions gives a deterministic exact algorithm in polynomial space and time*

$$O^*(|\mathcal{Y}|(2^{\tau n} + 2^n e^{-c_* n})).$$

1. ***Sequential comparisons.*** *Each prefix round is a singleton. Let* $R_{tw} = N_{tw}[w > t]$ *for* $t \le T$*, and assume* $\|R\|^2 \le KD_N$*. Set* $K_L = K$ *and* $c_* = g^2/(64\tau\alpha^2 K)$*.*
2. ***A fixed number of rounds.*** *There are* $\ell \ge 2$ *prefix rounds, with* ℓ *fixed, and* $\|N\|^2 \le KD_N$*. Set* $K_L = (\ell-1)K$ *and*

$$c_* = \min\left\{\frac{g^2}{32\tau\alpha^2 K_L}, \frac{g^2}{512\ell^2\tau K}, \frac{c_{\mathrm{HW}}g^2}{256\ell^2\tau K}, \frac{c_{\mathrm{HW}}g}{16\ell\sqrt{K}}\right\}.$$

Here $c_{\mathrm{HW}} > 0$ *is an absolute Hanson–Wright constant.*

In both cases, $\alpha = \max(2, 9\tau\sqrt{K_L}/(4g))$ *and* $n \ge 256\pi\tau K_L/g^2$*. Linear coefficients and units are arbitrary. For polynomially many grid points this is a fixed base below two.*

Exact count guessing is a special case: for $v_k \in \{-1,0,1\}^n$, use all integer counts and $\delta_k = 0$. Then $|\mathcal{Y}| \le (2n+1)^r$ and $s_t \le \sum_k |\lambda_k|$. Empirical mean centering takes $r = 1$, $v_1 = \mathbf{1}$, $\lambda_1 = \binom{n}{2}^{-1}\sum_{u<v} J_{uv}$, and only $n+1$ counts. If $P = 0$, there is one run and no comparison slack.

Lemma C.3 (A common gap estimate)**.** *Let* L *be the* $T \times T$ *matrix with* $L_{tw} = N_{tw}$ *when* $\rho(w) < \rho(t)$ *and zero otherwise. Suppose* $\|L\|^2 \le K_L D_N$ *and* (26)*'s gap bound holds. With* $M = \alpha\sqrt{K_L D_N}$ *and the theorem's lower bound on* n*,*

$$\mathbb{P}\left(\sum_{t\le T}\min\{|\Delta_t|, M\} < \tfrac{3}{4}gn\sqrt{D_N}\right) \le 4e^{-g^2 n/(16\tau K_L)}.$$

The probability is over a uniform Boolean assignment, with the formula fixed.

Proof. Put $u = 2x - \mathbf{1}$ and $H_t = a_t^{\hat{y}} + \frac{1}{2}\sum_{w \neq t} N_{tw}$. Write Lu for $L(u_1, \dots, u_T)^{\mathsf{T}}$, so $\Delta_t = H_t + \frac{1}{2}(Lu)_t$. Partition the indices into $B_0 = \{t : |H_t| \le M/2\}$ and its complement B_1. Since $\|Lu\|^2 \le K_L D_N T$, clipping the gaps at M loses at most $K_L D_N T/(8M)$ from $\sum_{B_0} |\Delta_t|$; use $(a-b)_+ \le a^2/(4b)$ with $a = |(Lu)_t|/2$, $b = M/2$. For $t \in B_1$, the clipped gap is at least $M/4$ whenever $|(Lu)_t| \le M/2$, and at most $4T/\alpha^2$ indices fail this condition. Thus, for $G(u) = \sum_{B_0} |\Delta_t|$,

$$\sum_{t \le T} \min\{|\Delta_t|, M\} \ge G(u) + \frac{M}{4}|B_1| - \frac{9T\sqrt{K_L D_N}}{8\alpha}.$$

A symmetric Rademacher sum of variance W_t has first absolute moment at least $\sqrt{W_t/3}$. Hence $\mathbb{E}|\Delta_t| \ge \frac{1}{2}\sqrt{W_t/3}$, uniformly in H_t. Also $W_t \le K_L D_N$ and $M/4 \ge \frac{1}{2}\sqrt{W_t/3}$. Consequently $\mathbb{E}G + M|B_1|/4 \ge 2gn\sqrt{D_N}$, and the displayed loss is at most $gn\sqrt{D_N}/2$. The convex function G is Lipschitz with constant at most $h = \frac{1}{2}\sqrt{TK_L D_N}$. Talagrand's inequality [11] gives

$$\mathbb{P}(G \le \mathbb{E}G - a - 8\sqrt{\pi}h) \le 4e^{-a^2/(16h^2)}.$$

The mean-to-median shift follows by integrating the median tail. The bound on n makes $8\sqrt{\pi}h \le gn\sqrt{D_N}/4$; set $a = gn\sqrt{D_N}/2$. □

Proof of Theorem C.2. Correctness is Lemma C.1. Fix a run and a uniform x, and write $u = 2x - \mathbf{1}$ and I for its creators. In the sequential case, L is the transpose of the leading square block of R, so $\|L\| \le \|R\|$. In the round case, each row block is $N_{B_j \times B_{<j}}$, whence $\|L\|^2 \le (\ell - 1)\|N\|^2$. The common gap estimate applies in both cases.

Accumulated gap. Given the earlier variables, each creator indicator is a fair coin and Δ_t is fixed; within a round all preferences were fixed before it started. For $Y_t = [t \in I]\min(|\Delta_t|, M)$, Azuma–Hoeffding gives

$$\mathbb{P}\left(\sum_t Y_t < \tfrac{1}{2}\sum_{t \le T} \min(|\Delta_t|, M) - a\right) \le e^{-2a^2/(TM^2)}.$$

Use $a = gn\sqrt{D_N}/8$. The gap lemma and the slack bound imply $S' = \sum_{t \in I}(|\Delta_t| - s_t) \ge gn\sqrt{D_N}/8$ outside these two exceptional events.

Sequential comparisons. Here the sum tested by the algorithm is

$$S' + Q, \qquad Q = -\tfrac{1}{2}\sum_w u_w \sum_{t \in I, t < w} u_t N_{tw}. \tag{27}$$

The inner coefficient is predictable, and its squared sum is at most $\|R^{\mathsf{T}}(u_t[t \in I])_{t \le T}\|^2 \le KD_N T$. The conditional cosh bound gives $\mathbb{P}(Q \le -b) \le e^{-2b^2/(KD_N T)}$. For $b = gn\sqrt{D_N}/16$, a passing assignment belongs to an exceptional event. Their total probability is at most $6e^{-c_* n}$ with the stated sequential constant (which conservatively retains the earlier bound).

Comparisons within rounds. The tested sum is $S' + Q_{\text{later}} + \sum_{j \le \ell} Q_j$, where

$$Q_{\text{later}} = -\tfrac{1}{2}\sum_w u_w \sum_{t \in I : \rho(t) < \rho(w)} u_t N_{tw}, \qquad Q_j = -\tfrac{1}{2}\sum_{t \in I \cap B_j} \sum_{w \in B_j \setminus \{t\}} u_t N_{tw} u_w.$$

For a later round the coefficients are fixed from earlier rounds. Each block of coefficients has squared norm at most KD_NT, and there are at most ℓ such blocks, including the suffix. Thus

$$\mathbb{P}(Q_{\text{later}} \le -b) \le e^{-2b^2/(\ell K D_N T)}.$$

Take $b = gn\sqrt{D_N}/32$.

Condition on the rounds preceding B_j. Put $c_t = 2p_t - 1$ and $\beta_t = -c_t u_t$. These are independent fair signs, and $t \in I$ precisely when $\beta_t = 1$. The matrix $M^{(j)}_{tw} = c_t c_w N_{tw}$ on B_j has zero diagonal and norm at most $\sqrt{KD_N}$. Direct expansion gives

$$Q_j = -\tfrac{1}{4}\beta^{\mathsf{T}} M^{(j)} \beta - \tfrac{1}{4}\langle M^{(j)}\mathbf{1}, \beta\rangle.$$

Both $\|M^{(j)}\mathbf{1}\|^2$ and $\|M^{(j)}\|_F^2$ are at most $|B_j|KD_N$. For $e = gn\sqrt{D_N}/(64\ell)$, the linear term is below $-e$ with probability at most $e^{-8e^2/(|B_j|KD_N)}$. Hanson–Wright [12] bounds the corresponding quadratic probability by

$$2\exp\left[-c_{\mathrm{HW}} \min\left\{\frac{16e^2}{|B_j|KD_N}, \frac{4e}{\sqrt{KD_N}}\right\}\right].$$

A union bound therefore gives $\sum_j Q_j \ge -gn\sqrt{D_N}/32$ outside these events. On the remaining assignments the tested sum is at least $(1/8 - 1/32 - 1/32)gn\sqrt{D_N} > 0$. Using $|B_j| \le T \le \tau n$, the total passing probability is at most $(6 + 3\ell)e^{-c_* n}$. Lemma 4.4 finishes both time bounds. Runs and candidates are streamed, so space is polynomial. □

C.3 Projection grids and discovering the low-rank part

Corollary C.4 (Polynomial projection grids)**.** *Suppose $\|v_k\|_2 \le 1$ and use $\mathcal{Y}_k = \delta\mathbb{Z} \cap [-\|v_k\|_1 - \delta, \|v_k\|_1 + \delta]$, where $\delta > 0$ is rational. Then*

$$|\mathcal{Y}| \le (2\sqrt{n}/\delta + 3)^r, \qquad \sum_{t \le T} s_t \le \sum_k |\lambda_k|(\delta\sqrt{n} + 1).$$

Thus Theorem C.2 applies whenever its norm and gap conditions hold and the last expression is at most $gn\sqrt{D_N}/8$.

Proof. Cauchy–Schwarz gives $\|v_k\|_1 \le \sqrt{n}$. Also $\sum_t s_t = \sum_k |\lambda_k|(\delta\|v_k\|_1 + \|v_k\|_2^2)$. □

The vectors need not be exact eigenvectors. One deterministic construction uses a rational, approximately diagonalizing orthogonal matrix, takes its r largest diagonal entries in absolute value, and uses the corresponding columns. Appendix D gives a polynomial-bit construction and a random-model guarantee. All pruning comparisons use the rational vectors actually computed, so numerical approximation affects the running bound, not exactness.

D Spectral preprocessing for hidden communities

D.1 A discovery lemma for the stochastic block model

Lemma D.1 (Discovering a bounded-rank mean)**.** *Fix a rank bound r and variance ratio $\chi \ge 1$. Suppose the integer coefficients J_{uv} are independent over $u < v$,*

$$|J_{uv} - \mathbb{E}J_{uv}| \le B, \quad 0 < \nu \le \operatorname{Var} J_{uv} \le \chi\nu,$$

where we take $B \ge 1$. Suppose their means are the off-diagonal entries of a symmetric matrix $\bar{A}$ with $\operatorname{rank}\bar{A} \le r$, $|\bar{A}_{uu}| \le B_0$ and $\|\bar{A}\| \le \mu n$. For a sufficiently large $C = C(r, \chi)$, if

$$n\nu \ge C(B^2 \log n + B_0^2 + \mu^2),$$

there is a deterministic exact polynomial-space algorithm, given the fixed class parameters r, χ but not $\bar{A}$ or ν, B, B_0, μ, taking $O^*(2^{(1-\gamma)n})$ time with probability $1 - o(1)$, for a fixed $\gamma > 0$. Arbitrary units, even chosen after the binary draw, are allowed.

We give the computational detail separately so that the general theorem requires only rational comparisons, with no exact-real arithmetic oracle.

Rational approximate diagonalization. For a rational symmetric matrix J and rational $\zeta > 0$, one can compute a rational orthogonal Q and rational diagonal D with $\|J - QDQ^{\mathsf{T}}\|_F \le \zeta$ in time polynomial in the input bit length, n, and $\log(1/\zeta)$. Here a slow Jacobi implementation suffices. For completeness, let $A = Q^{\mathsf{T}}JQ$ and let $E = \sum_{i \ne j} A_{ij}^2$. An exact plane rotation eliminating a largest entry A_{ij} reduces E by $2A_{ij}^2 \ge 2E/[n(n-1)]$. Approximate that rotation by a rational orthogonal plane rotation using

$$\cos\theta = \frac{1 - t^2}{1 + t^2}, \qquad \sin\theta = \frac{2t}{1 + t^2}, \qquad t \in \mathbb{Q}.$$

The eliminating angle can be chosen in $[-\pi/4, \pi/4]$, so its parameter $t = \tan(\theta/2)$ is approximated on a bounded interval by elementary root approximation. With rotation error $O(\zeta^2/(n^2(1+\|J\|_F)^2))$, the error in the decrease of E is at most $\zeta^2/[2n(n-1)]$. Therefore while $E > \zeta^2$, it contracts by at least a fixed multiple of $1/n^2$ per step. After $O(n^2 \log(2 + \|J\|_F/\zeta))$ steps, set $D = \operatorname{diag}(A)$. Every rotation is exactly orthogonal and rational; products and matrix updates are performed exactly. The bit lengths grow at most polynomially in the number of rotations and their polynomial bit lengths. This proves the claimed computational bound without an eigenvalue-separation assumption.

Let P retain the r entries of D of largest absolute value and the corresponding columns of Q. It has rank at most r and a rational orthonormal representation. Since QDQ^{T} is within ζ of J, the variational characterization of singular values gives

$$\|J - P\| \le \sigma_{r+1}(J) + 2\zeta. \tag{28}$$

This is sufficient; individual eigenvectors need not be close to any particular exact eigenbasis.

Proof of Lemma D.1. Put $X = \operatorname{offdiag}(J - \bar{A})$ and fix two prefix rounds of sizes $\lfloor n/4 \rfloor$ and $\lfloor n/2 \rfloor - \lfloor n/4 \rfloor$. Take $\zeta = n^{-4}$ in the preceding construction. Write $E_0 = \|X\| + B_0$, so $J = \bar{A} + X - \operatorname{diag}(\bar{A})$ and $\sigma_{r+1}(J) \le E_0$. Thus

$$\|J - P\| \le E_0 + 2\zeta, \qquad \|N\| \le 2E_0 + 4\zeta, \quad N = \operatorname{offdiag}(J - P).$$

For $H = \bar{A} - P$, we have $\operatorname{rank} H \le 2r$ and $\|H\| \le 2E_0 + 2\zeta$, hence $\|H\|_F \le \sqrt{2r}(2E_0 + 2\zeta)$.

The symmetric version of the random-matrix bound of Bandeira and van Handel [9], symmetrization and convex concentration as in Appendix B, imply $\|X\| \le 3\sqrt{\chi n\nu}$ with probability $1 - o(1)$ when $n\nu \ge CB^2 \log n$ with sufficiently large C. The leading symmetrized norm constant is $2\sqrt{2} < 3$; the $O(B\sqrt{\log n})$ term is absorbed by increasing C. Bernstein's inequality and a union bound also give, on the same event,

$$.99n(n-1)\nu \le \|X\|_F^2 \le 1.01\chi n(n-1)\nu, \qquad \sum_{s \in B_1} X_{ts}^2 \ge .99|B_1|\nu \quad (t \in B_2).$$

Here $B_0^2 \le n\nu/C$, and $B \ge 1$ ensures that the numerical error ζ is negligible on this scale. Because $N = X + \operatorname{offdiag} H$, these bounds and the triangle inequality show, for fixed r, χ and sufficiently large n,

$$.9n\nu \le D_N \le 1.1\chi n\nu, \qquad \|N\|^2 \le 100\chi D_N.$$

Furthermore the correction to the sum of cross-round row norms is at most

$$\sqrt{|B_2|}\,\|H\|_F = O(n\sqrt{r\chi\nu}) = o(n^{3/2}\sqrt{\nu}).$$

The uncorrected sum is at least $|B_2|\sqrt{.99|B_1|\nu}$. Consequently the gap condition holds with $g = 1/(256\sqrt{\chi})$, with room to spare, and $K = 100\chi$.

Let $\Lambda = \sum_{k\le r}|\lambda_k|$ for the retained diagonal entries. They are Rayleigh quotients of J, so $\Lambda \le r\|J\| \le r(\mu n + E_0)$. Increasing $C = C(r,\chi)$ and then n gives $\Lambda \le gn\sqrt{D_N}/64$. If $\Lambda = 0$, take $P = 0$ and no grid. Otherwise choose a dyadic rational δ within a factor two below $g\sqrt{nD_N}/(64\Lambda)$, and use Corollary C.4. The comparison slack is at most $gn\sqrt{D_N}/32$ and

$$|\mathcal{Y}| \le \left(\frac{256\Lambda}{g\sqrt{D_N}} + 3\right)^r = n^{O(r)}.$$

This proves the promised bound by the round case of Theorem C.2. Constants g, K can be replaced by slightly weaker rational bounds without affecting any conclusion.

The algorithm uses only J, the fixed class parameters r, χ, and its computed D_N, Λ. To define it on all inputs, use exhaustive search when $D_N = 0$; otherwise cap δ from below by n^{-2}. The cap keeps the number of runs polynomial on arbitrary inputs and is inactive on the probability event above for sufficiently large n. Every resulting grid still covers all projections, so these choices preserve exactness also outside that event. No planted mean, variance parameter, or spectral cutoff has to be supplied. □

E Development record and proof priorities

The following record keeps the first draft auditable while the research continues. It is intended to be removed or shortened before submission.

1. The bias theorem is consolidated from `max2sat_biased_and_direction.tex`; the elementary random-sign theorem from `max2sat_early_enumeration.tex`. Both are proved in this draft. The bias proof counts occurrences with multiplicity and includes the exact pair reductions and complementary-pair extension discussed in `max2sat_results_and_ideas.tex` and `max2sat_explicit_class.tex`. Corollary 3.8 states the dense frequency-gap consequence directly for the original input. Its simple-input statement also permits unit clauses, using the effective gap; the same scale-based proof covers the anchored partitioning example in Section 3.8. The representative-optimum argument now explicitly separates preferred values from actual child traversal order.

2. The group-count theorem comes from `max2sat_redundancy_and_group_counts.tex`. The average-defect extraction and explicit detection charge incorporate the subsequent discussion.

3. The general block certificate builds on `max2sat_prebranch_constraints.tex`. The bounded-exception block theorem generalizes the conflict-graph argument of `max2sat_near_twin_pairs.tex`; its proof is included here.

4. Theorem 4.3 and the round case of Theorem C.2 incorporate `max2sat_weighted_rounds.tex`. The former uses rational inverse-square-root weights, costing a factor four in its conservative exponent constant. The latter and the sequential count argument share one algorithm and one gap lemma. Max-Cut and the random models are corollaries. The Max-Cut proof includes integer round sizes and the bounded-degree range. Projection discovery uses rational orthogonal rotations and a fixed rank bound rather than a cutoff requiring unknown sampling parameters. The diagonal mean bound is charged

explicitly. The main statement is now Theorem 5.2 for the stochastic block model, including sparse and almost complete regimes; the more general bounded-rank assertion is a supporting lemma in Appendix D. The elementary specialization sets $\nu = 4v_{\min}$ and $B = B_0 = \mu = 2$.

5. Proposition B.1 retains Proposition 12 of `max2sat_count_guessing.tex` in an appendix. Its stronger unit-clause adversary and vanishing sampling probabilities distinguish it from the arbitrary-support weighted corollary. The previously sharper almost-regular constant is retained separately in Appendix A.

6. The regional comparison rule and bounded-cache proposition are recent additions. Their coverage invariant, retention policy, unweighted score range, and variable order are explicit. The literal-gap proof now includes an explicit count of assignments made by SIMPLIFY, as well as the loss of their contributions to the branching-gap sum. Additional simplifications in other algorithms still require their own analysis.

7. The immediate research priority is a quantitative bridge from poor cancellation to a recoverable family of approximate block comparisons or aggregate counts. Excluding exact short relations alone is not that bridge.

8. The literature review must compare class guarantees and memory bounds, including earlier exact algorithms on random inputs, before making any priority or novelty claim. Constants should be optimized only after the class statements and algorithmic model are settled.

Finite checks. The exact-arithmetic scripts retain the following reproducible checks. `verify_dense_max2sat_blocks.py`: 36 integral quadratic instances, 638 conditional block certificates, and 1,280 affine and regional comparisons, checked against exhaustive optimization. The objectives admit unweighted Max 2-SAT encodings with repeated clauses and units. Earlier enumeration checks cover 44 instances, 1,408 prefix counts and 65,024 assignment identities. `verify_dense_max2sat_random_models.py`: 19 pair distributions, 36 formulas with units, 39,936 count-shift identities, 50,688 comparisons, and optimum preservation for the original and adjusted gains. `verify_dense_max2sat_rounds.py`: 18 instances, 2,688 round-sum identities, 17,280 projection-error bounds, optimum preservation, and 36 rational Jacobi rotations. `verify_dense_max2sat_sbm.py`: 34 pair moments, 18 block-mean ranks, 8,640 supplied-group gain identities over 1,344 assignments, 35 optimum-survival cases, and seven clique optimizations with fields. Its models include unequal groups and dense, sparse, almost complete, and deterministic probabilities. Separately, the floating-point script `explore_weighted_rounds.py` loses no optimum on 200 weighted instances and 120 instances each with count and spectral grids. These finite consistency checks do not establish the asymptotic bounds.

References

[1] R. Williams. A new algorithm for optimal 2-constraint satisfaction and its implications. *Theoretical Computer Science*, 348(2–3):357–365, 2005. Author manuscript.

[2] A. S. Kulikov and K. Kutzkov. New bounds for MAX-SAT by clause learning. *Computer Science–Theory and Applications (CSR 2007)*, LNCS 4649, pp. 194–204, 2007. Institutional publication record.

[3] A. Golovnev and K. Kutzkov. New exact algorithms for the 2-constraint satisfaction problem. *Theoretical Computer Science*, 526:18–27, 2014. Author manuscript.

[4] M. B. Hastings. Local maxima and improved exact algorithm for MAX-2-SAT. 2016. arXiv:1610.07100.

[5] P. Beame, R. Impagliazzo, T. Pitassi and N. Segerlind. Memoization and DPLL: formula caching proof systems. *Proceedings of the 18th IEEE Conference on Computational Complexity*, pp. 225–236, 2003. Author manuscript.

[6] N. Cesa-Bianchi, C. Gentile, F. Vitale and G. Zappella. A correlation clustering approach to link classification in signed networks. *Proceedings of the 25th Annual Conference on Learning Theory*, PMLR 23, pp. 34.1–34.20, 2012. Conference article.

[7] R. Miyashiro and T. Matsui. Semidefinite programming based approaches to the break minimization problem. Technical Report METR 2003–28, University of Tokyo, 2003. Technical report.

[8] T. Erlebach, A. Hall and T. Schank. Classifying customer-provider relationships in the Internet. TIK Report 145, ETH Zürich, 2002. Institutional report.

[9] A. S. Bandeira and R. van Handel. Sharp nonasymptotic bounds on the norm of random matrices with independent entries. *Annals of Probability*, 44(4):2479–2506, 2016. Author manuscript.

[10] S. Boucheron, G. Lugosi and P. Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013.

[11] M. Talagrand. Concentration of measure and isoperimetric inequalities in product spaces. *Publications Mathématiques de l'IHÉS*, 81:73–205, 1995. Journal article.

[12] M. Rudelson and R. Vershynin. Hanson–Wright inequality and sub-gaussian concentration. *Electronic Communications in Probability*, 18(82):1–9, 2013. Author manuscript.

[13] A. E. Brouwer and W. H. Haemers. *Spectra of Graphs*. Springer, 2012. Authors' book page.

[14] K. Tikhomirov and P. Youssef. The spectral gap of dense random regular graphs. *Annals of Probability*, 47(1):362–419, 2019. Author manuscript.

[15] E. Abbe. Community detection and stochastic block models: recent developments. *Journal of Machine Learning Research*, 18(177):1–86, 2018. Journal article.

[16] *Count guessing: aligned structure plus a pseudo-random remainder*. Companion working note, October 2026. Local source: `max2sat_count_guessing.tex`.