Max 2-SAT Below 2^n on Unaligned Formulas: Theorem and Proof Sketch

8 October 2026

Working notes

Abstract

We define an explicit class of 2-CNF formulas by three conditions on the quadratic polynomial that counts satisfied clauses. Every formula in the class is solved exactly by a deterministic polynomial-space algorithm in time $O^*(c^n)$, with a constant $c < 2$ that depends only on a parameter of the class and not on the number of clauses. The class contains almost all formulas of density at least $C \log n$ and excludes the aligned formulas that defeat single-flip pruning. The algorithm is the search of [3]: pruning clauses and the sum check, along a fixed order. The proof is the argument of [4] for a fixed formula: random signs are replaced by the conditions of the class, and the only randomness left is the random partial assignment of Knuth's identity. This note states the theorem and sketches the proof; the details have not been written out.

1 The class and the theorem

Let F be a 2-CNF formula over $x_1, \dots, x_n$. Its clauses have one literal or two literals on distinct variables; several clauses on the same pair of variables are allowed, but no binary clause occurs twice. A literal on x_v has sign $s = +1$ if it is x_v and $s = -1$ if it is $\neg x_v$. For an assignment write $a_v = +1$ if x_v is true and $a_v = -1$ otherwise.

Definition 1 (Coefficients)**.** 1. The *parity matrix* of F is the symmetric $n \times n$ integer matrix $J = (J_{uv})_{u,v=1}^{n}$ with zero diagonal, $J_{vv} = 0$, and, for $u \neq v$,

$$J_{uv} = \sum_{\text{binary clauses on } \{x_u, x_v\}} s_u s_v,$$

where s_u, s_v are the signs of the two literals of the clause. So $J_{uv} = 0$ if no clause contains both x_u and x_v.

2. The *imbalance vector* of F is $h = (h_1, \dots, h_n) \in \mathbb{Z}^n$, where h_v is the sum of the signs of x_v over its binary clauses plus twice the sum of its signs over its unit clauses.

3. L is the strictly lower triangular part of J: the $n \times n$ matrix with $L_{uv} = J_{uv}$ for $u > v$ and $L_{uv} = 0$ for $u \leq v$. It depends on the numbering of the variables.

4. $d = \frac{1}{n} \sum_{u,v} J_{uv}^2$, the mean of the squared row norms of J.

A binary clause $(\ell_u \vee \ell_v)$ is satisfied with indicator $\frac{3}{4} + \frac{1}{4} s_u a_u + \frac{1}{4} s_v a_v - \frac{1}{4} s_u s_v a_u a_v$, and a unit clause (ℓ_v) with indicator $\frac{1}{2} + \frac{1}{2} s_v a_v$. Hence the number of satisfied clauses is

$$\mathrm{sat}_F(a) = c_0 + \frac{1}{4} \sum_{v} h_v a_v - \frac{1}{4} \sum_{u<v} J_{uv} a_u a_v,$$

where c_0 is $\frac{3}{4}$ times the number of binary clauses plus $\frac{1}{2}$ times the number of unit clauses. So J and h are the coefficients of the score, and every simplification that changes the score by a constant leaves them unchanged. For example, $(x \vee y)$ and $(x \vee \neg y)$ together score $1 + [x]$: they give $J_{xy} = 0$ and add 2 to h_x, exactly as the unit clause (x) does. A complementary pair $(x \vee y), (\neg x \vee \neg y)$ gives $J_{xy} = 2$ and adds nothing to h. If every pair carries at most one clause, J is the signed adjacency matrix with entries ± 1 and $d = 2m/n$ is the average degree.

Definition 2 (Unaligned formulas)**.** For $K \geq 1$, a formula is *K-unaligned* if

$$\|J\| \leq K\sqrt{d}, \qquad \|L\| \leq K\sqrt{d}, \qquad \sum_v h_v^2 \leq K\, nd,$$

where $\|\cdot\|$ is the spectral norm.

Theorem 3. *For every $K \geq 1$ there is $\varepsilon_K > 0$, for instance $\varepsilon_K = 10^{-5}K^{-4}$, such that Max 2-SAT on K-unaligned formulas can be solved by a deterministic algorithm in time $O^*(2^{(1-\varepsilon_K)n})$ using polynomial space, for every number of clauses.*

The algorithm is the same for all K; it does not need to know K or to test membership in the class.

Remark 4 (About the class). The class has the following properties.

1. Membership can be checked in polynomial time.
2. Random formulas lie in the class with small K (Table 1). If each pair carries a clause with probability $d/(n-1)$, with uniform signs, and $d \geq C \log n$, then the formula is 3-unaligned with probability $1 - o(1)$. This follows from norm bounds for random matrices with independent entries [5] and from Chernoff bounds for $\sum_v h_v^2$.
3. The aligned formulas of [1, 2] lie far outside. In the circulant formulas $C_{n,c}$ all entries of J equal 1, and in Max-Cut on K_m all equal 2, so $\|J\|$ is of order d and K would have to be of order $\sqrt{d}$.
4. The numbering of the variables enters only through L; the algorithm branches in this order.
5. Sparse random formulas with d below about $\log n$ are not K-unaligned for a fixed K, because a few variables of high degree make $\|J\|$ large. For them only the average-case result of [4] is available.
6. The definition is spectral; the algorithm is not.

2 The algorithm

Preprocessing. Exact pair reductions [1] bring F into canonical form: unit clauses, kept with multiplicity, and on each pair of variables nothing, one binary clause, or one complementary pair. They change the score by a constant, so J, h, L and d do not change.

The search [3]. Branch on $x_1, \ldots, x_n$ in this order. At a node where $x_1, \ldots, x_{t-1}$ are assigned, let Δ_t be the effective gap of x_t: the sum of its signs over its remaining binary clauses plus twice the sum over its unit clauses. Let D_t be the number of its remaining binary clauses, ℓ its preferred literal (x_t if $\Delta_t \geq 0$, else $\neg x_t$), and $p = (D_t - |\Delta_t|)/2$. The child in which ℓ is false receives the pruning clause $\langle N_t \leq p - 1 \rangle$. Here N_t counts the true

	d	$\|J\|/\sqrt{d}$	$\|L\|/\sqrt{d}$	$\sum_v h_v^2/(nd)$
random, one clause per pair	30	2.02	1.67	1.02
	100	2.03	1.66	1.01
	999	2.00	1.64	0.94
random, pairs and unit clauses	189	2.00	1.66	0.38
circulant $C_{401,100}$	400	20.0	12.8	0.00
Max-Cut on K_{400}	1596	20.0	12.7	0.00

Table 1: The three parameters of Definition 2; random formulas have $n = 1000$. In the fourth row, 30% of the interacting pairs are complementary pairs of random type, and each variable has up to two random unit clauses.

literals among $\neg m$ for each binary clause $(\ell \vee m)$ and m for each binary clause $(\neg\ell \vee m)$, with multiplicity; a complementary pair contributes the same literal twice. The bad value of a variable in a pruning clause is the value that makes its literal true. A node is cut when a pruning clause has more true literals than it allows, or when the *sum check* fires:

$$\sum_{w \text{ free}} \min\{m_w^0, m_w^1\} > \sum_c k_c,$$

over the live pruning clauses that are not vacuous. Here k_c is the remaining allowance of c, and m_w^b counts, with multiplicity, the clauses in which b is the bad value of w.

Soundness and cost. Each branching adds a clause to one child only, and the flip identity holds for canonical formulas with multiplicities. So the path-order argument of [1, Proposition 20] shows that an optimal assignment survives, and the sum check is implied by the clauses [2]. Let π_t be the probability that a uniformly random $\sigma \in \{0,1\}^n$ survives the checks at depths $1, \dots, t$. By Knuth's identity [6, 3], the search has at most $1 + 2n \max_t 2^t \pi_t$ nodes. It uses polynomial space and polynomial time per node.

3 Proof sketch

Fix a K-unaligned formula; all probabilities are over the uniform σ. Put $a = 2\sigma - 1$ and $T = (1-\delta)n$, with $\delta = \delta(K)$ chosen in Step 6.

Step 1: one depth. $\max_t 2^t \pi_t \le \max\{2^T, 2^n \pi_T\}$, because π_t is non-increasing. It suffices to show $\pi_T \le 6e^{-cn}$.

Step 2: the race. Run σ through the search without cutting and keep every pruning clause created up to depth T. Put $Z_T = S_T + Q_T - U_T$, with

$$S_T = \sum_c \Big(\frac{|\Delta_c|}{2} + 1\Big), \quad Q_T = \sum_c \Big(j_c - \frac{e_c}{2}\Big), \quad U_T = \frac{1}{2} \sum_{w>T} |m_w^1 - m_w^0|.$$

Here Δ_c is the gap at the creation of c, e_c is the number of its assigned support variables, and j_c is the number of those that took their bad value, both with multiplicity. By the race identity and the remark on vacuous clauses [3, Lemma 2, Remark 3], whose algebra is unchanged by multiplicities, a σ that survives depth T has $Z_T \le 0$.

Step 3: the race in terms of J and h.

Lemma 5 (Representation)**.** *Let C be the set of $s \le T$ at which σ_s is the non-preferred value, and $y_s = a_s$ for $s \in C$, $y_s = 0$ otherwise. Then*

(a) $\Delta_t = h_t - \sum_{u<t} J_{tu} a_u$, and $t \in C$ if and only if $a_t = -\operatorname{sign} \Delta_t$ (with $\operatorname{sign} 0 = +1$);

(b) the balance $m^1_w - m^0_w$ of a variable x_w, taken when x_w is assigned (or at depth T if $w > T$), equals $\sum_{s<\min\{w,T+1\}} J_{sw} y_s$;

(c) $S_T = \sum_{t\in C}(|\Delta_t|/2 + 1)$, $Q_T = \frac{1}{2}\sum_{t\le T} a_t (Ly)_t$ and $U_T = \frac{1}{2}\sum_{w>T} |\sum_{s\le T} J_{sw} y_s|$.

Why. After $a_1, \dots, a_{t-1}$ are fixed, the coefficient of a_t in the score polynomial is $\frac{1}{4}(h_t - \sum_{u<t} J_{tu} a_u)$. The remaining formula is canonical, and there the same coefficient is $\frac{1}{4}\Delta_t$; this gives (a). For (b), let x_s create a clause and x_w be in its support. A case distinction shows that the bad value of x_w, written as ± 1 and taken with its multiplicity, is $J_{sw} a_s$. For a single clause this is ± 1. For a complementary pair both clauses give the same bad value, σ_s for the pair $(x \vee y), (\neg x \vee \neg y)$ and $1 - \sigma_s$ for the other type, which is $J_{sw} a_s$ with $J_{sw} = \pm 2$. Part (c) follows: Q changes at step t by $\frac{1}{2} a_t (m^1_{x_t} - m^0_{x_t})$. All parts were checked numerically, including complementary pairs and unit clauses (Section 5).

Step 4: uncertainty, for every σ. By (c) and Cauchy–Schwarz, using $\|y\|^2 \le T$,

$$U_T \le \tfrac{1}{2}\sqrt{n-T}\,\|J\|\,\|y\| \le \tfrac{1}{2}\sqrt{T(n-T)}\,\|J\| \le \tfrac{1}{2} K \sqrt{\delta}\, n\sqrt{d}.$$

Step 5: the martingale. By (c), Q changes at step t by $\pm\frac{1}{2}(Ly)_t$. The sign is a_t, which is independent of $\sigma_1, \dots, \sigma_{t-1}$, while $(Ly)_t$ is determined by them. The squared steps sum to at most $\frac{1}{4}\|L\|^2\|y\|^2 \le \frac{1}{4}K^2 dn$ for every σ. Azuma's argument gives $\Pr[Q_T \le -\kappa] \le \exp(-2\kappa^2/(K^2 dn))$.

Step 6: the shortfall. Let $G = \sum_{t\le T} |\Delta_t|$.

(a) *Coin flips.* Given $\sigma_1, \dots, \sigma_{t-1}$, the gap Δ_t is fixed and $t \in C$ with probability $\frac{1}{2}$. So $X_t = [t \in C]\,|\Delta_t|/2 \ge 0$ has conditional mean $|\Delta_t|/4$ and conditional second moment $\Delta_t^2/8$. Moreover $\sum_t \Delta_t^2 \le 2\|h\|^2 + 2\|L\|^2 n \le 2(K + K^2)nd$ by (a). The one-sided bound for nonnegative increments [7, 4] gives $S_T \ge \frac{1}{4}G - \alpha$, except with probability $\exp(-\alpha^2/(2V))$, where $V = (K + K^2)nd/4$.

(b) *Mean.* By (a), Δ_t is a fixed number plus the symmetric sum $-\sum_{u<t} J_{tu} a_u$. So by symmetry and Szarek's inequality [8], $\mathbb{E}|\Delta_t| \ge (W_t/2)^{1/2}$ with $W_t = \sum_{u<t} J_{tu}^2$. Every row of J has $\sum_u J_{tu}^2 \le \|J\|^2 \le K^2 d$, so $\sum_{t\le T} W_t \ge \frac{1}{2}nd - \delta n K^2 d$ and $W_t \le K^2 d$. Hence

$$\mathbb{E}G \ge \sum_{t\le T} \sqrt{W_t/2} \ge \frac{\sum_t W_t}{\sqrt{2}\,K\sqrt{d}} \ge \frac{\frac{1}{2} - \delta K^2}{\sqrt{2}\,K}\, n\sqrt{d}.$$

(c) *Concentration.* G is a convex function of $a \in [-1, 1]^n$, Lipschitz with constant $\sqrt{T}\|L\| \le K\sqrt{nd}$. By Talagrand's inequality [9], $\Pr[G \le \mathbb{E}G - 15K\sqrt{nd} - \eta] \le 4\exp(-\eta^2/(16K^2 nd))$. Here $15K\sqrt{nd}$ bounds the distance between the mean and the median.

Step 7: the constants. Take $\delta = 1/(128K^4)$. Steps 4 and 6 give a margin

$$A = \frac{\frac{1}{2} - \delta K^2}{4\sqrt{2}\,K} - \frac{K}{2}\sqrt{\delta} \;\geq\; \frac{0.0428}{K}$$

in units of $n\sqrt{d}$. Take $\alpha = \kappa = \eta = \frac{1}{4}An\sqrt{d}$. Outside the three exceptional events,

$$Z_T \;\geq\; \tfrac{1}{4}(\mathbb{E}G - 15K\sqrt{nd} - \eta) - \alpha - \kappa - U_T \;\geq\; \tfrac{7}{16}An\sqrt{d} - 4K\sqrt{nd} \;>\; 0$$

for $n \geq n_0(K)$. The three events have probabilities at most $e^{-A^2n/(16K^2)}$, $4e^{-A^2n/(256K^2)}$ and $e^{-A^2n/(8K^2)}$. So $\pi_T \leq 6e^{-cn}$ with $c = A^2/(256K^2) \geq 7 \cdot 10^{-6}K^{-4}$. By Step 1 the search has $O^*(2^{(1-\varepsilon)n})$ nodes, with $\varepsilon = \min\{\delta, c/\ln 2\} \geq 10^{-5}K^{-4}$. □

4 What has to be written out

1. The race identity and the remark on vacuous clauses with multiplicities, and Lemma 5 in full. All are checked numerically, but only the version with one clause per pair is written [3, 4].

2. The exact form of Talagrand's inequality used in Step 6(c), with its constants.

3. The claim of Remark 4(2) on random formulas, with the constants of [5].

The constants are not optimized. The measured saving of the same search on random formulas is about $0.5n$ [3], many orders of magnitude more than ε_K. Three questions stay open: whether $\|L\|$ can be replaced by $\|J\|$ (it is at most $O(\log n)$ times larger, and that factor would cost too much); whether the class can be widened to sparse formulas by discounting high-degree variables; and whether repeated clauses (weights) can be admitted.

5 Checks

`explore_explicit_class.py`, each mode in a few seconds:

- `identities`: 128,826 checks on random canonical formulas (one clause per pair, complementary pairs, unit clauses) and on $C_{21,6}$, $C_{31,9}$, Max-Cut on K_{12} and K_{16}. Covered: the gap formula and the bad-value formula of Lemma 5, the race identity with multiplicities, the formula for U_T, and $Z_t \leq 0$ on every surviving prefix.

- `params`: Table 1.

- `constant`: Step 7 for $K = 1, \ldots, 4$, giving $\varepsilon_K K^4 \approx 1.04 \cdot 10^{-5}$ to $1.10 \cdot 10^{-5}$.

References

[1] Max 2-SAT and the domination rule: results and ideas so far. Working note, 5 October 2026 (`max2sat_results_and_ideas.tex`).

[2] Max 2-SAT with pruning clauses: pair decisions, sums, and a measure. Working note, 6–7 October 2026 (`max2sat_pairs_sums_measure.tex`).

[3] A simplified plan for Max 2-SAT: one search, one count, one Key Lemma. Working note, 8 October 2026 (`max2sat_plan_key_lemma.tex`).

[4] The Key Lemma for random signs. Working note, 8 October 2026 (`max2sat_key_lemma_proof.tex`).

[5] A. S. Bandeira and R. van Handel. Sharp nonasymptotic bounds on the norm of random matrices with independent entries. *The Annals of Probability* 44(4):2479–2506, 2016.

[6] D. E. Knuth. Estimating the efficiency of backtrack programs. *Mathematics of Computation* 29(129):121–136, 1975.

[7] A. Maurer. A bound on the deviation probability for sums of non-negative random variables. *Journal of Inequalities in Pure and Applied Mathematics* 4(1), Article 15, 2003.

[8] S. J. Szarek. On the best constants in the Khinchin inequality. *Studia Mathematica* 58(2):197–208, 1976.

[9] M. Talagrand. Concentration of measure and isoperimetric inequalities in product spaces. *Publications Mathématiques de l'IHÉS* 81:73–205, 1995.