The Key Lemma for Random Signs

8 October 2026

Working notes

Abstract

We prove the Key Lemma of the plan note [2]. Consider the search of [2]: a fixed variable order, pruning clauses, and the sum check. Let π_t be the probability that a uniformly random partial assignment survives depth t. For any graph with independent uniform signs we bound $\mathbb{E}\,\pi_T$ by four explicit terms, each exponentially small in n when a simple margin of the graph is of order $n\sqrt{d}$. Two consequences follow. On random formulas $R(n,d)$ of every density $d_0 \le d \le n-1$, the search solves Max 2-SAT in time $O^*(2^{(1-\gamma)n})$ for every $\gamma < 0.0236$, with probability $1-o(1)$. On every d-regular graph with random signs, in any order, every $\gamma < 0.0132$ works. The proof looks at the signs of a variable only when the search reaches it. Then the three terms of the race identity of [2] are sums of fair coins that have not been looked at, and the inequalities of Szarek, Azuma and McDiarmid finish the proof. The constant is far below the measured saving of about $0.5n$. Random signs are used in exactly three places, and these three places sharpen the worst-case target.

1 Statement

We use the search of [2, Section 2]. The formula is canonical with $B = 1$, and the order $x_1, \dots, x_n$ is fixed in advance. Branching on x_t gives the non-preferred child the pruning clause $\langle L_t \le p-1 \rangle$ over the free neighbours of x_t. A node is cut when a clause is violated (this includes $p = 0$) or when the sum check fires. Let $\pi_t(F)$ be the probability that a uniform $\sigma \in \{0,1\}^n$ survives the checks at depths $1, \dots, t$; the plan note calls it S_t, but here S denotes the shortfall. By Knuth's identity [2, Proposition 1], the search has at most $1 + 2n \max_t 2^t \pi_t$ nodes.

Graph notation. G is the graph of F, with one edge per clause. For a variable x_t, D_t^+ counts its neighbours after it in the order (the support of its pruning clause) and D_t^- its neighbours before it. For a depth T, $e_{\le T} = \sum_{t \le T} D_t^-$ is the number of edges among $x_1, \dots, x_T$ and e_T the number of edges between them and the rest. For a free variable x_w, $w > T$, $\deg_T(w)$ is its number of neighbours among $x_1, \dots, x_T$. *Random signs* on G means that each clause $(\ell_u \vee \ell_v)$ gets independent uniform signs for ℓ_u and for ℓ_v. $R(n,d)$ is the random graph $G(n, d/(n-1))$ with random signs.

Theorem 1. *Let G be a graph on n vertices, $1 \le T \le n$, and put*

$$\mu_G = \frac{1}{4\sqrt{2}} \sum_{t \le T} \sqrt{D_t^+ + 2D_t^-}, \qquad U_G = \frac{1}{2} \sum_{w > T} \sqrt{\deg_T(w)}, \qquad V_G = \sum_{t \le T} \Big(\frac{A_t}{8} + \frac{\sqrt{A_t}}{2} + \frac{1}{2} \Big),$$

where $A_t = D_t^+ + 4D_t^-$. If $\alpha, \xi, \kappa, \upsilon \ge 0$ and $\alpha + \xi + \kappa + \upsilon < \mu_G - U_G$, then for random signs on G

$$\mathbb{E}\,\pi_T \le e^{-\alpha^2/(2V_G)} + e^{-16\xi^2/e_{\le T}} + e^{-2\kappa^2/e_{\le T}} + e^{-2\upsilon^2/e_T}.$$

For $R(n,d)$ and for d-regular graphs, at a suitable depth T, the margin $\mu_G - U_G$ is of order $n\sqrt{d}$ while V_G, $e_{\le T}$ and e_T are $O(nd)$ (Section 4). Each term is then $e^{-\Omega(n)}$, independently of d. The theorem bounds a single depth; that is enough.

Lemma 2 (One depth suffices)**.** *For every T, $\max_t 2^t\pi_t \le \max\{2^T,\, 2^n\pi_T\}$.*

Proof. $2^t\pi_t \le 2^t \le 2^T$ for $t \le T$, and $\pi_t \le \pi_T$ for $t > T$. □

Corollary 3 (Key Lemma)**.** *For every $\gamma < 0.0236$ there is d_0 such that for all n and all $d_0 \le d \le n-1$, with probability $1 - o(1)$ over $F \in R(n,d)$,*

$$\max_{0\le t\le n} 2^t\pi_t(F) \le 2^{(1-\gamma)n}.$$

Hence the search solves Max 2-SAT on $R(n,d)$ in time $O^*(2^{(1-\gamma)n})$ and polynomial space, with probability $1-o(1)$, at every density.

Corollary 4 (Regular graphs)**.** *For every $\gamma < 0.0132$ there is d_0 such that for every d-regular graph G with $d \ge d_0$ and every order of its vertices, with probability $1 - e^{-\Omega(n)}$ over random signs on G, $\max_t 2^t\pi_t \le 2^{(1-\gamma)n}$.*

So the graph does not matter; only the signs do. This is the form we want for the worst case (Section 6).

2 Looking at signs late

The proof rests on one observation in clause terms: when the search reaches x_t, it has never used the signs of x_t. We state it for the *uncut process*, which runs σ through all steps, creates every clause and keeps it. On the event that σ survives depth T, the uncut process and the search agree up to depth T.

For a pruning clause c and a variable x_w of its support, the *bad value* of x_w in c is the value that makes the literal of x_w in c true; j_c in [2, Lemma 2] counts the support variables that took their bad value. Concretely, let c belong to x_s with preferred literal ℓ, and let $(\ell_s \vee \ell_w)$ be the clause of x_s and x_w. If $\ell_s = \ell$, the literal of x_w in c is $\neg\ell_w$: since ℓ is false in this child, the clause is falsified exactly when ℓ_w is false. If $\ell_s = \neg\ell$, the literal is ℓ_w, since a flip of x_s back to ℓ does not lose the clause if ℓ_w is true.

Lemma 5 (Late signs)**.** *Let H_t consist of the graph, the signs of $x_1,\dots,x_t$ in all their clauses, and $\sigma_1,\dots,\sigma_t$.*

(a) H_t determines, for every $s \le t$, the effective gap Δ_s, the preferred value of x_s, whether x_s creates a clause, its allowance, and the values S_s and Q_s.

(b) Given H_{t-1}, the signs of x_t and the value σ_t are independent and uniform.

(c) Let c be the clause of x_s, and x_w in its support. Given H_s, the bad value of x_w in c is a function of the sign of x_w in its clause with x_s, and the two signs give the two bad values.

Proof. (a) The effective gap of x_s sums over the clauses of x_s. A clause with a later variable is still a 2-clause and contributes ± 1, by the sign of x_s. A clause with an earlier x_u contributes ± 2, by the sign of x_s, if the literal of x_u is false (the clause is now a unit clause on x_s), and 0 if it is true. Whether it is true depends on the sign of x_u and on σ_u. So Δ_s depends only on the signs of $x_1,\dots,x_s$ and on $\sigma_1,\dots,\sigma_{s-1}$. The preferred value follows. A clause is created if σ_s is the other value, and its allowance is $p_s - 1 = (D_s^+ - |\Delta_s|)/2 - 1$. S_s adds $|\Delta_u|/2 + 1$

over the created clauses. Q_s adds $[\sigma_u \text{ bad in } c] - \frac{1}{2}$ over the assigned support variables x_u of the created clauses c, and by (c) these bad values depend on the signs of x_u, $u \le s$. (b) All signs and all values are independent, and H_{t-1} contains no sign of x_t. (c) This is the case distinction above: ℓ and ℓ_s are determined by H_s, and the literal of x_w in c is $\neg\ell_w$ or ℓ_w. □

Example 6. Let x_1 have the clauses $(x_1 \vee x_2)$, $(x_1 \vee x_3)$, $(x_1 \vee \neg x_4)$ and $(\neg x_1 \vee x_5)$, and no others. Then $\Delta_1 = 1 + 1 + 1 - 1 = 2$, $D_1^+ = 4$, $p = (4-2)/2 = 1$, and the preferred literal is x_1. The child $x_1 = 0$ receives

$$P_1 = \langle \neg x_2 + \neg x_3 + x_4 + x_5 \le 0 \rangle,$$

with bad values $x_2 = 0$, $x_3 = 0$, $x_4 = 1$, $x_5 = 1$. Each bad value is read off from the sign of x_w in its clause with x_1. The uncut process needs it only when x_w is assigned, or at depth T through U_T. When the search reaches x_2, its clause with x_1 has become the unit clause (x_2). It adds $+2$ to Δ_2, and the same positive sign of x_2 makes $x_2 = 0$ its bad value in P_1. So the preferred value of x_2 tends to be good for P_1. Such correlations exist between the three terms below, and we never need independence between them: we add their failure probabilities. Had $\sigma_1 = 1$, the clause $(x_1 \vee x_2)$ would be satisfied and would add 0 to Δ_2.

In words, Lemma 5 gives three sources of unseen fair coins.

1. At step t, the gap Δ_t is ± 1 over the 2-clauses of x_t plus ± 2 over its unit clauses, with unseen signs.
2. At step t, the bad values of x_t in the pruning clauses that contain it are unseen fair coins.
3. At depth T, the bad values of the free variables are unseen fair coins.

3 The three estimates

By [2, Lemma 2, Remark 3], applied to the set of all clauses created up to depth T, a partial assignment that survives depth T has

$$Z_T = S_T + Q_T - U_T \le 0,$$

with $S_T = \sum_c (|\Delta_c|/2 + 1)$, $Q_T = \sum_c (j_c - a_c/2)$ and $U_T = \frac{1}{2}\sum_{w>T} |b_w|$, all taken over the uncut process. The three lemmas below bound the three terms.

Lemma 7 (Shortfall)**.** $\Pr[S_T < \mu_G - \alpha - \xi] \le e^{-\alpha^2/(2V_G)} + e^{-16\xi^2/e_{\le T}}$.

Proof. Write $S_T = \sum_{t \le T} X_t$ with $X_t = [x_t \text{ creates a clause}]\,(|\Delta_t|/2 + 1) \ge 0$. Let r_t be the number of unit clauses on x_t, i.e. of earlier neighbours whose literal in the shared clause is false. Given H_{t-1}, Δ_t is a sum of D_t^+ unseen signs with weight 1 and r_t unseen signs with weight 2. A clause is created exactly when σ_t is the non-preferred value, which happens with probability $\frac{1}{2}$ independently of Δ_t. By Szarek's inequality [8], $\mathbb{E}|\sum_i a_i \varepsilon_i| \ge (\sum_i a_i^2)^{1/2}/\sqrt{2}$ for independent uniform signs ε_i, and therefore

$$\mathbb{E}[X_t \mid H_{t-1}] = \tfrac{1}{2}\big(\tfrac{1}{2}\mathbb{E}[|\Delta_t| \mid H_{t-1}] + 1\big) \ge \frac{\sqrt{D_t^+ + 4r_t}}{4\sqrt{2}} + \frac{1}{2} =: \mu_t.$$

Also $\mathbb{E}[X_t^2 \mid H_{t-1}] = \frac{1}{2}\mathbb{E}[(|\Delta_t|/2+1)^2 \mid H_{t-1}] \le \frac{1}{2}(\sqrt{D_t^+ + 4r_t}/2 + 1)^2 \le v_t$, where $v_t = A_t/8 + \sqrt{A_t}/2 + \frac{1}{2}$ uses $r_t \le D_t^-$.

Lower tail. For $X \ge 0$ and $\lambda \ge 0$, $e^{-\lambda X} \le 1 - \lambda X + \lambda^2 X^2/2$, so $\mathbb{E}[e^{-\lambda X_t} \mid H_{t-1}] \le \exp(-\lambda\mu_t + \lambda^2 v_t/2)$. Hence $\exp(-\lambda \sum_{s\le t}(X_s - \mu_s) - \frac{\lambda^2}{2}\sum_{s\le t} v_s)$ is a supermartingale. With $\lambda = \alpha/V_G$,

$$\Pr\Big[S_T \le \sum_{t\le T} \mu_t - \alpha\Big] \le e^{-\alpha^2/(2V_G)}.$$

This is the standard one-sided bound for nonnegative increments [6]; it needs no upper bound on X_t.

The sum of the μ_t. Given the graph, the events "the literal of x_u in its clause with x_t is false", $u < t \le T$, are $e_{\le T}$ independent fair coins: the literal of x_u is false when its sign disagrees with σ_u, and the signs of x_u are independent and uniform, also given σ_u. r_t counts the coins of the clauses between x_t and earlier variables, and changing one coin changes $\sum_t \sqrt{D_t^+ + 4r_t}$ by at most 2. By McDiarmid's inequality [7], this sum falls more than x below its mean with probability at most $e^{-x^2/(2e_{\le T})}$. Its mean is at least $\sum_t(\sqrt{D_t^+ + 2D_t^-} - 1)$: the elementary inequality $\sqrt{y} \ge \sqrt{\mu} + \frac{y-\mu}{2\sqrt{\mu}} - \frac{(y-\mu)^2}{2\mu^{3/2}}$ for $y \ge 0$ gives $\mathbb{E}\sqrt{Y} \ge \sqrt{\mathbb{E}Y} - \mathrm{Var}(Y)/(2(\mathbb{E}Y)^{3/2})$, and here $\mathrm{Var}(D_t^+ + 4r_t) = 4D_t^- \le 2\,\mathbb{E}Y$. The terms $\frac{1}{2}$ in μ_t pay for the -1, because $\frac{1}{4\sqrt{2}} < \frac{1}{2}$. So $\sum_t \mu_t \ge \mu_G - x/(4\sqrt{2})$ except with probability $e^{-x^2/(2e_{\le T})}$. Taking $x = 4\sqrt{2}\,\xi$ proves the lemma. □

Lemma 8 (Martingale)**.** $\Pr[Q_T < -\kappa] \le e^{-2\kappa^2/e_{\le T}}$.

Proof. At step t, Q changes by $\sum_c([\sigma_t \text{ bad in } c] - \frac{1}{2})$ over the $m_t \le D_t^-$ created clauses that contain x_t; m_t is determined by H_{t-1}. Given H_{t-1}, the bad values of x_t in these clauses are independent fair coins (Lemma 5(b),(c)), and σ_t is independent of them. So the indicators $[\sigma_t \text{ bad in } c]$ are independent fair coins, and the step of Q is $\frac{1}{2}$ times a sum of m_t independent signs. So Q is a walk with fair steps $\pm\frac{1}{2}$, at most $e_{\le T}$ of them, and $\mathbb{E}[e^{-\lambda(Q_t - Q_{t-1})} \mid H_{t-1}] = \cosh(\lambda/2)^{m_t} \le e^{\lambda^2 D_t^-/8}$. Azuma's argument [3, 5] with $\lambda = 4\kappa/e_{\le T}$ gives the claim. □

Lemma 9 (Uncertainty)**.** $\Pr[U_T > U_G + \upsilon] \le e^{-2\upsilon^2/e_T}$.

Proof. Condition on H_T. A free variable x_w lies in $m_w \le \deg_T(w)$ created clauses, one for each neighbour among $x_1, \dots, x_T$ that created a clause, and b_w is the sum of ± 1 over them according to its bad values. By Lemma 5(c) these are unseen fair coins, independent over all pairs (c, w). So $\mathbb{E}[|b_w| \mid H_T] \le \sqrt{m_w} \le \sqrt{\deg_T(w)}$ by Jensen's inequality, and $\mathbb{E}[U_T \mid H_T] \le U_G$. U_T depends on at most e_T unseen signs, and changing one of them changes U_T by at most 1. McDiarmid's inequality [7] finishes the proof. □

Proof of Theorem 1. If σ survives depth T, the search and the uncut process agree up to depth T, so $Z_T \le 0$. Outside the three events of Lemmas 7–9,

$$Z_T \ge (\mu_G - \alpha - \xi) - \kappa - (U_G + \upsilon) > 0.$$

So survival lies in the union of the three events, and averaging over the signs gives $\mathbb{E}\,\pi_T \le$ the sum of the four failure probabilities. □

4 The constants

Proof of Corollary 3. Let $p = d/(n-1)$, $T = \tau n$, and $\eta > 0$. In $G(n,p)$ the following hold with probability $1 - o(1)$ for $d \ge d_0(\eta)$.

- $e_{\le T} \le (1+\eta)\tau^2 nd/2$ and $e_T \le (1+\eta)\tau(1-\tau)nd$, by Chernoff bounds.

- $U_G \le \frac{1}{2}\sqrt{(n-T)e_T} \le (1+\eta)\frac{1}{2}(1-\tau)\sqrt{\tau}\, n\sqrt{d}$, by Cauchy–Schwarz.
- $V_G \le (1+\eta)(\tau + \frac{3}{2}\tau^2)nd/8$, since $\sum_{t\le T} A_t = e_T + 5e_{\le T}$ and $\sum_t \sqrt{A_t} \le \sqrt{T\sum_t A_t}$.
- $\mu_G \ge (1-\eta)\frac{1}{6\sqrt{2}}((1+\tau)^{3/2}-1)\, n\sqrt{d}$. The expected value of $D_t^+ + 2D_t^-$ is $(n+t-2)p$, so the mean of $\sum_t \sqrt{D_t^+ + 2D_t^-}$ is at least $\sum_t \sqrt{(n+t-2)p} - T$ by the inequality used in Lemma 7. One edge changes the sum by at most $1+\sqrt{2}$, so the Efron–Stein inequality [4] bounds its variance by $3.2\,nd$, and Chebyshev's inequality finishes.

Put $M(\tau) = \frac{1}{6\sqrt{2}}((1+\tau)^{3/2}-1) - \frac{1}{2}(1-\tau)\sqrt{\tau}$, measure $\alpha, \xi, \kappa, \upsilon$ in units of $n\sqrt{d}$, and let $\eta \to 0$. The four exponents of Theorem 1 become $k_i n$ times the square of the corresponding parameter, with

$$k_1 = \frac{4}{\tau + \frac{3}{2}\tau^2}, \quad k_2 = \frac{32}{\tau^2}, \quad k_3 = \frac{4}{\tau^2}, \quad k_4 = \frac{2}{\tau(1-\tau)}.$$

The best split of $M(\tau)$ makes all four exponents equal to $c(\tau) = M(\tau)^2/(\sum_i k_i^{-1/2})^2$. At $\tau = 0.976$ we get $M = 0.1976$ and $c = 0.01638$, so $\mathbb{E}\,\pi_T \le 4e^{-(c-\varepsilon)n}$ for typical graphs. By Markov's inequality, $\pi_T \le e^{-(c-2\varepsilon)n}$ except with probability $4e^{-\varepsilon n}$ over the signs. Lemma 2 gives $\max_t 2^t \pi_t \le 2^{(1-\gamma)n}$ for every $\gamma < \min\{1-\tau,\, c/\ln 2\} = \min\{0.024,\, 0.02364\}$. □

Proof of Corollary 4. In a d-regular graph $D_t^+ + 2D_t^- = d + D_t^- \ge d$ in every order, so $\mu_G \ge T\sqrt{d}/(4\sqrt{2})$. Moreover $U_G \le \frac{1}{2}(n-T)\sqrt{d}$, $V_G \le T(d/2 + \sqrt{d} + \frac{1}{2})$, $e_{\le T} \le Td/2$ and $e_T \le (n-T)d$. The same optimization, with $M(\tau) = \tau/(4\sqrt{2}) - \frac{1}{2}(1-\tau)$ and $k = (1/\tau,\, 32/\tau,\, 4/\tau,\, 2/(1-\tau))$, gives $c = 0.00915$ at $\tau = 0.986$ and $c/\ln 2 = 0.0132$. □

5 How far from the truth

	S_T	$\sum_t \mu_t$	μ_G	U_T	U_G	Q_T	$\mathbb{E}\lvert\Delta_t\rvert$ ratio
$n = 1000$, $d = 100$	0.253	0.235	0.190	0.026	0.047	−0.001	1.14
$n = 2000$, $d = 200$	0.249	0.222	0.191	0.027	0.048	−0.002	1.14
$n = 2000$, $d = 1000$	0.238	0.205	0.191	0.027	0.048	+0.001	1.13
$n = 1500$, complete	0.224	0.202	0.191	0.028	0.047	−0.002	1.14

Table 1: The terms at depth $T = 0.9n$ in units of $n\sqrt{d}$, averaged over four assignments; the bounds of the proof are μ_G and U_G. The last column is the average of $|\Delta_t|/\sqrt{(D_t^+ + 4r_t)/2}$, against 1 in Szarek's inequality.

Table 1 compares the proof with the process. Every bound holds with room, and the race is won by a wide margin: Z_T is between $0.19\,n\sqrt{d}$ and $0.23\,n\sqrt{d}$, where the proof uses $\mu_G - U_G = 0.143\,n\sqrt{d}$. The constant $\gamma \approx 0.02$ is small for three reasons.

1. Four separate tails, each with a second-moment rate, are added.
2. Szarek's constant $1/\sqrt{2}$ is used where the measured ratio is $\sqrt{4/\pi} \approx 1.13$ times larger.
3. Only one depth is used. The proof places the cut before depth $0.976n$, while the measured median cut is at $0.43n$ and the measured saving is $0.50n$ to $0.62n$ [2, Tables 1 and 2].

A single exponential moment of Z_t at a depth near $0.5n$ would remove the first and the third loss; the Gaussian heuristic of [2, Section 6] suggests a saving of about $0.2n$. We leave this; the point here is that the saving is independent of the density.

6 What the proof uses

Random signs enter only through Lemma 5, in three places.

(S) The signs of x_t in its 2-clauses and unit clauses make $|\Delta_t|$ large and the shortfall concentrated.

(Q) The bad values of x_t in the clauses that contain it make Q a walk with fair steps $\pm\frac{1}{2}$.

(U) The bad values of the free variables make $|b_w|$ of order $\sqrt{m_w}$ instead of m_w.

For a fixed formula, randomness comes only from σ. Two things remain. Q is still a martingale, but its step at x_t is $\pm|b_{x_t}|/2$. The shortfall still has mean at least $\sum_{t\le T}(\sqrt{D_t^-/2}/4 + \frac{1}{2})$, because each earlier neighbour of x_t satisfies or falsifies its literal by an independent fair coin, so Δ_t is a fixed number plus ±1 over the D_t^- earlier neighbours. What is lost is the size of the bad-value sums. The proof of Theorem 1 goes through verbatim for a fixed formula along every path on which

(W1) U_T is a small constant times $n\sqrt{d}$,

(W2) $\sum_{t\le T} b_{x_t}^2 = O(nd)$, and

(W3) $\sum_{t\le T} |\Delta_t|$ is at least a constant times $n\sqrt{d}$ and $\sum_{t\le T} \Delta_t^2 = O(nd)$,

provided these hold except with probability $e^{-\Omega(n)}$ over σ. (W1) and (W2) are cancellation of bad values, at the free and at the assigned variables. (W3) says that effective gaps keep their typical size $\sqrt{d}$; it makes the shortfall concentrate by the argument of Lemma 7. This replaces the single cancellation condition of [2, Section 7], which named only (W1); the claim there that the shortfall part holds for every formula is true only for its mean. Where (W1) or (W2) fails, many clauses agree about the values of many variables. That is the aligned structure of [1, Section 9], which is the case for aggregation.

7 Checks

`explore_key_lemma_proof.py` (a few seconds each):

- `inclusion`: on random formulas with $n \le 400$, all 51,041 pairs (σ, t) in which σ survives depth t have $Z_t \le 0$ for the uncut process, and the identity of [2, Lemma 2] holds exactly at every step.
- `terms`: Table 1.
- `constant`: the optimizations in the proofs of Corollaries 3 and 4.

References

[1] Max 2-SAT with pruning clauses: pair decisions, sums, and a measure. Working note, 6–7 October 2026 (`max2sat_pairs_sums_measure.tex`).

[2] A simplified plan for Max 2-SAT: one search, one count, one Key Lemma. Working note, 8 October 2026 (`max2sat_plan_key_lemma.tex`).

[3] K. Azuma. Weighted sums of certain dependent random variables. *Tôhoku Mathematical Journal* 19(3):357–367, 1967.

[4] B. Efron and C. Stein. The jackknife estimate of variance. *The Annals of Statistics* 9(3):586–596, 1981.

[5] W. Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association* 58(301):13–30, 1963.

[6] A. Maurer. A bound on the deviation probability for sums of non-negative random variables. *Journal of Inequalities in Pure and Applied Mathematics* 4(1), Article 15, 2003.

[7] C. McDiarmid. On the method of bounded differences. In *Surveys in Combinatorics 1989*, London Mathematical Society Lecture Note Series 141, pages 148–188. Cambridge University Press, 1989.

[8] S. J. Szarek. On the best constants in the Khinchin inequality. *Studia Mathematica* 58(2):197–208, 1976.