Count Guessing Makes Constraints Useful

Max 2-SAT for aligned blocks plus a pseudo-random remainder

11 October 2026

Abstract

The note on useful constraints and group counts [8] names two exits for the branching algorithm: learned constraints that give useful additional pruning, and compression of redundant structure. We combine them. If the quadratic part of the score is a few *count terms* plus a pseudo-random remainder, the algorithm first guesses the values of the counts. Every count term then becomes a field, the learned constraints come from the remainder only, and early enumeration [5, 4] prunes as for a pseudo-random formula. The guess costs a polynomial factor. We prove a deterministic theorem: such formulas are solved exactly in time $O^*(c^n)$, $c < 2$, and polynomial space, at every density, with arbitrary fields. The class contains the pseudo-random formulas of [4]. With probability at least $1 - 1/n$ it also contains random formulas whose pairs carry independent clause parities: with any parity bias, at every density $d \geq C \log n$, with any number of clauses per pair and with arbitrary unit clauses (Proposition 12). Random planted bisections and random formulas with a bounded number of aligned groups are members as well, together with their planted decompositions. The aligned family of [8, Section 4] is its degenerate case. The simplest members need no unknowns: the parity matrix minus its mean is pseudo-random. On strongly biased random formulas, plain early enumeration keeps a constant fraction of 10^{-4} to 10^{-3} of all assignments up to $n = 480$; with count guessing no sampled assignment survives.

1 The idea in clause terms

A learned constraint is redundant when its bad values agree with those of the constraints already present [8]. Agreement on a large scale has a simple source: many clauses whose parities follow one pattern. In the aligned family of [8, Section 4], every pair carries $(u \vee v)$ or $(\neg u \vee \neg v)$, and the score is a linear function minus $\binom{k}{2}$, where k is the number of true variables. All learned constraints then push the same way, and the algorithm of [8] instead sorts for every value of k.

Now add noise. Let a fraction of the pairs carry the opposite parity at random. The count term is still there, so the constraints still largely agree, and the summed constraint stays weak. But sorting no longer solves the problem either. The two exits combine. For a *guessed* final count k, the count term contributes to every single flip only a shift by a common field, which depends on k alone. Replace it by that shift. The learned constraints then come from the noise, which is pseudo-random, and they prune as in [4]. One of the $n + 1$ guesses is right, and in that run the optimum survives.

2 The class

Write the score of a 2-CNF formula on $x \in \{0, 1\}^n$ as

$$F(x) = c_0 + \sum_u h_u x_u + \sum_{u<v} J_{uv} x_u x_v.$$

Fix the order $1, \dots, n$ and a depth $T = \lfloor \tau n \rfloor$, $0 < \tau < 1$.

Definition 1 (Aligned decomposition)**.** Let $s_1, \dots, s_r \in \{-1, 0, 1\}^n$ and $\beta_1, \dots, \beta_r \in \mathbb{Q}$. Put $A_{uv} = \sum_k \beta_k s_k(u) s_k(v)$ for $u \neq v$, the *aligned part*, and $N = J - A$ off the diagonal, the *remainder*. The *counts* of x are $c_k(x) = \langle s_k, x \rangle$, and $\sigma = \sum_k |\beta_k|$.

If the s_k are indicators of blocks, the aligned part is constant inside each block; sums and differences of two indicators add a constant bias between two blocks. The counts are then the group counts of [8], and the number of runs below is governed by the number of distinct group counts.

Definition 2 (Pseudo-random remainder)**.** Let $D = \frac{1}{n} \sum_{u \neq v} N_{uv}^2$, let R be the $T \times n$ matrix with $R_{tw} = N_{tw}$ for $w > t$ and 0 otherwise, and $W_t = \sum_{j<t} N_{tj}^2$. For constants $K, g > 0$, the remainder is (K, g)-*pseudo-random* if

$$\|R\|^2 \leq K\, D, \tag{R1}$$

$$\sum_{t \leq T} \sqrt{W_t} \geq 4\sqrt{3}\, g\, n \sqrt{D}, \tag{R2}$$

$$\sigma \leq g\sqrt{D}/(8\tau). \tag{R3}$$

Condition (R1) is the aggregate pair-count condition of [4] for the remainder. Condition (R2) asks the earlier neighbours to carry weight, as in [3]. Condition (R3) compares the strength of the aligned part per pair with the scale $\sqrt{D}$ of the remainder; it holds as soon as D is large. No condition is placed on the fields h or on the aligned part otherwise.

Examples.

1. $r = 0$: the pseudo-random formulas of [4, 3].
2. *Mean-centred*: $r = 1$, $s = \mathbf{1}$, and β the mean of J_{uv} over all pairs $u \neq v$. Nothing is unknown, and all conditions can be checked in polynomial time (Corollary 9). Random formulas whose pairs carry independent clause parities with a common mean are members with probability at least $1 - 1/n$: with any parity bias, at every density $d \geq C \log n$, with any number of clauses per pair and with arbitrary unit clauses (Section 6, Corollary 13). Both standard models are of this kind: one clause per pair, and each of the four clauses of a pair present independently.
3. *Planted bisection*: $r = 1$, $s \in \{\pm 1\}^n$: pairs on the same side have one parity bias, pairs across the other. Random planted bisections, whose pairs carry independent parities with mean $\beta s(u) s(v)$, are members by Proposition 12(ii) when s is given.
4. *Few groups*: s_k the indicators of $r = O(1)$ groups and β_k a bias per pair of groups, written as a sum of terms of this form. Random formulas whose pairs carry independent parities with a mean that depends only on the groups of the two variables are members by Proposition 12(ii) when the groups are given.
5. *Degenerate*: $N = 0$. Then every run below optimizes a linear function, which is the sorting argument of [8, Section 4]. The theorem does not need this case.

3 The algorithm

For every vector $c = (c_1, \dots, c_r)$ of integers with $|c_k| \leq \|s_k\|_1$ run the following, and return the best assignment found over all runs.

1. Shift the fields: $h^c = h + \sum_k \beta_k c_k s_k$, and let $F'_c(x) = c_0 + \langle h^c, x\rangle + \sum_{u<v} N_{uv} x_u x_v$.

2. Branch on $x_1, \ldots, x_T$, visiting all 2^T prefixes. At step t let

$$\Delta_t = h^c_t + \sum_{j<t} N_{tj} x_j + \tfrac{1}{2} \sum_{w>t} N_{tw},$$

   the gain of $x_t = 1$ over $x_t = 0$ under F'_c when the later variables are set to $\frac{1}{2}$. The preferred value is $p_t = [\Delta_t \geq 0]$. If $x_t \neq p_t$, then t is a *creator*, and $\gamma_t(x)$ denotes the gain of flipping x_t to p_t under F'_c, a linear function of the other variables.

3. At each prefix, enumerate the completions satisfying the summed constraint

$$\sum_{t \text{ creator}} \gamma_t(x) \leq \sigma \cdot \#\{\text{creators}\}, \tag{1}$$

   with the output-sensitive generator of [5], and evaluate F on each.

There are at most $(2n+1)^r$ runs. Every run uses polynomial space, and its time is $O^*(2^T + W_c)$, where W_c is the number of assignments x that satisfy (1) for their own prefix [5, Lemma 2].

4 Exactness

Lemma 3 (Gains)**.** *Let $c(x) = c$. Flipping x_t from 0 to 1 changes F and F'_c equally. Flipping x_t from 1 to 0 changes F by $\sum_k \beta_k s_k(t)^2 \geq -\sigma$ more than it changes F'_c.*

Proof. The aligned part of F is $\sum_k \frac{\beta_k}{2}(c_k(x)^2 - \sum_u s_k(u)^2 x_u)$, since $x_u^2 = x_u$. Flipping x_t up changes it by $\sum_k \beta_k c_k s_k(t)$, the same as the shift in F'_c. Flipping x_t down changes it by $\sum_k \beta_k (s_k(t)^2 - c_k s_k(t))$, while the shift changes by $-\sum_k \beta_k c_k s_k(t)$. The remainder and h contribute equally to both. □

Lemma 4 (Exactness)**.** *The algorithm returns an optimal assignment.*

Proof. Add to F a generic perturbation $\langle \theta, x\rangle$ with $\|\theta\|_1$ below $1/q$, where q is a common denominator of all coefficients. Let x° be the unique maximizer of $F + \langle \theta, x\rangle$; it maximizes F. Consider the run $c = c(x^\circ)$, add the same perturbation to F'_c, and let t be a creator on the path of x°. If the perturbed gain of flipping x_t to p_t at x° were at least σ, Lemma 3 would make the perturbed gain under F nonnegative, contradicting uniqueness. Summing the strict inequalities over the creators and removing the perturbation gives $\sum_t \gamma_t(x^\circ) < \sigma \cdot \#\{\text{creators}\} + 1/q$. The left side minus $\sigma \cdot \#\{\text{creators}\}$ is a multiple of $1/q$, so it is at most zero. Hence x° satisfies (1) at its own prefix and is enumerated. □

The argument needs no property of N: the algorithm is exact for every decomposition, and the conditions of Definition 2 only bound its running time.

5 Running time

Fix a run c and let x be uniform; put $u = 2x - 1 \in \{\pm 1\}^n$. Write $H_t = h^c_t + \frac{1}{2}\sum_{w \neq t} N_{tw}$ and let L be the $T \times T$ matrix with $L_{tj} = N_{tj}$ for $j < t$. Then $\Delta_t = H_t + \frac{1}{2}(Lu)_t$, and L^{T} is the leading block of R, so $\|L\| \leq \|R\|$.

Lemma 5 (Passing)**.** *With I the set of creators, x satisfies (1) if and only if $S' + Q \leq 0$, where*

$$S' = \sum_{t \in I} (|\Delta_t| - \sigma), \qquad Q = -\tfrac{1}{2} \sum_{w=1}^{n} u_w \sum_{t \in I,\, t<w} u_t N_{tw}.$$

Proof. For a creator, $\gamma_t(x) = (2p_t - 1)(\Delta_t + \frac{1}{2}\sum_{w>t} N_{tw}u_w)$, $(2p_t - 1)\Delta_t = |\Delta_t|$ and $u_t = -(2p_t - 1)$. □

Lemma 6 (Gap profile with arbitrary fields)**.** *Assume* (R1), (R2) *and* $n \geq 256\pi\tau K/g^2$. *Let* $\alpha = \max\{2, \frac{9}{4}\tau\sqrt{K}/g\}$ *and* $M = \alpha\sqrt{KD}$. *Then*

$$\mathbb{P}\Big[\sum_{t\leq T} \min\{|\Delta_t|, M\} < \tfrac{3}{4}g\,n\sqrt{D}\Big] \leq 4e^{-g^2 n/(16\tau K)}.$$

Proof. Split $t \leq T$ into Small $= \{|H_t| \leq M/2\}$ and Large $= \{|H_t| > M/2\}$. Since $\|Lu\|^2 \leq \|L\|^2 T \leq KDT$, two deterministic facts hold.

- For $t \in$ Small, $|\Delta_t| - M \leq \frac{1}{2}|(Lu)_t| - \frac{M}{2}$, and $(a-b)_+ \leq a^2/(4b)$ gives $\min\{|\Delta_t|, M\} \geq |\Delta_t| - (Lu)_t^2/(8M)$. Summed over Small, the loss is at most $KDT/(8M) = T\sqrt{KD}/(8\alpha)$.
- For $t \in$ Large with $|(Lu)_t| \leq M/2$, $|\Delta_t| \geq M/4$. At most $4KDT/M^2 = 4T/\alpha^2$ indices have $|(Lu)_t| > M/2$.

So $\sum_t \min\{|\Delta_t|, M\} \geq G(u) - T\sqrt{KD}/(8\alpha) + \frac{M}{4}(|\text{Large}| - 4T/\alpha^2)$ with $G(u) = \sum_{t\in\text{Small}} |\Delta_t|$. Each $(Lu)_t$ is a symmetric sum of independent terms with variance W_t, so $\mathbb{E}|\Delta_t| \geq \mathbb{E}|\frac{1}{2}(Lu)_t| \geq \frac{1}{2}\sqrt{W_t/3}$ by symmetry and the fourth-moment bound $\mathbb{E}|Z| \geq \sqrt{\operatorname{Var} Z/3}$ [5]. For $t \in$ Large, $\frac{M}{4} \geq \frac{1}{2}\sqrt{W_t/3}$ because $W_t \leq \|L\|^2 \leq KD$ and $\alpha \geq 2$. With (R2),

$$\mathbb{E}G + \tfrac{M}{4}|\text{Large}| \;\geq\; \tfrac{1}{2}\sum_{t\leq T}\sqrt{W_t/3} \;\geq\; 2g\,n\sqrt{D},$$

and the two losses $T\sqrt{KD}/\alpha + T\sqrt{KD}/(8\alpha)$ are at most $\frac{1}{2}g\,n\sqrt{D}$ by the choice of α. The function G is convex in $u \in [-1,1]^T$ and Lipschitz with constant $\ell = \frac{1}{2}\sqrt{T}\|L\| \leq \frac{1}{2}\sqrt{TKD}$. Talagrand's inequality [10] gives $\mathbb{P}[G \leq \mathbb{E}G - s - 8\sqrt{\pi}\,\ell] \leq 4e^{-s^2/(16\ell^2)}$; the assumption on n makes $8\sqrt{\pi}\,\ell \leq \frac{1}{4}g\,n\sqrt{D}$. Take $s = \frac{1}{2}g\,n\sqrt{D}$. □

Proposition 7 (One run)**.** *Under* (R1)–(R3) *and* $n \geq 256\pi\tau K/g^2$, *a uniform* x *satisfies* (1) *for its own prefix with probability at most* $6e^{-c_* n}$, *where* $c_* = g^2/(64\tau\alpha^2 K)$. *The bound does not depend on the fields, hence not on the run.*

Proof. Shortfall. Let $Y_t = [t \in I]\min\{|\Delta_t|, M\} \in [0, M]$. Given $u_1, \ldots, u_{t-1}$, the value Δ_t is fixed and t is a creator with probability $\frac{1}{2}$. So $\mathbb{E}[Y_t \mid \cdot] = \frac{1}{2}\min\{|\Delta_t|, M\}$ and $\mathbb{E}[Y_t^2 \mid \cdot] \leq M^2/2$. The lower tail for nonnegative increments [5, 6] gives $\sum Y_t \geq \frac{1}{2}\sum_t \min\{|\Delta_t|, M\} - a$ except with probability $e^{-a^2/(TM^2)}$. Take $a = \frac{1}{8}g\,n\sqrt{D}$. By Lemma 6 and (R3), $S' \geq \sum Y_t - \sigma T \geq (\frac{3}{8} - \frac{1}{8} - \frac{1}{8})g\,n\sqrt{D}$ outside the two exceptional events.

Martingale. In the order $w = 1, \ldots, n$, the coefficient $\kappa_w = \sum_{t\in I, t<w} u_t N_{tw}$ of u_w in Q is fixed before u_w is drawn, and $\sum_w \kappa_w^2 = \|R^{\mathsf{T}} v\|^2 \leq KDT$ with $v_t = u_t[t \in I]$. The conditional cosh bound gives $\mathbb{P}[Q \leq -b] \leq e^{-2b^2/(KDT)}$; take $b = \frac{1}{8}g\,n\sqrt{D}$.

Outside the three events $S' + Q > 0$. The exponents are $g^2 n/(64\tau\alpha^2 K)$, $g^2 n/(16\tau K)$ and $g^2 n/(32\tau K)$, and the first is the smallest. □

Theorem 8 (Aligned blocks plus a pseudo-random remainder)**.** *Let* F *be a 2-CNF formula on* n *variables with an aligned decomposition of rank* r *whose remainder is* (K, g)*-pseudo-random, and let* $n \geq 256\pi\tau K/g^2$. *Given the decomposition, Max 2-SAT on* F *is solved exactly by a deterministic algorithm in polynomial space and time*

$$O^*\Big((2n+1)^r\big(2^{\tau n} + 2^n e^{-c_* n}\big)\Big), \qquad c_* = \frac{g^2}{64\tau\alpha^2 K}, \quad \alpha = \max\Big\{2, \frac{9\tau\sqrt{K}}{4g}\Big\}.$$

For $r = O(1)$ *this is* $O^*(c^n)$ *with* $c = 2^{\max\{\tau,\, 1 - c_*/\ln 2\}} < 2$, *independent of the density and of the fields.*

Proof. Lemma 4 gives exactness. Each run costs $O^*(2^T + W_c)$, and $W_c \le 2^n \cdot 6e^{-c_* n}$ by Proposition 7. □

Corollary 9 (Mean-centred formulas)**.** *Let β be the mean of J_{uv} over all pairs $u \ne v$ and $N = J - \beta(\mathbf{1}\mathbf{1}^{\mathsf{T}} - I)$. If N satisfies* (R1)–(R3) *with $\sigma = |\beta|$, then Max 2-SAT on F is solved in polynomial space and time $O^*(n(2^{\tau n} + 2^n e^{-c_* n}))$. All conditions can be verified in polynomial time; there are $n + 1$ runs, one for each number of true variables.*

Remark 10 (Constants and other decompositions). The constants are very small: for the complete parity-biased formulas with $\tau = \frac{1}{2}$, (R2) holds with $g \approx 0.034$ (Table 1). The measured behaviour is far better (Table 2). For a planted bisection s is unknown. When $|\beta| n \gg \|N\|$, rounding the top eigenvector of J recovers s except for $O(\|N\|^2/(\beta^2 n))$ coordinates, by the Davis–Kahan theorem. The misplaced coordinates change N by a matrix of rank at most two and norm $O(\|N\|)$. This recovery is sketched, not proved here.

6 Random formulas

The class of Definition 2 is deterministic. This section shows that random formulas belong to it with probability at least $1 - 1/n$, with any parity bias, at every density $d \ge C \log n$, and with arbitrary unit clauses. Recall that J_{uv} is the number of mixed clauses on the pair $\{u, v\}$ minus the number of equal clauses on it: $(u \vee v)$ and $(\neg u \vee \neg v)$ contribute -1, while $(u \vee \neg v)$ and $(\neg u \vee v)$ contribute $+1$.

Definition 11 (Independent pairs)**.** A random 2-CNF formula has *independent pairs* if the numbers J_{uv}, $u < v$, are independent random variables. Nothing else is assumed: the unit clauses, and which literals are negated in a clause of a given parity, may be arbitrary and may even depend on the J_{uv}.

Two models. In both, $\varepsilon \in [-1, 1]$ is the parity bias, and d is the expected number of binary clauses that contain a given variable.

(a) *One clause per pair.* Each pair independently receives no clause with probability $1 - p$. Otherwise it receives exactly one clause: its literal on u is positive or negative with probability $\frac{1}{2}$ each, and its literal on v has the same polarity with probability $\frac{1+\varepsilon}{2}$ and the opposite polarity otherwise. Here $d = p(n-1)$, $J_{uv} \in \{-1, 0, 1\}$, $\mathbb{E}J_{uv} = -\varepsilon p$ and $\operatorname{Var} J_{uv} = p(1 - \varepsilon^2 p)$. For $\varepsilon = 0$ these are the usual random formulas on $G(n, p)$; for $p = 1$ and $\varepsilon = 1$ they form the aligned family of [8, Section 4].

(b) *Independent clauses.* Each of the four clauses on each pair is present independently: $(u \vee v)$ and $(\neg u \vee \neg v)$ with probability $q(1+\varepsilon)$ each, $(u \vee \neg v)$ and $(\neg u \vee v)$ with probability $q(1-\varepsilon)$ each. A pair carries between 0 and 4 clauses, $d = 4q(n-1)$, $|J_{uv}| \le 2$ and $\mathbb{E}J_{uv} = -4q\varepsilon$; if $q(1 + |\varepsilon|) \le \frac{1}{2}$, then $\operatorname{Var} J_{uv} \ge 2q$. For $\varepsilon = 0$ and $q = \frac{1}{4}$ every clause is present with probability $\frac{1}{4}$.

Proposition 12 (Random formulas are pseudo-random)**.** *Let $0 < \tau < 1$, $\rho \ge 1$, $C = 10^9 \rho/\tau$ and $n \ge n_0 = 10^9 \rho^2/\tau^2$. Let F have independent pairs with $|J_{uv} - \mathbb{E}J_{uv}| \le B$ and $\nu \le \operatorname{Var} J_{uv} \le \rho\nu$ for all $u < v$, and take one of the following aligned parts.*

(i) Mean-centred*: $\mathbb{E}J_{uv} = \mu$ for all pairs, and $A = \beta(\mathbf{1}\mathbf{1}^{\mathsf{T}} - I)$ with β the mean of the J_{uv}, as in Corollary 9; put $\sigma_0 = |\mu|$.*

(ii) Expected*: a fixed aligned decomposition with $A_{uv} = \mathbb{E}J_{uv}$ for all $u \ne v$; put $\sigma_0 = \sigma$.*

If $n\nu \ge C(B^2 \log n + \rho\sigma_0^2)$, then with probability at least $1 - 1/n$ the remainder $N = J - A$ is (K, g)-pseudo-random with $K = 9\rho$ and $g = \tau^{3/2}/(12\sqrt{\rho})$. Moreover $n \ge 256\pi\tau K/g^2$.

Proof. Let $X_{uv} = J_{uv} - \mathbb{E}J_{uv}$. These are independent and centred, with $|X_{uv}| \le B$ and $\nu \le \mathbb{E}X_{uv}^2 \le \rho\nu$. Let $\delta = \beta - \mu$ in case (i) and $\delta = 0$ in case (ii); in both cases $N_{uv} = X_{uv} - \delta$ for $u \ne v$. Put $\eta = 1/100$. The assumption gives $B\sqrt{\log n} \le \sqrt{n\nu/C}$ and $\sigma_0 \le \sqrt{n\nu/(C\rho)}$. Three standard tools are used.

- *Bernstein's inequality* [2]: if $Y_1, \dots, Y_m$ are independent, $Y_i - \mathbb{E}Y_i \le \bar{c}$ and $\sum_i \operatorname{Var} Y_i \le \lambda$, then $\mathbb{P}[\sum_i (Y_i - \mathbb{E}Y_i) \ge s] \le \exp(-s^2/(2\lambda + 2\bar{c}s/3))$.
- *Talagrand's inequality* [10], as in Lemma 6: if f is convex and 1-Lipschitz in independent variables with values in $[-B, B]$, then $f \le \mathbb{E}f + 8\sqrt{\pi}\,B + s$ except with probability $4e^{-s^2/(16B^2)}$.
- *The norm bound of Bandeira and van Handel* [1, Theorem 3.1 and Corollary 3.2]: for a fixed $T \times n$ matrix (a_{tw}) and independent random signs ξ_{tw},

$$\mathbb{E}\|(\xi_{tw}a_{tw})\| \le (1+\vartheta)\Big(\sigma_1 + \sigma_2 + \frac{5\max_{t,w}|a_{tw}|\sqrt{\log n}}{\sqrt{\log(1+\vartheta)}}\Big), \qquad 0 < \vartheta \le \tfrac{1}{2},$$

  where σ_1 and σ_2 are the largest Euclidean norms of a row and of a column of (a_{tw}).

Each event below fails with probability at most n^{-3}, and there are at most $n + 4$ of them, so they all hold with probability at least $1 - 1/n$.

The mean and D. In case (i), $\binom{n}{2}\delta = \sum_{u<v} X_{uv}$, and Bernstein's inequality for both tails (with $\bar{c} = B$ and $\lambda = \binom{n}{2}\rho\nu$) gives $|\delta| \le \delta_0 = 6\sqrt{\rho\nu\log n}/n$: the assumption makes $\bar{c}s/3 \le \lambda$, and the exponent is then at least $4.5(1 - \frac{1}{n})\log n$. As $n \ge n_0$, $\delta_0 \le 10^{-3}\sqrt{\nu}$ and $n\delta_0 \le 10^{-3}\sqrt{n\rho\nu}$. Bernstein's inequality for the sum of the X_{uv}^2 (with $\bar{c} = B^2$ and $\lambda = \binom{n}{2}B^2\rho\nu$) gives $(1-\eta)n(n-1)\nu \le \sum_{u\ne v} X_{uv}^2 \le (1+\eta)n(n-1)\rho\nu$. Since $\sum_{u\ne v} X_{uv} = n(n-1)\delta$, we have $D = \frac{1}{n}\sum_{u\ne v} X_{uv}^2 - (n-1)\delta^2$, and hence

$$(1-2\eta)(n-1)\nu \le D \le (1+\eta)(n-1)\rho\nu. \tag{2}$$

Condition (R1). Let R_X be the $T \times n$ matrix with entries X_{tw} for $w > t$ and 0 otherwise, and U the 0/1 matrix with the same pattern. Then $R = R_X - \delta U$ and $\|U\| \le \|U\|_F \le n$. Let X' be an independent copy of X and $\hat{X} = X - X'$. By Jensen's inequality, $\mathbb{E}\|R_X\| \le \mathbb{E}\|R_X - R_{X'}\| = \mathbb{E}\|R_{\hat{X}}\|$. The entries of $\hat{X}$ are independent and symmetric, so $R_{\hat{X}}$ has the law of $(\xi_{tw}\hat{X}_{tw})$. The norm bound, applied conditionally on $\hat{X}$ with $\vartheta = \eta$ and $|\hat{X}_{tw}| \le 2B$, gives

$$\mathbb{E}\|R_X\| \le (1+\eta)(\mathbb{E}\sigma_1(\hat{X}) + \mathbb{E}\sigma_2(\hat{X})) + 102\,B\sqrt{\log n},$$

where $\sigma_1(\hat{X})^2$ and $\sigma_2(\hat{X})^2$ are the largest row sum and the largest column sum of the $\hat{X}_{tw}^2$ over the pattern. A row has fewer than n entries and a column at most $T \le \tau n$, each with $\hat{X}_{tw}^2 \le 4B^2$ and $\mathbb{E}\hat{X}_{tw}^2 \le 2\rho\nu$. Bernstein's inequality (with $\bar{c} = 4B^2$, $\lambda = 8nB^2\rho\nu$ and $s = 2\eta n\rho\nu$) and a union bound over the at most $2n$ rows and columns show that all row sums are at most $2(1+\eta)n\rho\nu$ and all column sums at most $2(\tau+\eta)n\rho\nu$, except with probability $2n\,e^{-\eta^2 n\rho\nu/(5B^2)} \le n^{-3}$. Every such sum is at most $4nB^2$, so $\mathbb{E}\sigma_1(\hat{X})^2 \le 2(1+2\eta)n\rho\nu$ and $\mathbb{E}\sigma_2(\hat{X})^2 \le 2(\tau+2\eta)n\rho\nu$. With $102/\sqrt{C} \le 0.004$,

$$\mathbb{E}\|R_X\| \le \Big(1.01\sqrt{2}(\sqrt{1.02} + \sqrt{\tau + 0.02}) + 0.004\Big)\sqrt{n\rho\nu} \le 2.89\sqrt{n\rho\nu}.$$

The norm $\|R_X\|$ is a convex function of the entries X_{tw} and is 1-Lipschitz for the Euclidean norm. Talagrand's inequality with $s = \eta\sqrt{n\rho\nu}$ gives $\|R_X\| \le \mathbb{E}\|R_X\| + 8\sqrt{\pi}\,B + \eta\sqrt{n\rho\nu}$ except with probability $4n^{-\eta^2 C/16} \le n^{-3}$. As $8\sqrt{\pi}\,B \le 10^{-3}\sqrt{n\rho\nu}$,

$$\|R\| \le \|R_X\| + n|\delta| \le 2.902\sqrt{n\rho\nu}, \qquad \|R\|^2 \le 8.43\,n\rho\nu \le 9\rho(1-2\eta)(n-1)\nu \le KD.$$

Condition (R2). Let $t_0 = \lceil \tau n/10 \rceil$. For $t > t_0$ the sum $V_t = \sum_{j<t} X_{tj}^2$ has mean at least $(t-1)\nu$, and Bernstein's inequality for its lower tail (with $\bar{c} = B^2$, $\lambda = (t-1)B^2\rho\nu$ and $s = \eta(t-1)\nu$) gives $V_t \geq (1-\eta)(t-1)\nu$ except with probability $e^{-\eta^2\tau n\nu/(30B^2\rho)} \leq n^{-3}$. Then $\sqrt{W_t} \geq \sqrt{V_t} - |\delta|\sqrt{t-1} \geq 0.993\sqrt{(t-1)\nu}$. Using $\sum_{s=t_0}^{T-1} \sqrt{s} \geq \frac{2}{3}((T-1)^{3/2} - t_0^{3/2})$, $T - 1 \geq \tau n - 2$ and $\tau n \geq 10^3$,

$$\sum_{t \leq T} \sqrt{W_t} \geq 0.993\sqrt{\nu} \cdot \tfrac{2}{3}((T-1)^{3/2} - t_0^{3/2}) \geq 0.635\, \tau^{3/2} n^{3/2} \sqrt{\nu}.$$

By (2), $4\sqrt{3}\, g\, n\sqrt{D} \leq \frac{\sqrt{3}}{3}\sqrt{1.01}\, \tau^{3/2} n^{3/2}\sqrt{\nu} \leq 0.581\, \tau^{3/2} n^{3/2}\sqrt{\nu}$.

Condition (R3). By (2), $g\sqrt{D}/(8\tau) \geq \sqrt{\tau}\sqrt{0.98(n-1)\nu}/(96\sqrt{\rho}) \geq 0.0102\sqrt{\tau n\nu/\rho}$. In case (ii), $\sigma = \sigma_0 \leq \sqrt{n\nu/(C\rho)} \leq 10^{-4}\sqrt{\tau n\nu/\rho}$. In case (i), $\sigma = |\beta| \leq \sigma_0 + \delta_0 \leq 0.002\sqrt{\tau n\nu/\rho}$.

Finally, $256\pi\tau K/g^2 = 256\pi \cdot 1296\, \rho^2/\tau^2 < 1.1 \cdot 10^6 \rho^2/\tau^2 \leq n$. □

Corollary 13 (Random formulas)**.** *Let $0 < \tau < 1$, $C = 10^9/\tau$ and $n \geq 10^9/\tau^2$. In each of the following cases the formula satisfies, with probability at least $1 - 1/n$, the conditions of Corollary 9 with $K = 9$ and $g = \tau^{3/2}/12$, whatever its unit clauses are.*

- *(a) Model (a) with $p \leq \frac{1}{2}$ and $d \geq 9C \log n$, for every $\varepsilon \in [-1, 1]$; and model (a) with $p = 1$ if $n(1-\varepsilon^2) \geq 5C \log n$.*
- *(b) Model (b) with $q(1 + |\varepsilon|) \leq \frac{1}{2}$ and $d \geq 34C \log n$, for every $\varepsilon \in [-1, 1]$.*

On such a formula the algorithm of Corollary 9 solves Max 2-SAT in polynomial space and time $O^*(c^n)$, with $c = 2^{\max\{\tau,\, 1 - c_*/\ln 2\}} < 2$ and $c_* \geq \tau^3/(5.5 \cdot 10^8)$. The algorithm is exact on every input; only its running time uses the probability.

Proof. All pairs have the same distribution, so Proposition 12(i) applies with $\rho = 1$ and $\nu = \operatorname{Var} J_{uv}$. In model (a), $B = 2$ and $\sigma_0 = |\varepsilon| p \leq p$. If $p \leq \frac{1}{2}$, then $\nu \geq p/2$ and $n\nu \geq d/2 \geq C(4 \log n + 1)$. If $p = 1$, then $\nu = 1 - \varepsilon^2$ and $\sigma_0 \leq 1$. In model (b), $B = 4$, $\sigma_0 = 4q|\varepsilon| \leq 2$ and $n\nu \geq 2qn \geq d/2 \geq C(16 \log n + 4)$. The running time is Theorem 8 with $\alpha = 81/\sqrt{\tau}$ and $c_* = \tau^3/(144 \cdot 64 \cdot 81^2 \cdot 9)$. □

Remark 14 (Constants and density). No attempt was made to optimize C and n_0, and c_* is tiny. The proved values $K = 9$ and $g = \tau^{3/2}/12 \approx 0.029$ (for $\tau = \frac{1}{2}$) are not far from the measured ones in Table 1, which `explore_count_guessing.py membership` produces. The condition $d \geq C \log n$ is needed for (R1): $\|R\|$ is at least the largest norm of a row or a column of R, which is of the order of the square root of the largest number of neighbours of a variable. For $d = o(\log n)$ some variables have far more than d neighbours, so K is unbounded. Densities between $\omega(1)$ and $C \log n$ would need a separate treatment of these few variables, for instance branching on them first.

7 Experiments

`explore_count_guessing.py exact` runs the algorithm on 60 formulas with $6 \leq n \leq 11$. They are complete parity-biased formulas, planted bisections, and random formulas with an arbitrary two-group decomposition. In every case the algorithm found the optimum, with $T = n/2$, and it evaluated 12,375 of the 46,400 assignments. `scaling` estimates the fraction of all assignments that early enumeration evaluates at $T = n/2$, by sampling (Table 2).

With bias $\varepsilon \leq 0.8$ plain early enumeration also falls below the sampling resolution from $n = 240$ on, consistent with a floor of size $e^{-\Theta(1/\varepsilon^2)}$. A heuristic for this floor: on a constant fraction of prefixes, the creators favour one value by $\Theta(\varepsilon n)$. The count term then adds to

model	n	ε	K	g	(R3) ratio
(a), $p = 1$	2000	0	2.42	0.034	3450
(a), $p = 1$	2000	0.5	2.44	0.034	0.66
(a), $p = 1$	2000	0.9	2.46	0.034	0.18
(a), $d = 5$	4000	1	3.45	0.028	12.5
(a), $d = 20$	4000	0.5	2.62	0.033	14.7
(a), $d = 80$	4000	1	2.47	0.034	3.74
(b), $q = \frac{1}{4}$	2000	0	2.39	0.034	1380
(b), $q = \frac{1}{4}$	2000	0.5	2.42	0.034	0.63
(b), $d = 20$	4000	0	2.68	0.032	3360
(b), $d = 20$	4000	1	2.61	0.033	7.26

Table 1: Measured constants of random formulas for the mean-centred decomposition, $\tau = \frac{1}{2}$, one formula per row: $K = \|R\|^2/D$, the largest g allowed by (R2), and the ratio of the right side of (R3) to σ (at least 1 is needed). Proposition 12 proves $K = 9$ and $g \approx 0.029$. For a dense, strongly biased formula (R3) needs a larger n; the ratio grows like $\sqrt{n(1-\varepsilon^2)}/|\varepsilon|$.

	$n = 60$	120	240	480
parity bias $\varepsilon = 0.9$, plain	$6.3 \cdot 10^{-4}$	$1.1 \cdot 10^{-3}$	$6 \cdot 10^{-5}$	$3.3 \cdot 10^{-4}$
with count guessing	$1.3 \cdot 10^{-4}$	0	0	0
parity bias $\varepsilon = 1$, plain	$6.3 \cdot 10^{-4}$	$1.4 \cdot 10^{-3}$	$2.5 \cdot 10^{-4}$	$8.3 \cdot 10^{-4}$
with count guessing	0	0	0	0
planted bisection, $\beta = 0.9$, plain	$5.6 \cdot 10^{-3}$	$9.4 \cdot 10^{-4}$	$2.5 \cdot 10^{-4}$	0
with count guessing	$7.5 \cdot 10^{-4}$	0	0	0

Table 2: Fraction of the 2^n assignments evaluated by early enumeration at $T = n/2$; 16,000 samples (6,000 for $n = 480$), so 0 means below about 10^{-4}. One formula per entry.

Q a product of this imbalance and the imbalance of the completion, which is as large as the shortfall with constant probability. The bound of [4] for plain early enumeration degenerates on these formulas, because the aligned part has norm $|\beta|n \gg \sqrt{D}$. Count guessing removes the count term from Q altogether.

8 Relation to the note on group counts

- *Group counts.* The solver of [8, Section 5] exploits exact repetition across one cut, with arbitrary structure on both sides, by enumerating the sides separately. Count guessing exploits aligned terms that are global but noisy, and hands the rest to pruning. The two combine: a run of Theorem 8 may end in the group-count solver instead of enumeration whenever its remaining formula has repeated profiles.

- *Useful constraints.* In a run, the learned constraints are those of F'_c. They differ from those of F exactly by the count term, which is what made them redundant. Guessing the counts is therefore a way of *making* constraints useful, not only of choosing among them.

- *One-literal-neighbourhood constraints.* The constraint $A_\ell \geq G_\ell + 1$ of [8, Section 2] excludes a fraction $\mathbb{P}[\mathrm{Bin}(d_2(\ell), \frac{1}{2}) \leq G_\ell]$ of the completions of its support. This fraction is non-negligible only when $G_\ell \geq d_2(\ell)/2 - O(\sqrt{d_2(\ell)})$, that is, for strongly biased

variables. In the regime of Theorem 8 the combined constraint is the one that carries the shortfall.

- *A by-product for near twins.* If every variable has an r-near-twin in the sense of [7], then a spanning star forest of the near-twin graph yields at least $n/3$ disjoint $(2r+1)$-near-twin pairs: two leaves of a star differ at most at the centre and at their own discrepancies. So [7, Theorem 3] gives time $O^*((2(3/4)^{1/(6r+6)})^n)$ for this natural class, at every density.

What is not covered. Aligned structure of high rank: a blow-up with groups of size L needs n/L group counts. The resulting $(L+1)^{n/L} = 2^{n\log_2(L+1)/L}$ runs are affordable only when $\log_2(L+1)/L$ is below the saving of a run, while for bounded L the plain criterion of [4] applies with K of order L. Between the two, and for aligned parts that are not low-rank at all, a different exit is needed. The decomposition must also be known, except in the mean-centred case. Random formulas are covered only at density $d \geq C \log n$ (Remark 14). Finally, an adaptive version would measure the runs with the soft count [9] rather than enumerate them.

References

[1] A. S. Bandeira and R. van Handel. Sharp nonasymptotic bounds on the norm of random matrices with independent entries. *Ann. Probab.* 44(4):2479–2506, 2016.

[2] S. Boucheron, G. Lugosi and P. Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013.

[3] *A checkable replacement for the gap-profile condition*, addendum, 8 October 2026 (`max2sat_counting_addendum.tex`).

[4] *Replacing random signs by neighborhood counts*, working note, 8 October 2026 (`max2sat_counting_conditions.tex`).

[5] *Max 2-SAT by early constraint-guided enumeration*, working note, 8 October 2026 (`max2sat_early_enumeration.tex`).

[6] A. Maurer. A bound on the deviation probability for sums of non-negative random variables. *J. Inequal. Pure Appl. Math.* 4(1), Article 15, 2003.

[7] *Near-twin neighborhoods as a structural exit*, working note, 8 October 2026 (`max2sat_near_twin_pairs.tex`).

[8] *Useful constraints or compression by group counts*, working note, 10 October 2026 (`max2sat_redundancy_and_group_counts.tex`).

[9] *The soft-count measure for adaptive branching*, working note, 8 October 2026 (`max2sat_soft_count.tex`).

[10] M. Talagrand. Concentration of measure and isoperimetric inequalities in product spaces. *Publ. Math. IHÉS* 81:73–205, 1995.